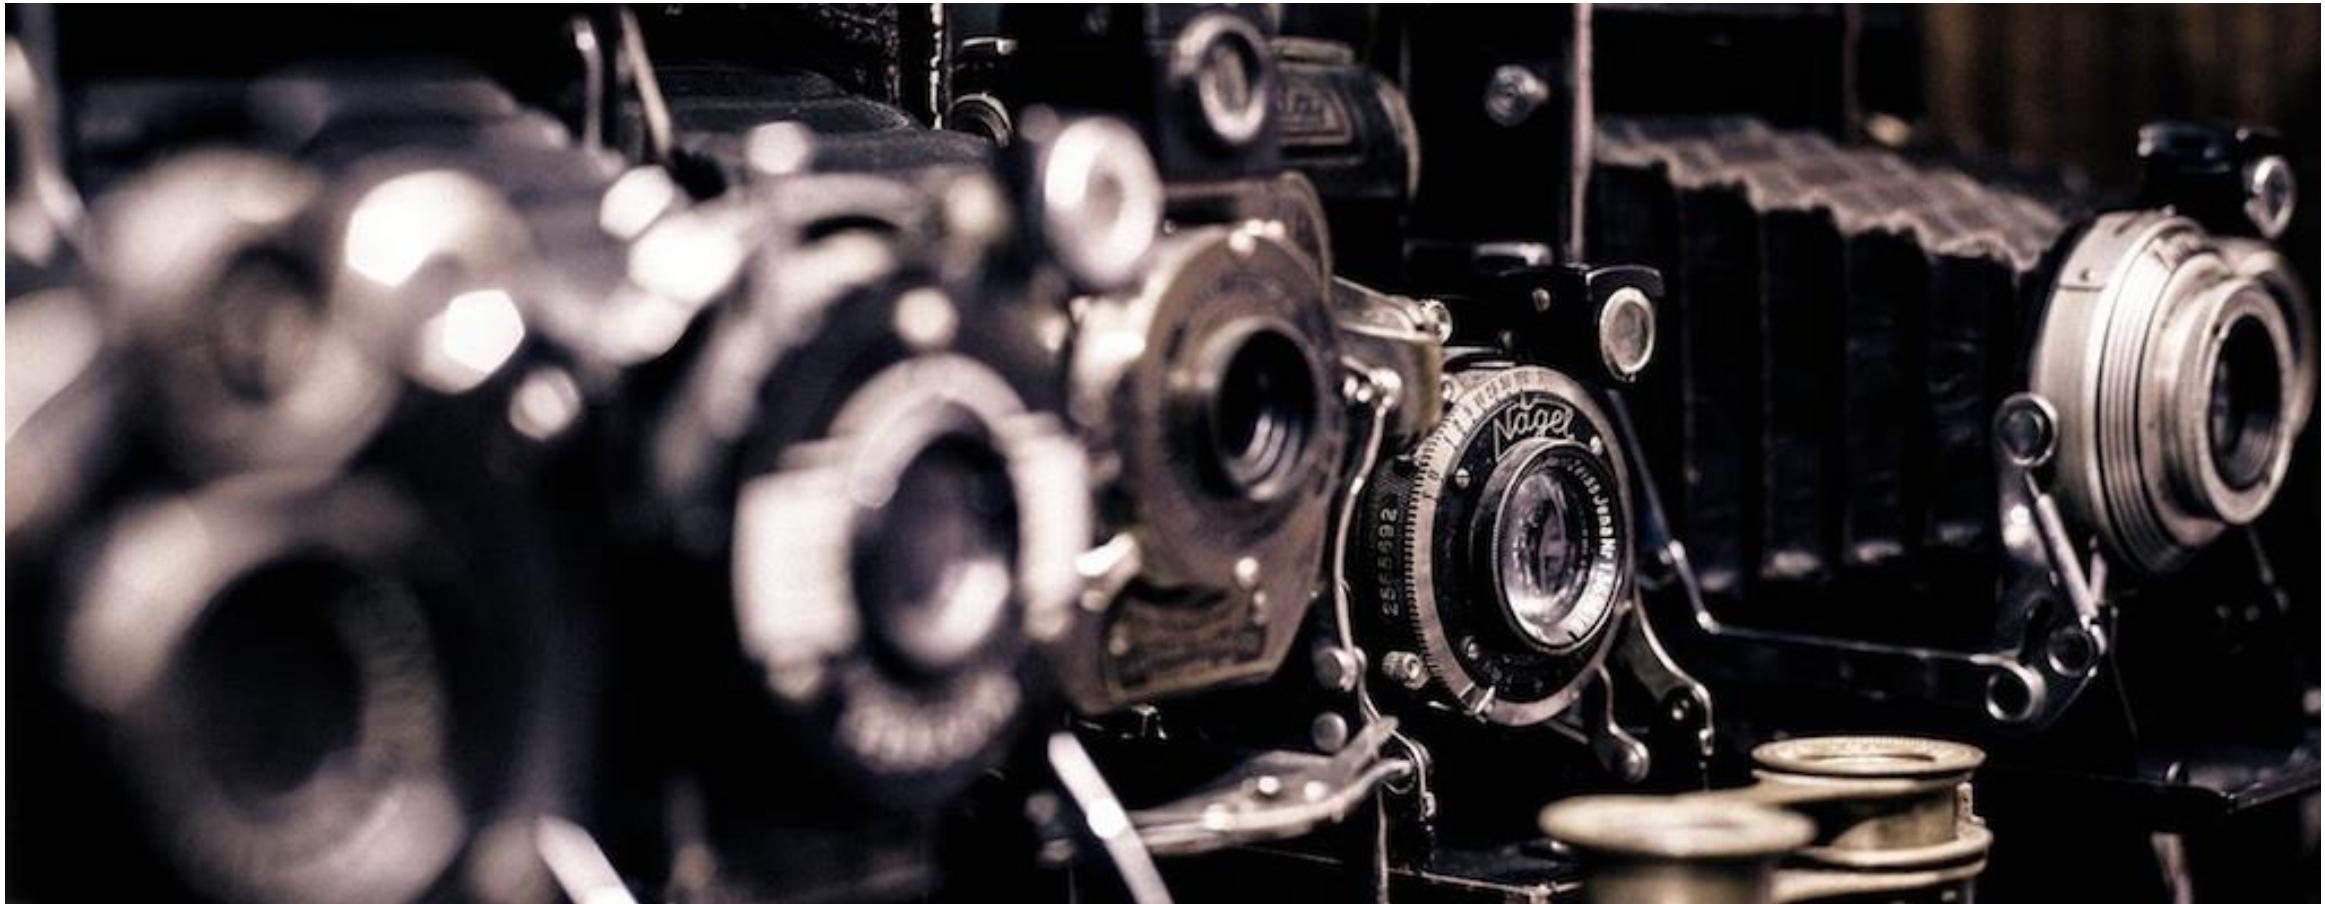


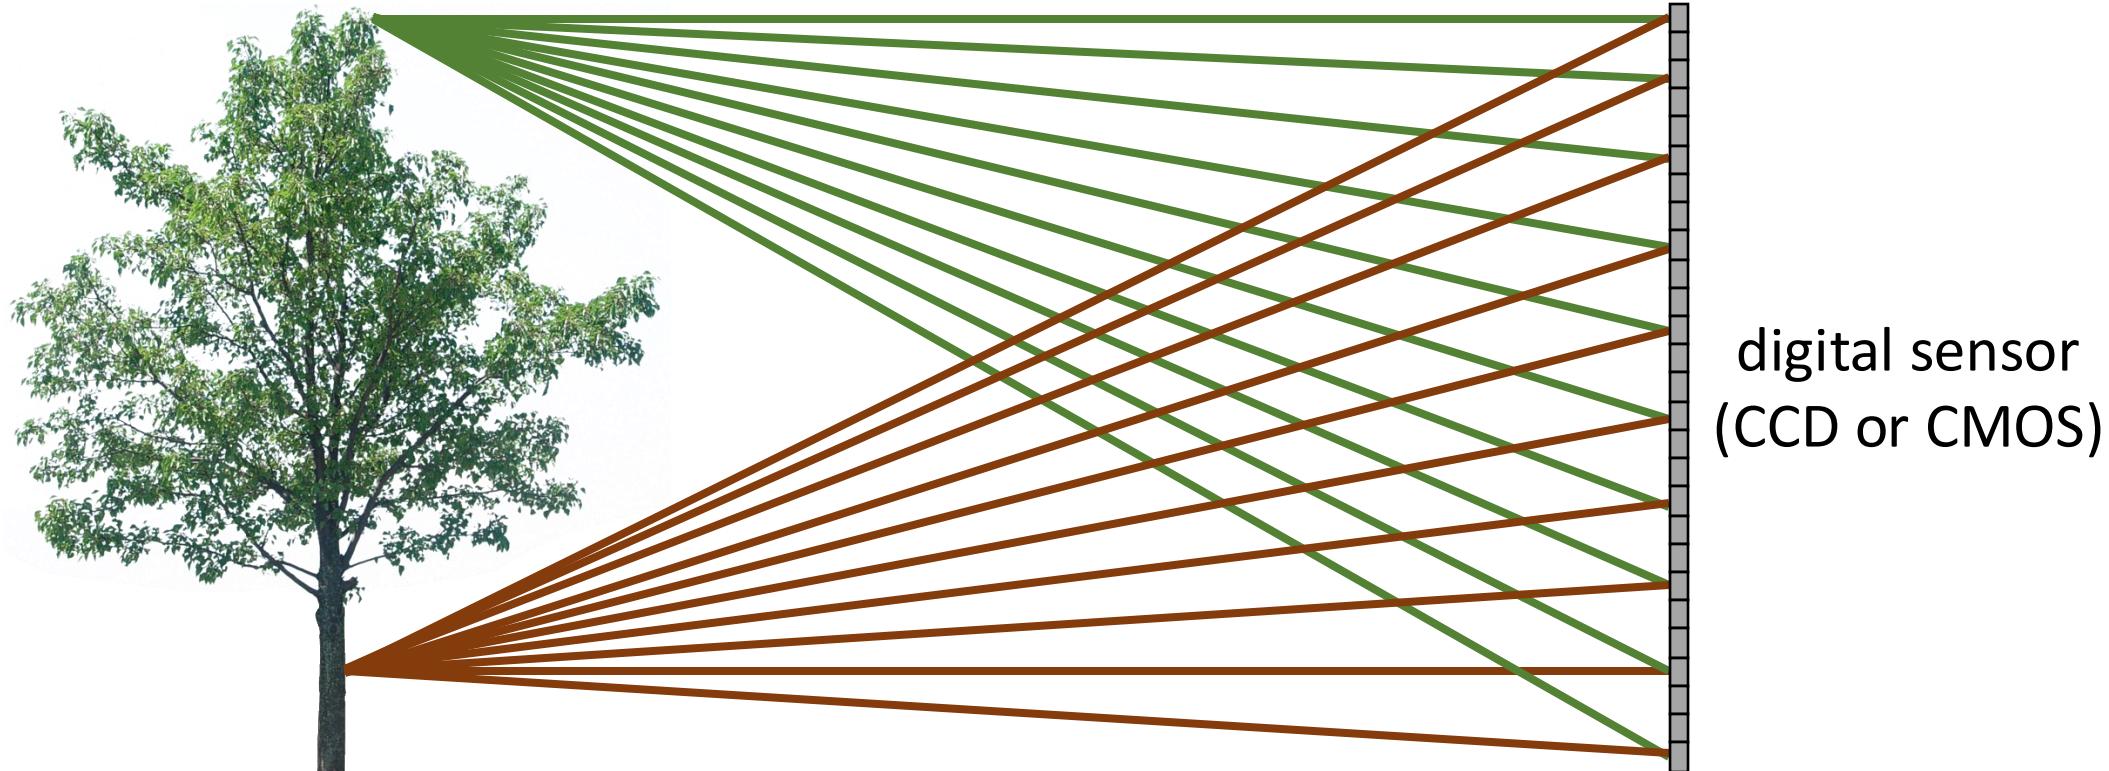
Geometric camera models



ICS 483 Computer Vision
Fall 2025, Lecture 15

Recap: Bare-sensor imaging

real-world
object



All scene points contribute to all sensor pixels

What does the
image on the
sensor look like?

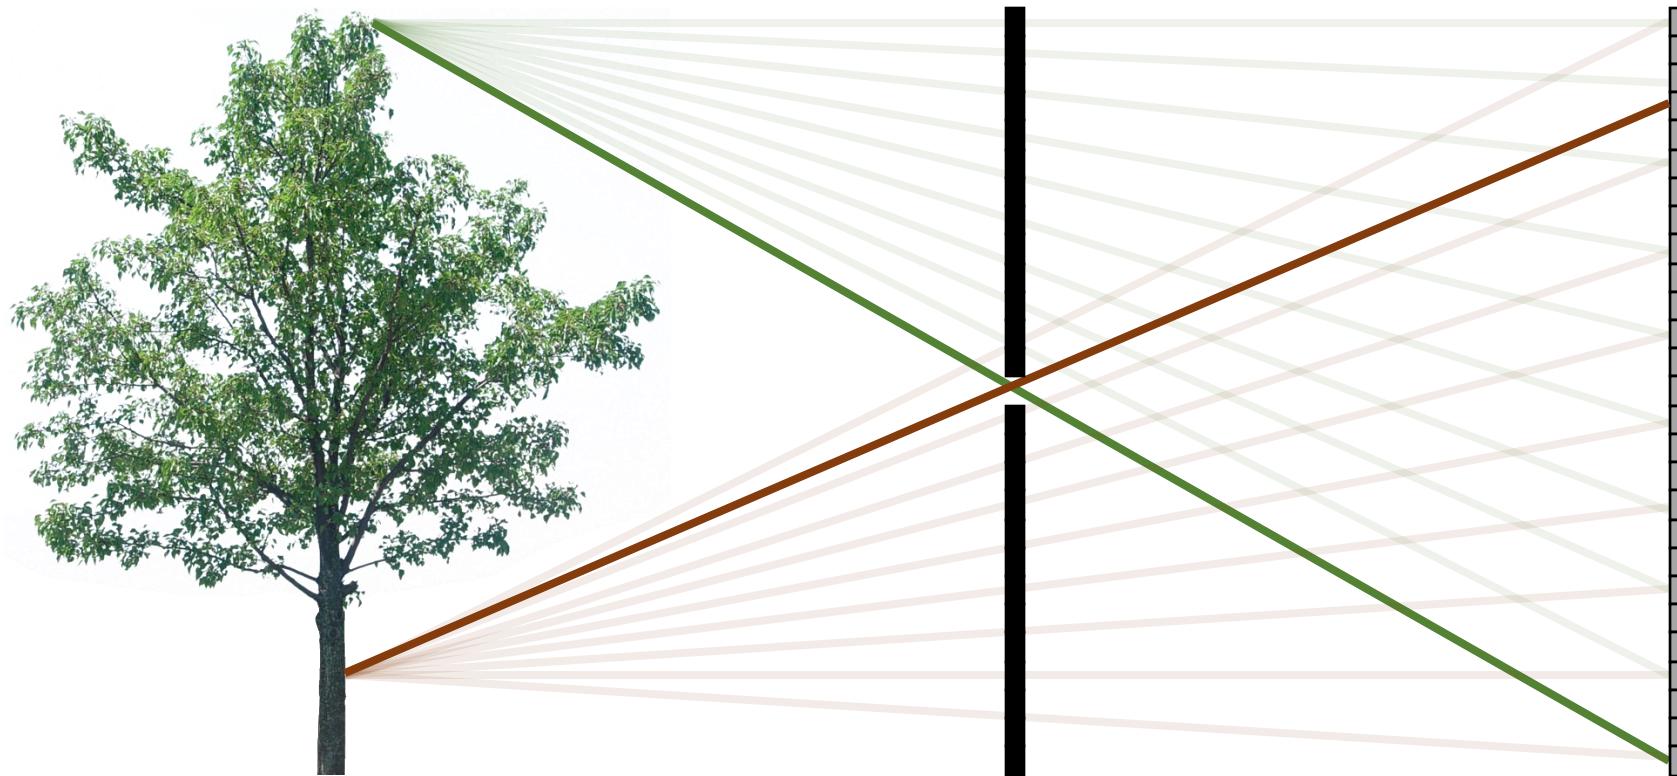
Recap: Bare-sensor imaging



All scene points contribute to all sensor pixels

Recap: Pinhole imaging

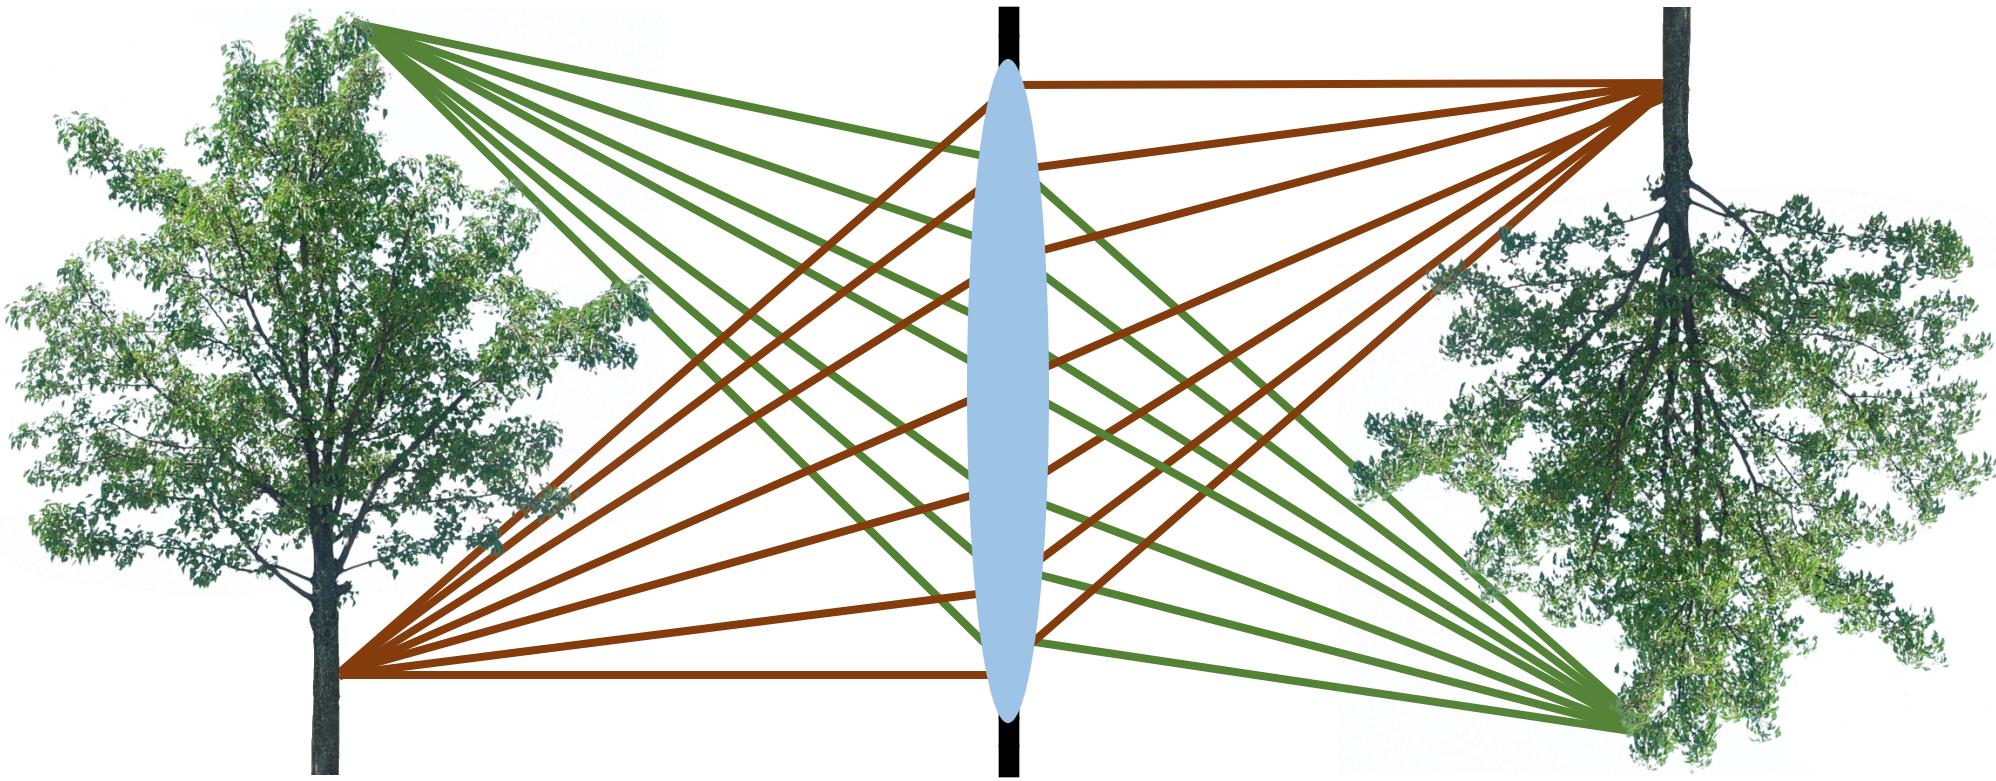
real-world
object



Each scene point contributes to only one sensor pixel

What does the
image on the
sensor look like?

Recap: The lens camera

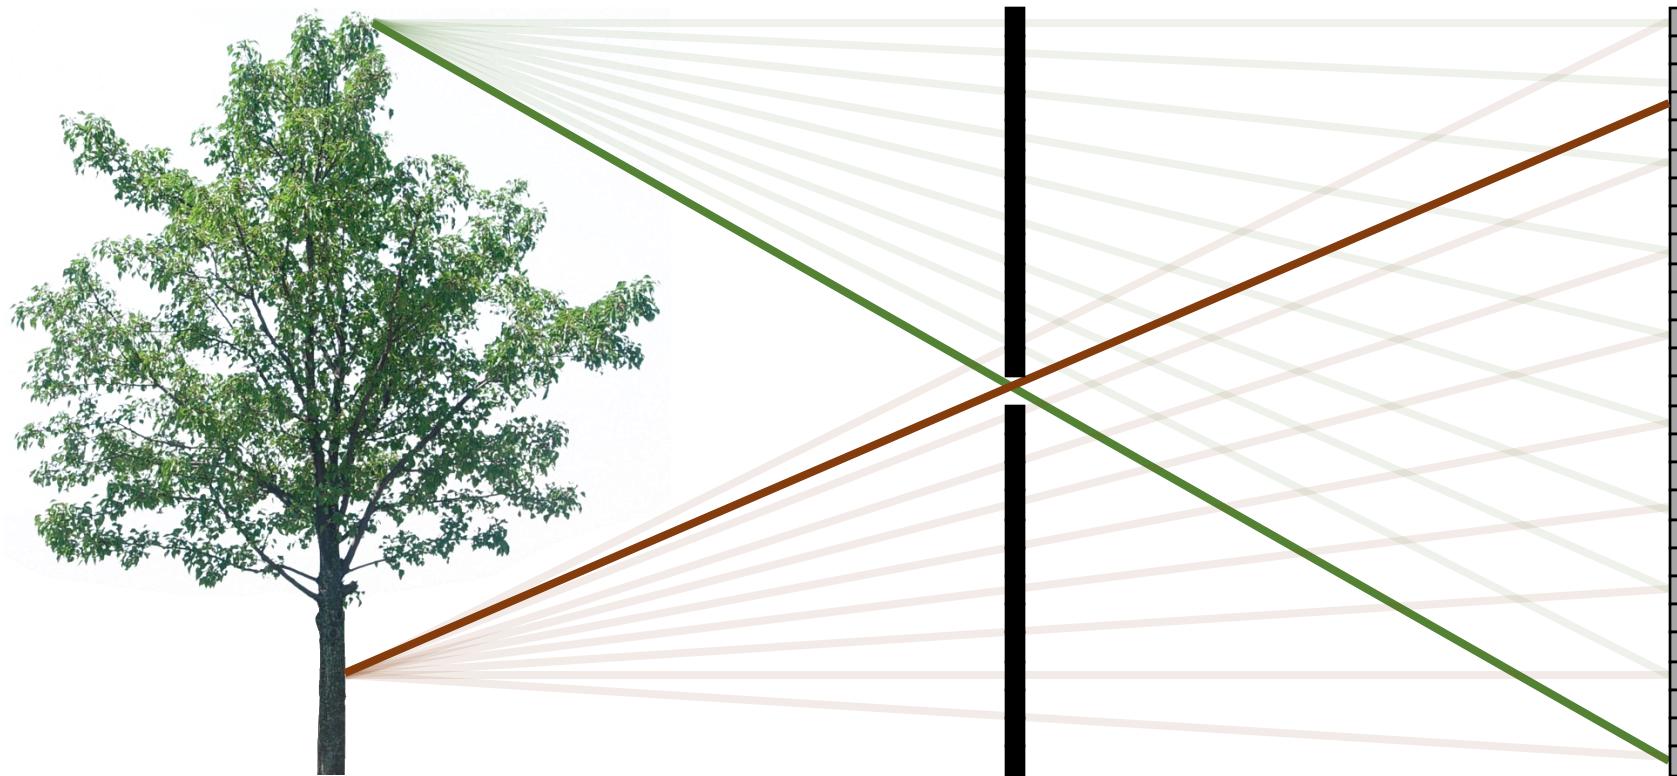


Lenses map “bundles” of rays from points on the scene to the sensor.

How does this mapping work exactly?

Recap: Pinhole imaging

real-world
object



Each scene point contributes to only one sensor pixel

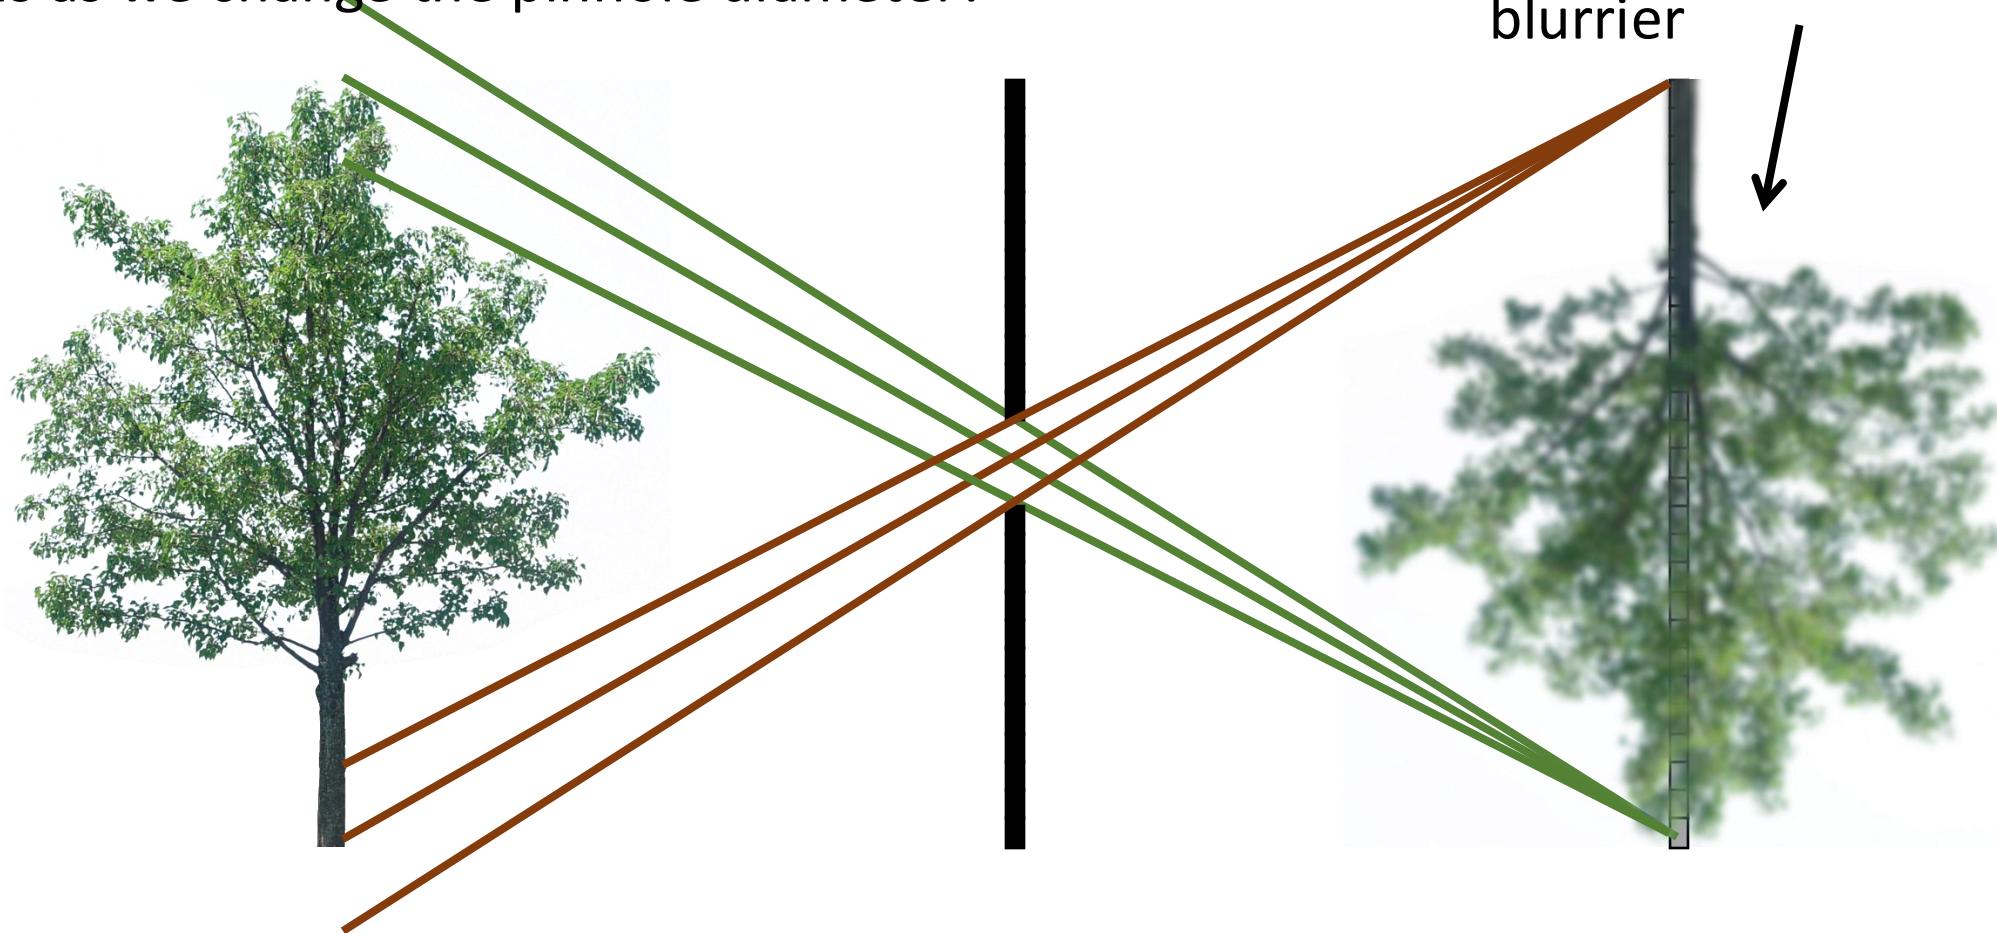
What does the
image on the
sensor look like?

Pinhole size

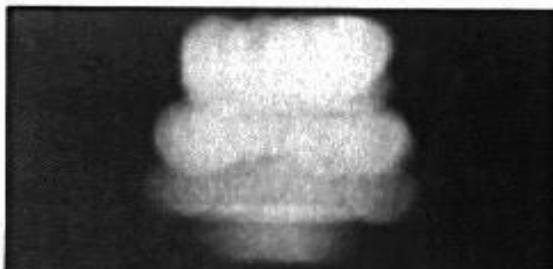
What happens as we change the pinhole diameter?

real-world
object

object projection becomes
blurrier



Shrinking the aperture



2 mm



1 mm



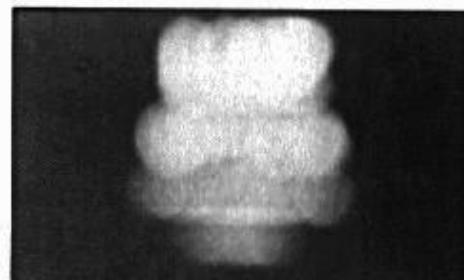
0.6mm



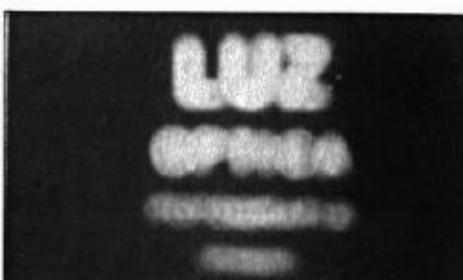
0.35 mm

- Why not make the aperture as small as possible?
 - Less light gets through
 - *Diffraction* effects...

Shrinking the aperture



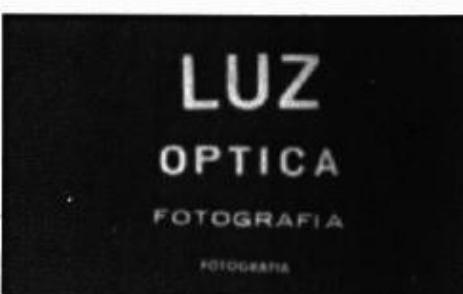
2 mm



1 mm



0.6 mm



0.35 mm



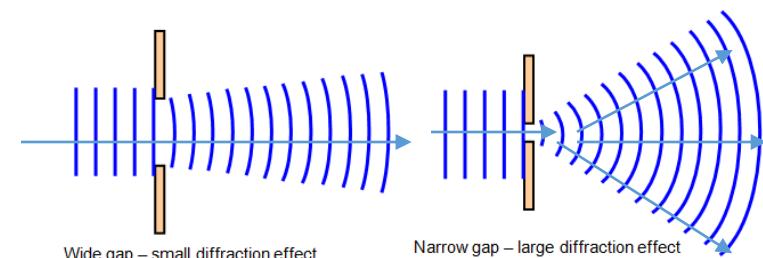
0.15 mm



0.07 mm



<https://physics.highpoint.edu/~jregester/potl/Waves/DiffractionInterference/Diffraction.html>



Wide gap – small diffraction effect

Narrow gap – large diffraction effect

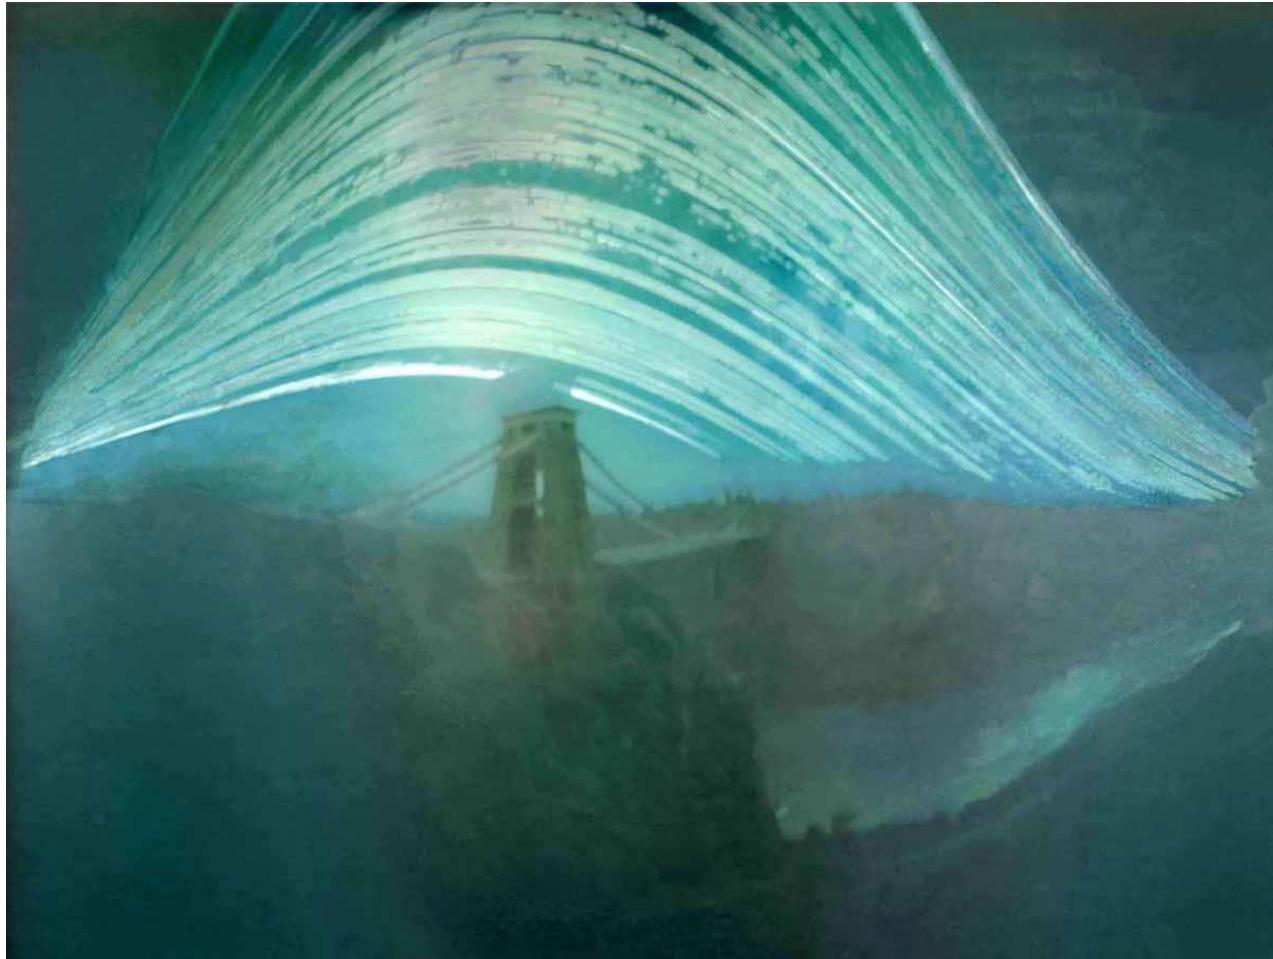
Let's try the DIY pinhole camera!



Home-made pinhole camera



Pinhole photography for solargraphy



Justin Quinnell, The Clifton Suspension Bridge. December 17th 2007 - June 21st 2008
6-month exposure

Moving subjects

Shutter speed choice becomes more important when you photograph moving objects. The quicker the subject is moving, the faster the shutter speed you need to freeze the subject. Go for a slower speed and the moving elements will appear blurred – but get the right degree of blur and your shot can look great.

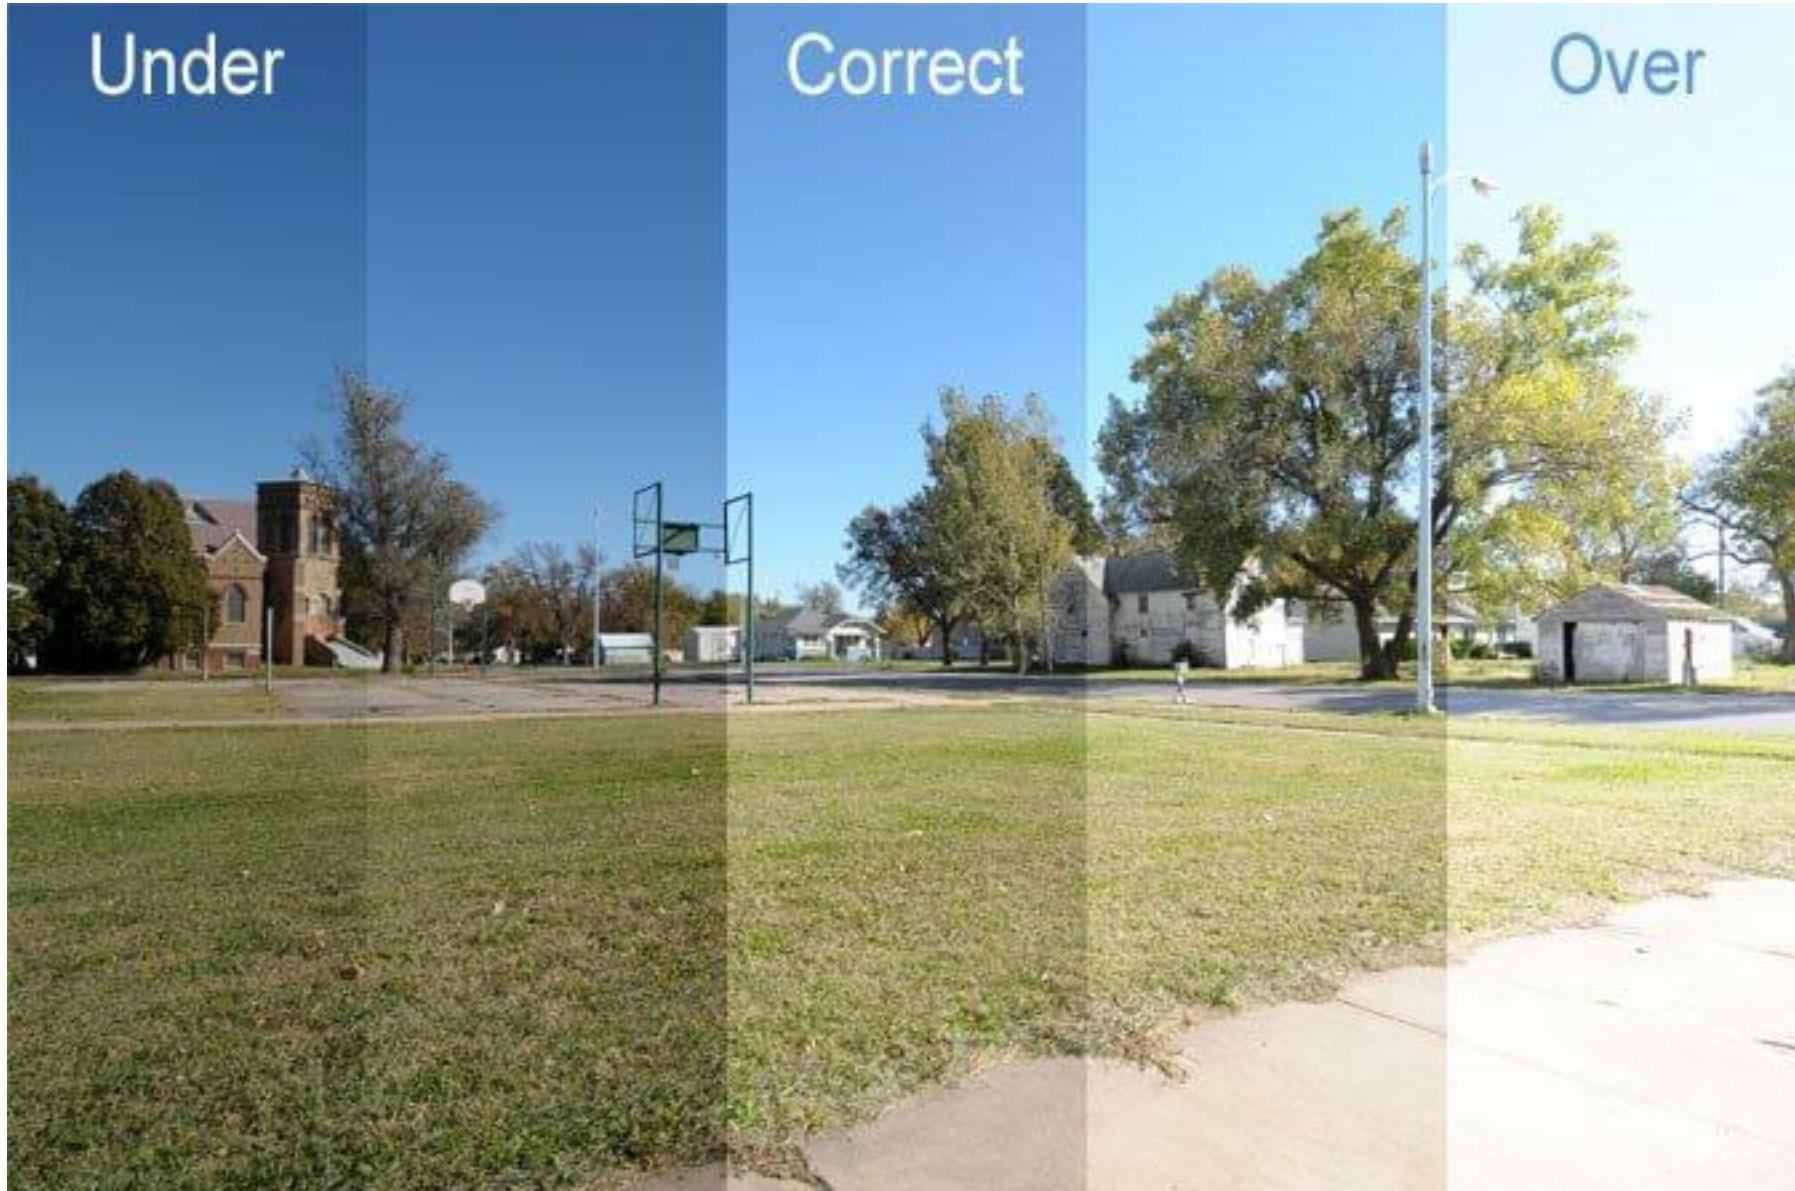


Motion blur



Long exposure time





<https://www.studiobinder.com/blog/what-is-overexposure-in-photography/>



<https://petapixel.com/2019/07/17/this-is-the-worlds-first-solargraphy-timelapse/>

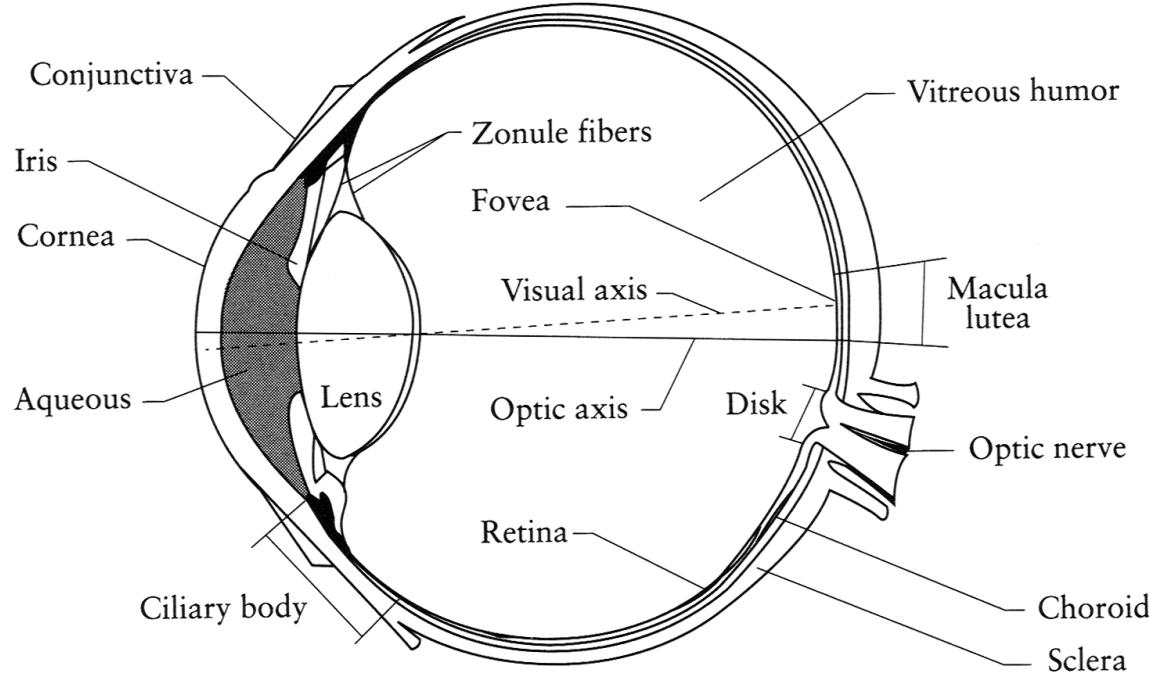




How's time-lapse video made?

https://www.huffpost.com/entry/time-lapse-gifs_n_4421810

The eye

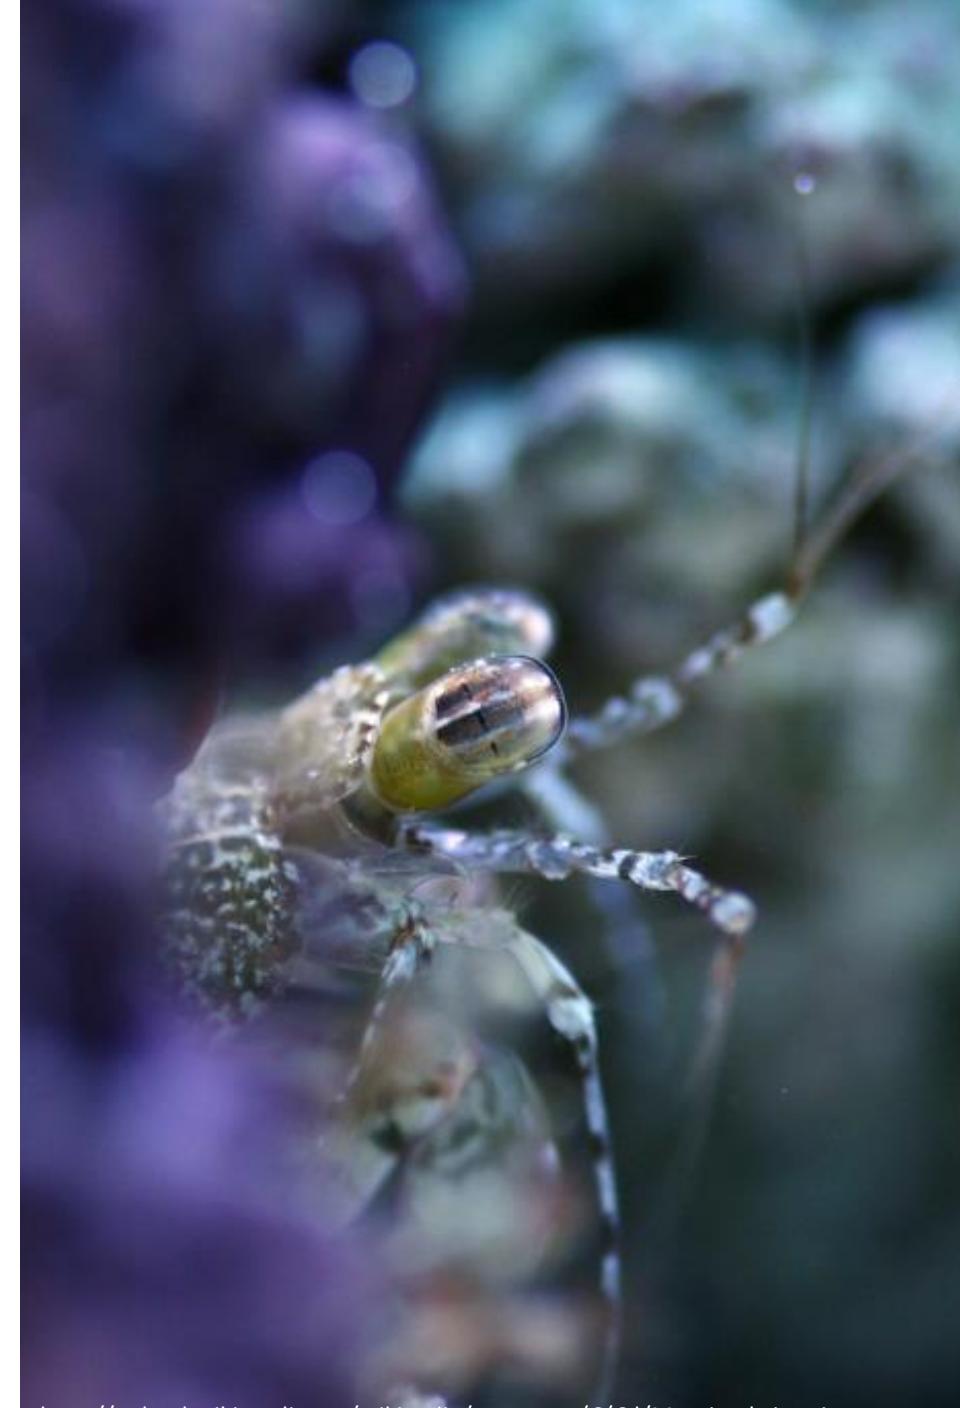
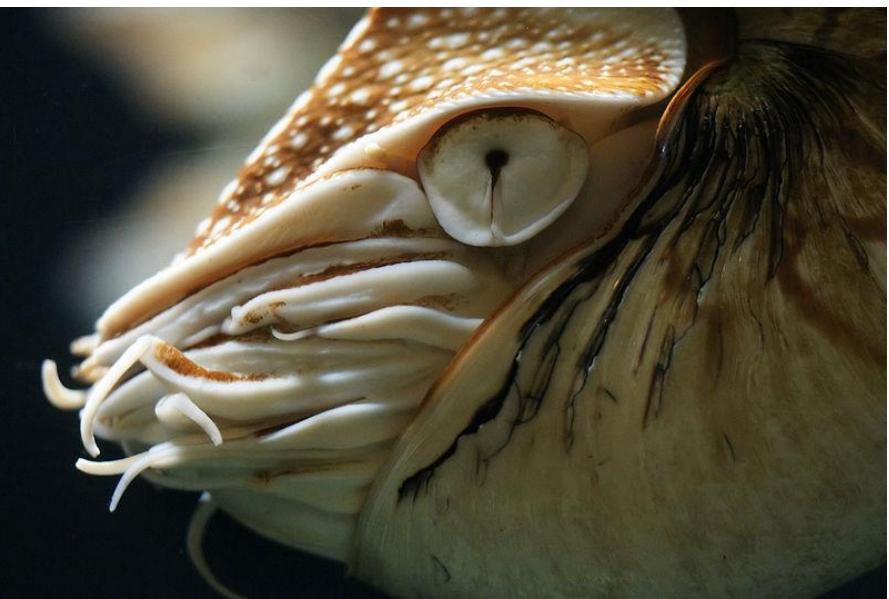


The human eye is a camera

- **Iris** - colored annulus with radial muscles
- **Pupil** - the hole (aperture) whose size is controlled by the iris
- What's the “film”?
 - photoreceptor cells (rods and cones) in the **retina**



Eyes in nature: eyespots to pinhole camera and compound eyes

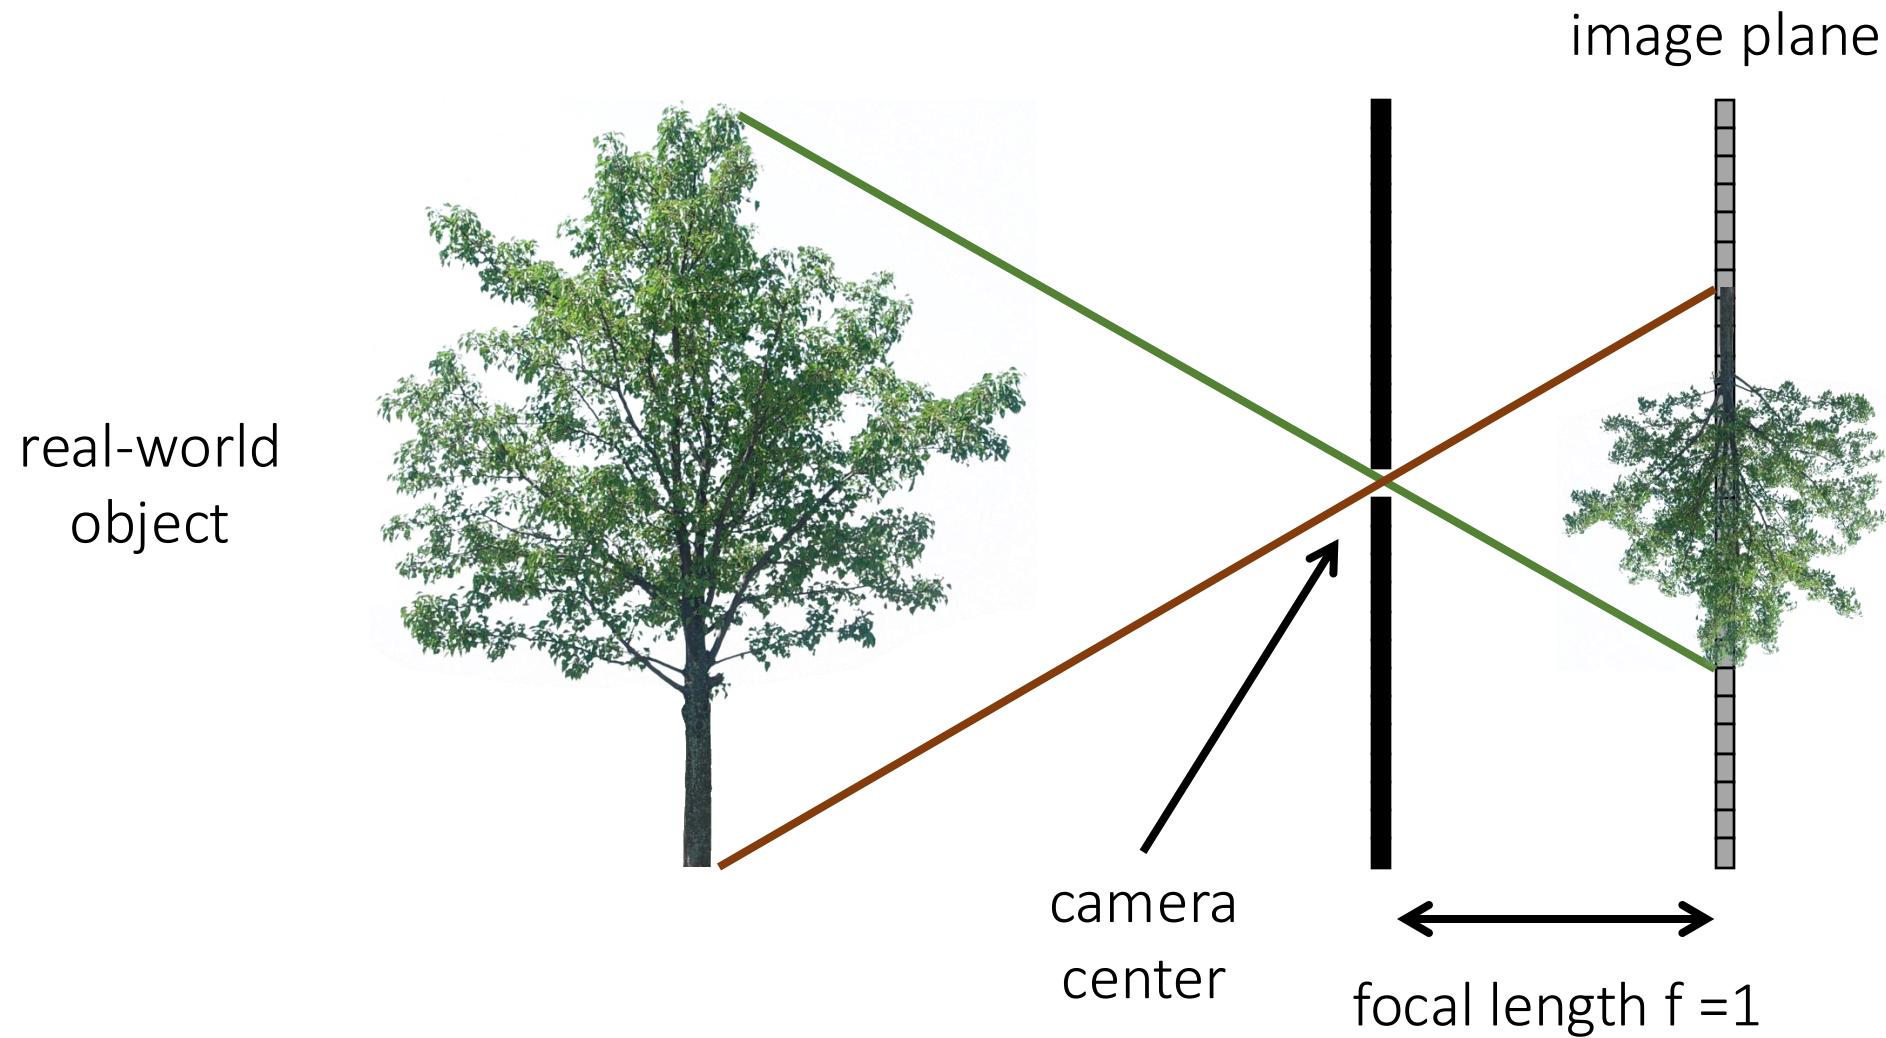


Great Camera Demo

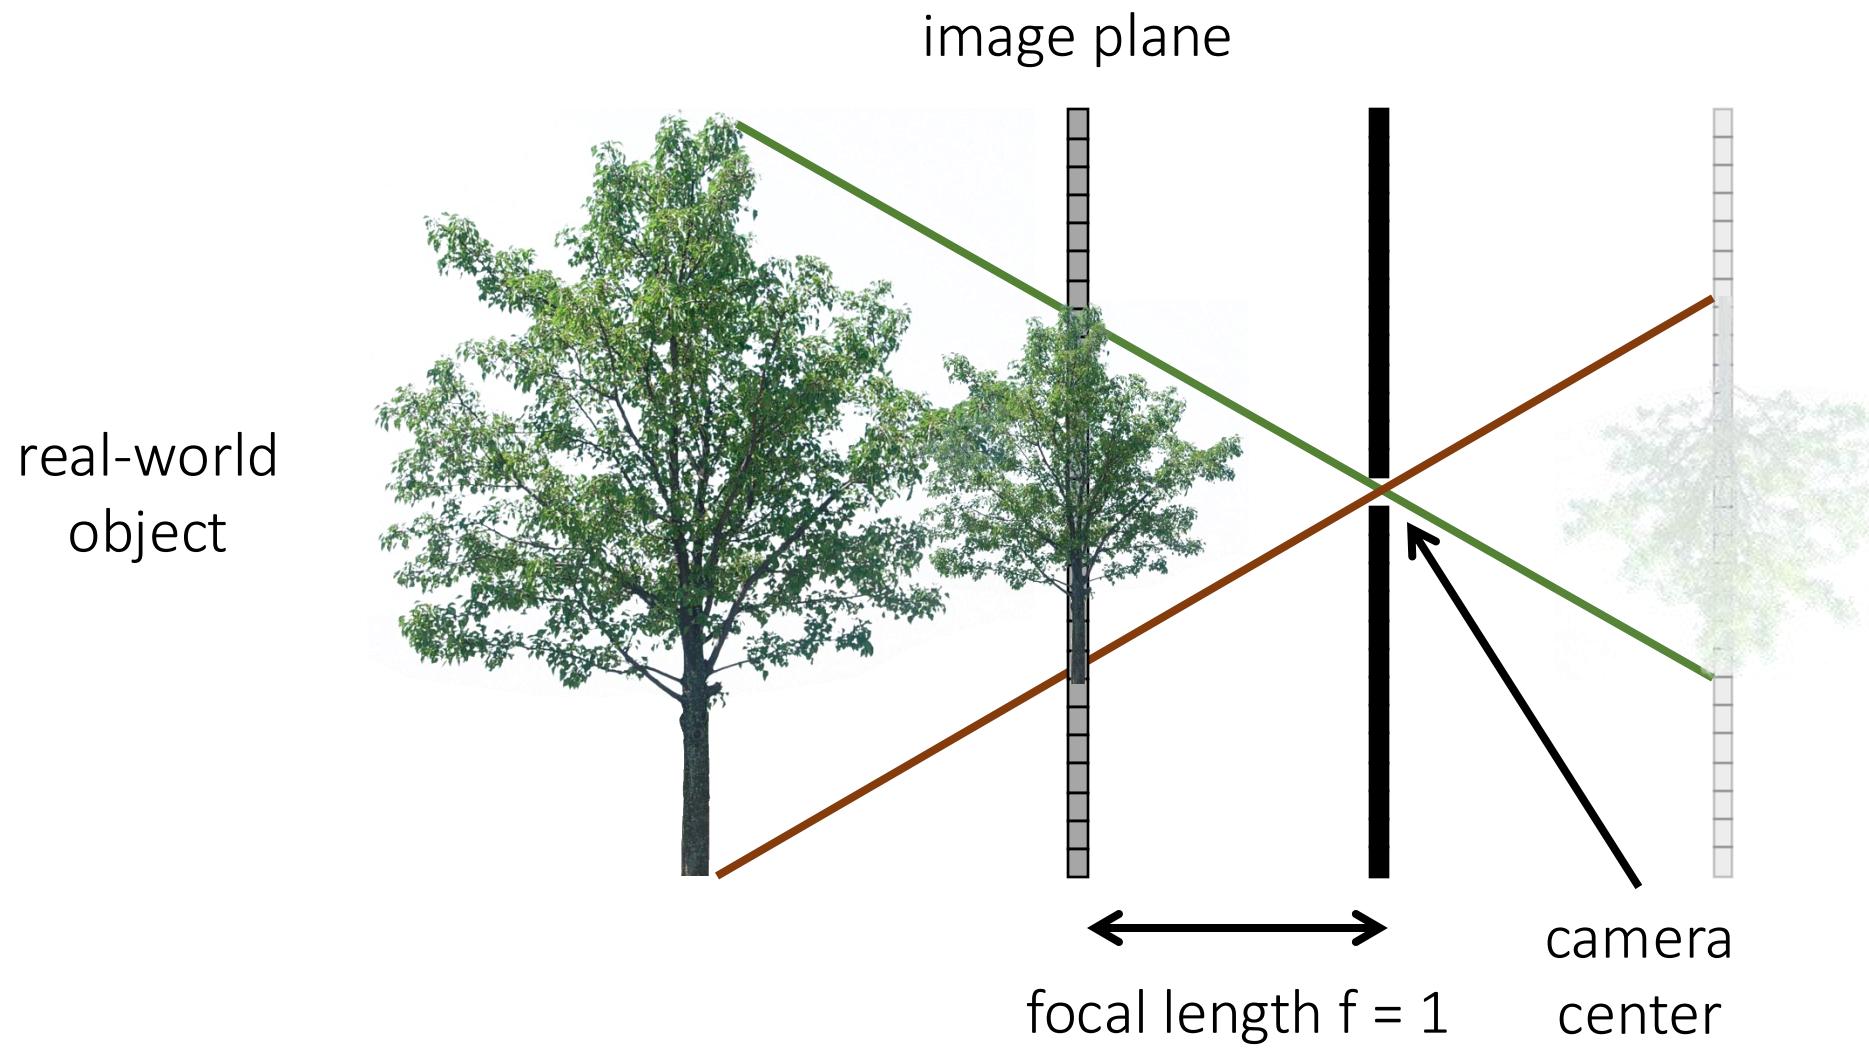
<https://ciechanow.ski/cameras-and-lenses/>

Make sure you check it out!

The pinhole camera



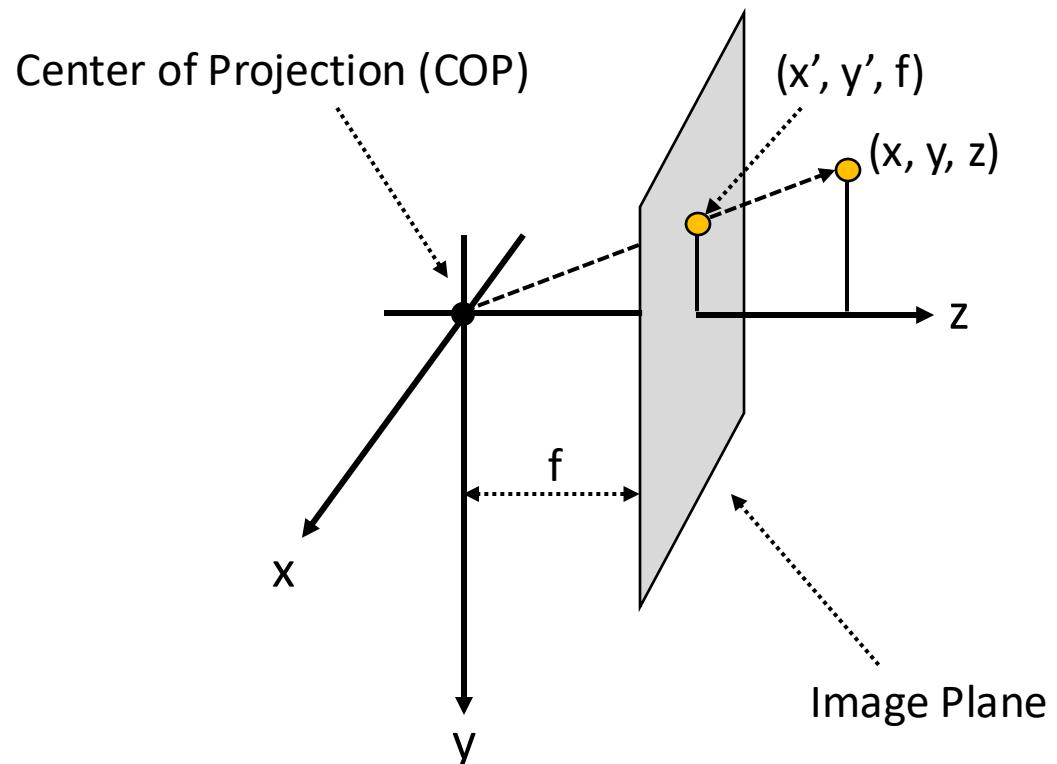
The (rearranged) pinhole camera



Modeling projection

- Projection equations
 - Compute intersection with PP of ray from (x, y, z) to COP
 - Derived using similar triangles
- We get the projection by throwing out the last coordinate:

$$(x, y, z) \rightarrow \left(f \frac{x}{z}, f \frac{y}{z}, f \right)$$



Modeling projection

- Is this a linear transformation?
 - no—division by z is nonlinear

Homogeneous coordinates to the rescue—again!

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image
coordinates

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

homogeneous scene
coordinates

Converting from homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

$$\begin{bmatrix} x \\ y \\ z \\ w \end{bmatrix} \Rightarrow (x/w, y/w, z/w)$$

Perspective Projection

Projection is a matrix multiply using homogeneous coordinates:

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z/f \end{bmatrix} \Rightarrow \left(f \frac{x}{z}, f \frac{y}{z} \right)$$

divide by third coordinate

This is known as **perspective projection**

- The matrix is the **projection matrix**
- (Can also represent as a 4x4 matrix – OpenGL does something like this)

Perspective Projection

How does scaling the projection matrix change the transformation?

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z/f \\ 1 \end{bmatrix} \Rightarrow \left(f\frac{x}{z}, f\frac{y}{z} \right)$$

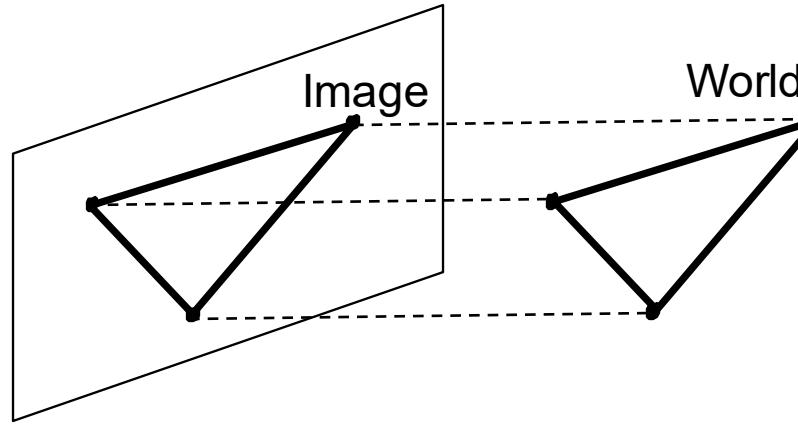
Scale by f :

$$\begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} fx \\ fy \\ z \\ 1 \end{bmatrix} \Rightarrow \left(f\frac{x}{z}, f\frac{y}{z} \right)$$

Scaling a projection matrix produces an equivalent projection matrix!

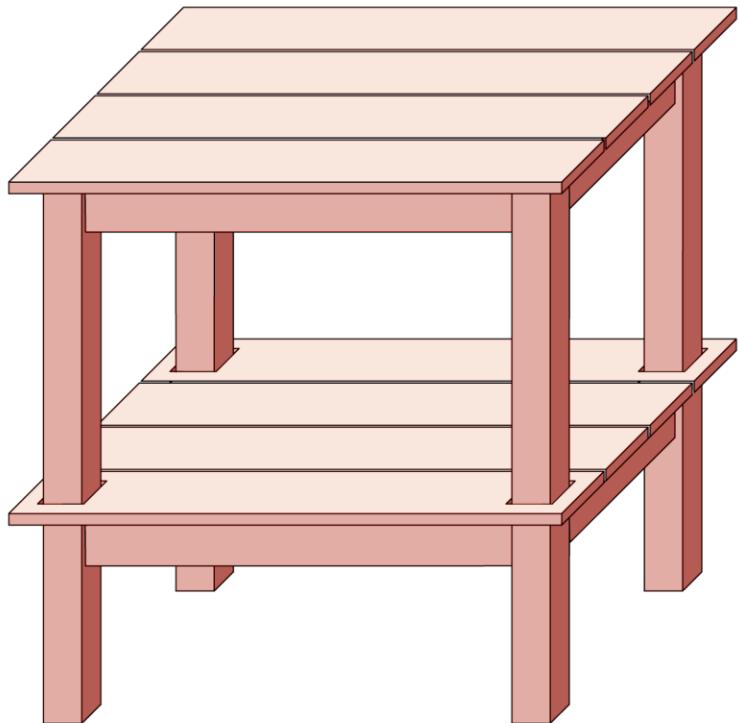
Orthographic projection

- Special case of perspective projection
 - Distance from the COP to the PP is infinite

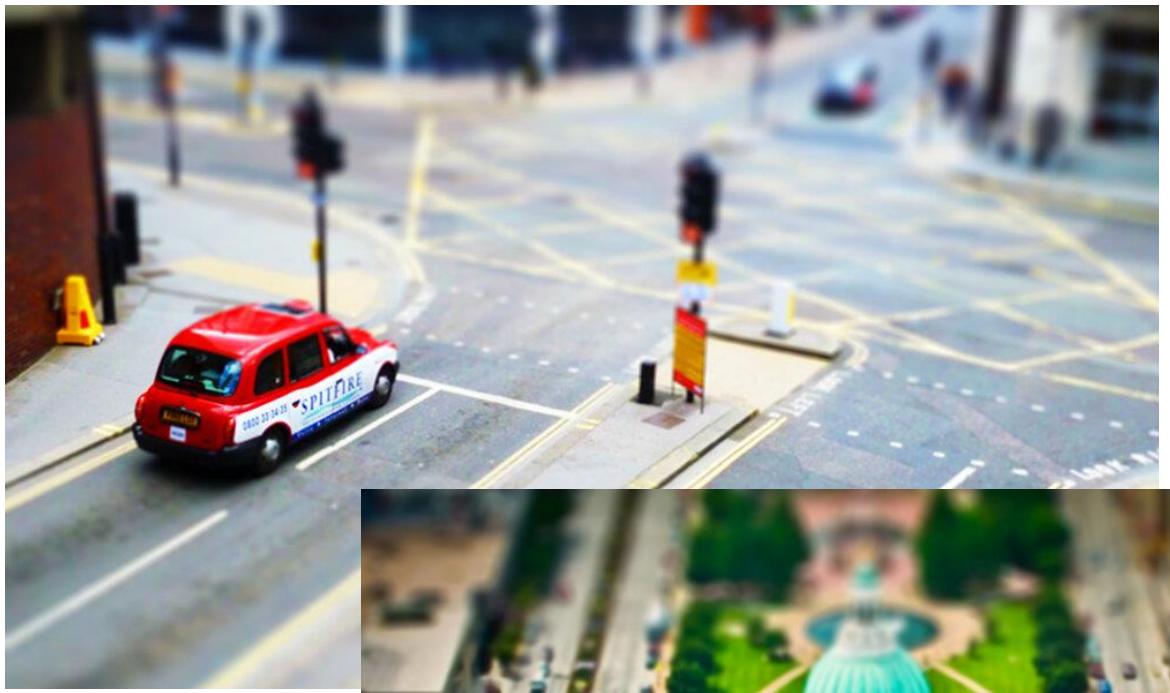


$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \Rightarrow (x, y)$$

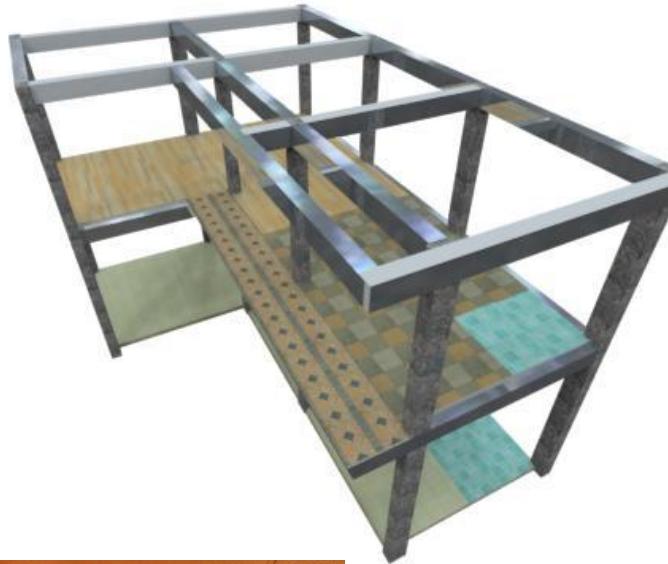
Orthographic projection



How's this related to
"miniature" effect by
tilt-shift photography?



Perspective projection



Variants of orthographic projection

- Scaled orthographic
 - Also called “weak perspective”

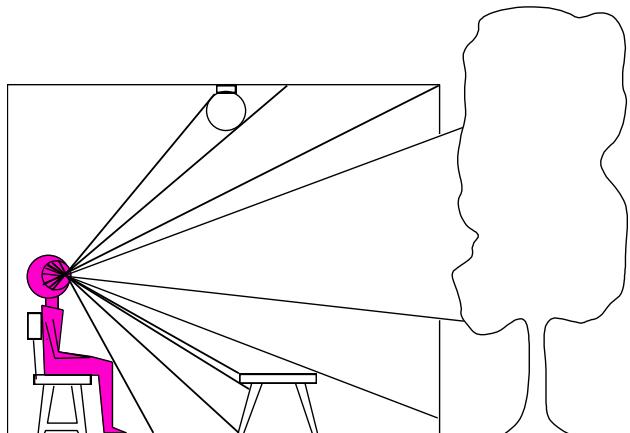
$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1/d \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ 1/d \end{bmatrix} \Rightarrow (dx, dy)$$

- Affine projection
 - Also called “paraperspective”

$$\begin{bmatrix} a & b & c & d \\ e & f & g & h \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

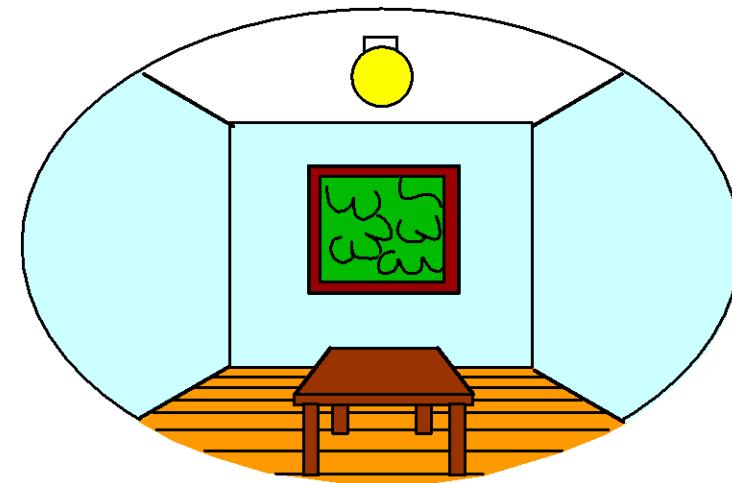
Dimensionality Reduction Machine (3D to 2D)

3D world



Point of observation

2D image



What have we lost?

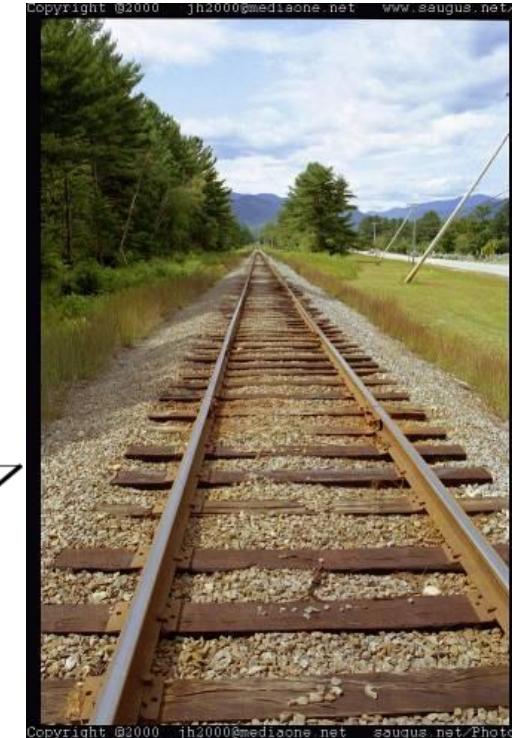
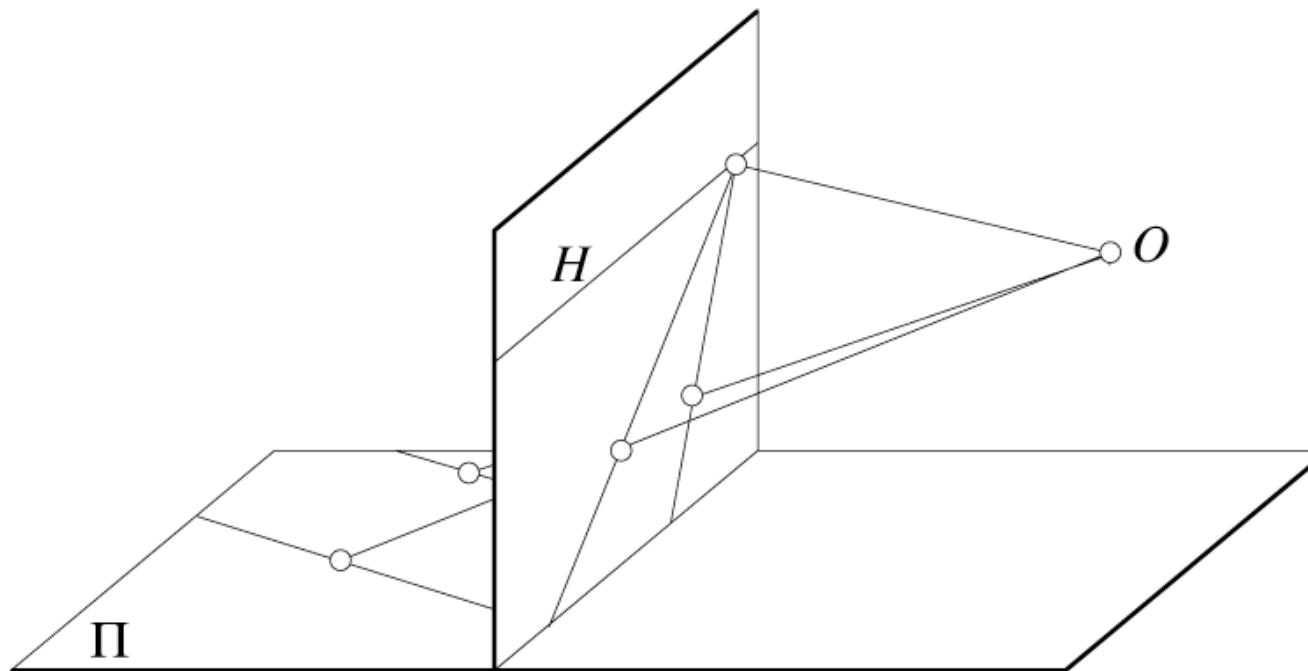
- Angles
- Distances (lengths)

Projection properties

- Many-to-one: any points along same ray map to same point in image
- Points → points
- Lines → lines (collinearity is preserved)
 - But line through focal point projects to a point
- Planes → planes (or half-planes)
 - But plane through focal point projects to line

Projection properties

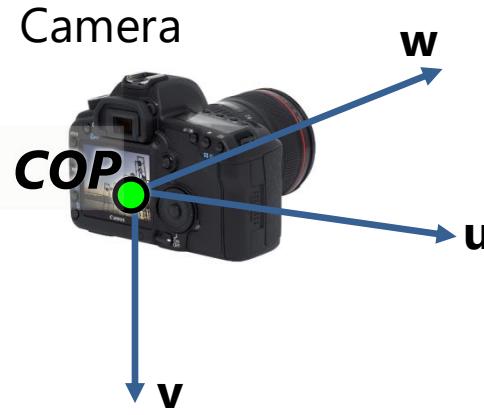
- Parallel lines converge at a vanishing point
 - Each direction in space has its own vanishing point
 - But parallel lines parallel to the image plane remain parallel



Questions?

Camera parameters

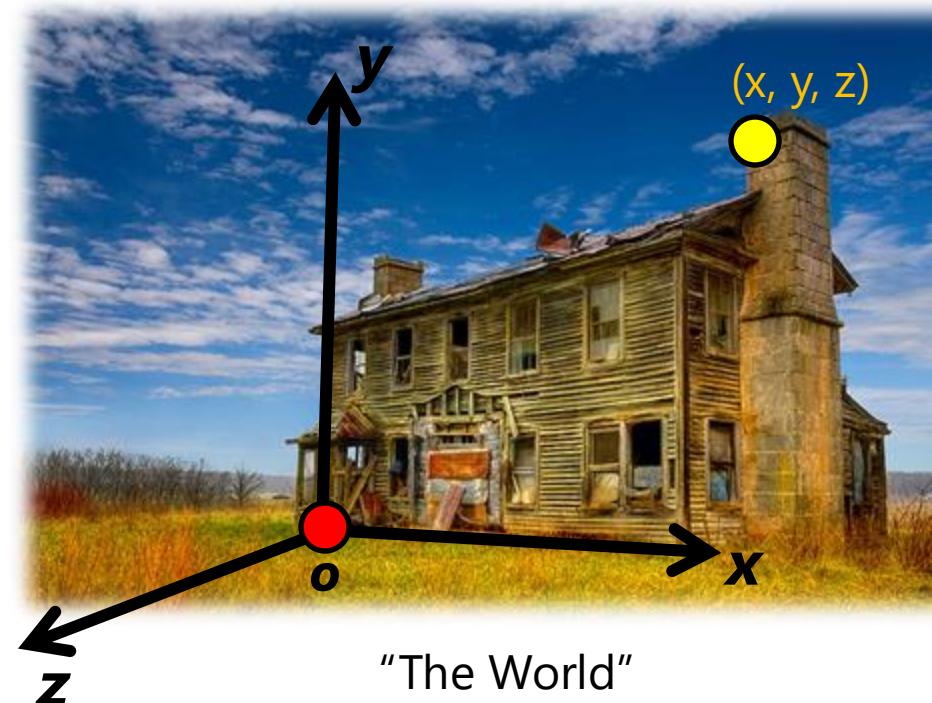
- How can we model the geometry of a camera?



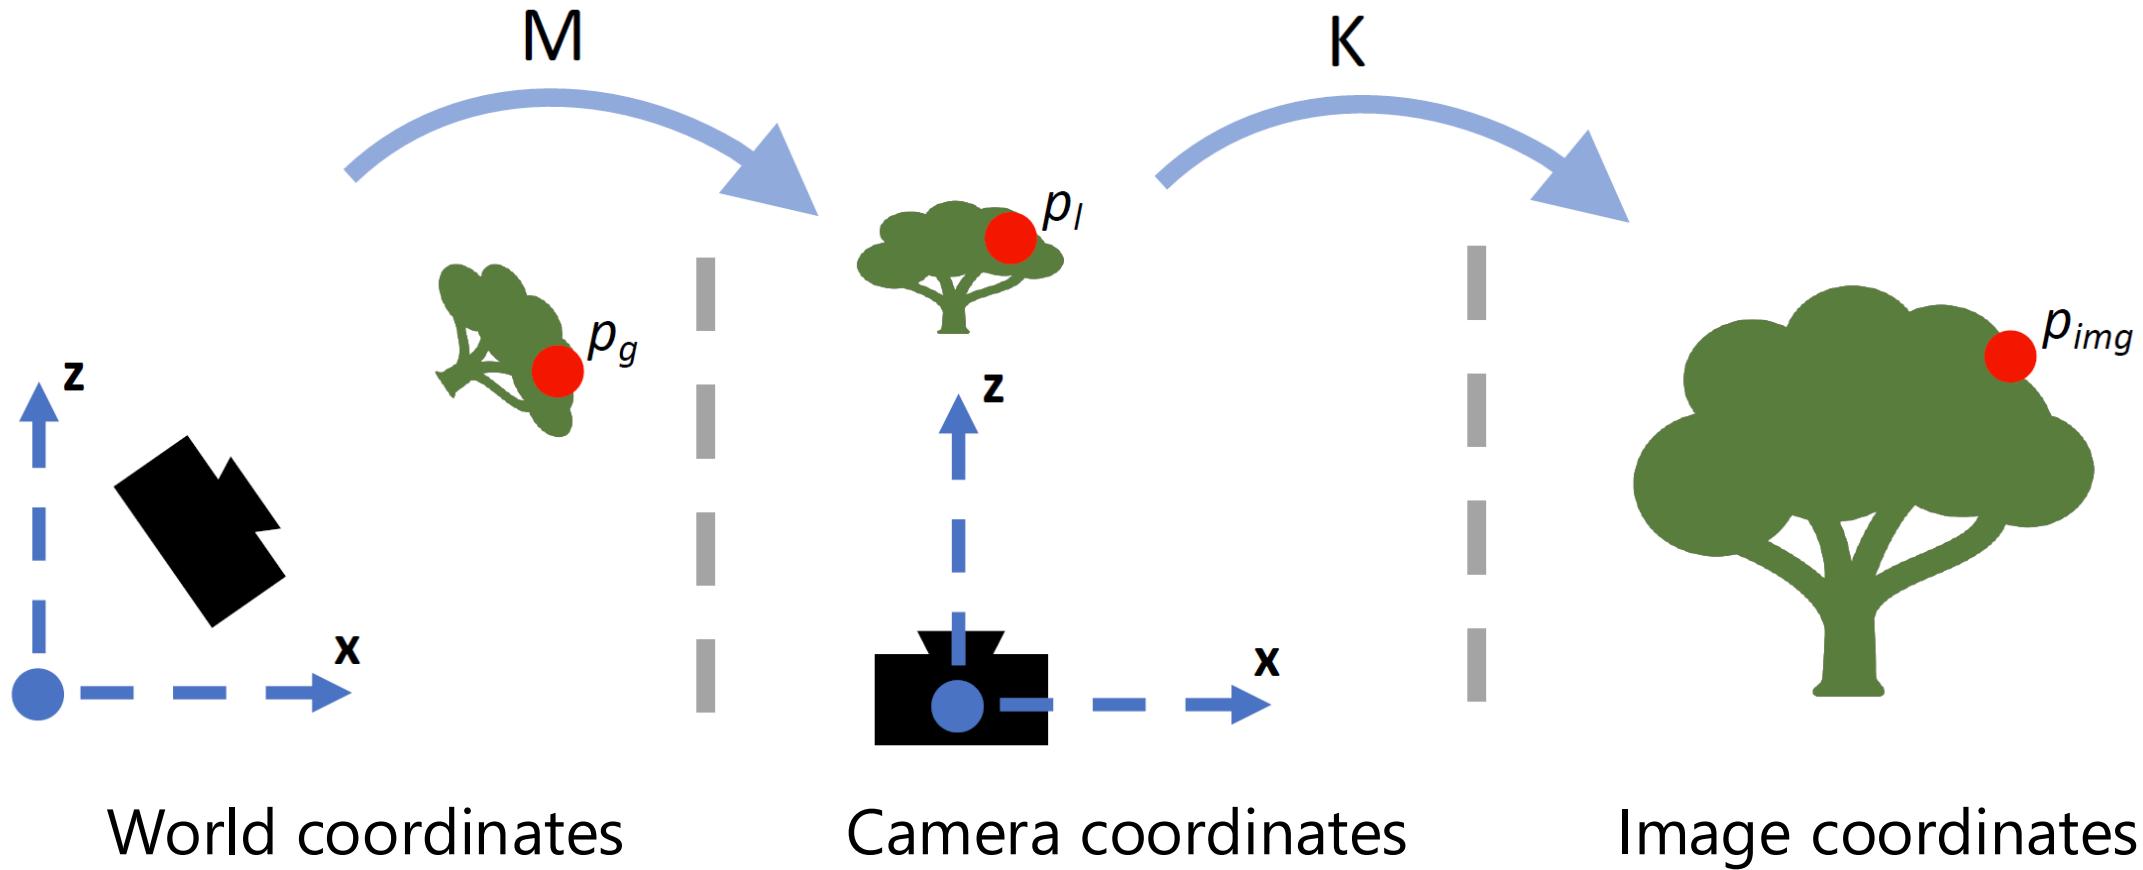
Three important coordinate systems:

1. *World* coordinates
2. *Camera* coordinates
3. *Image* coordinates

How do we project a given world point (x, y, z) to an image point?



Coordinate frames

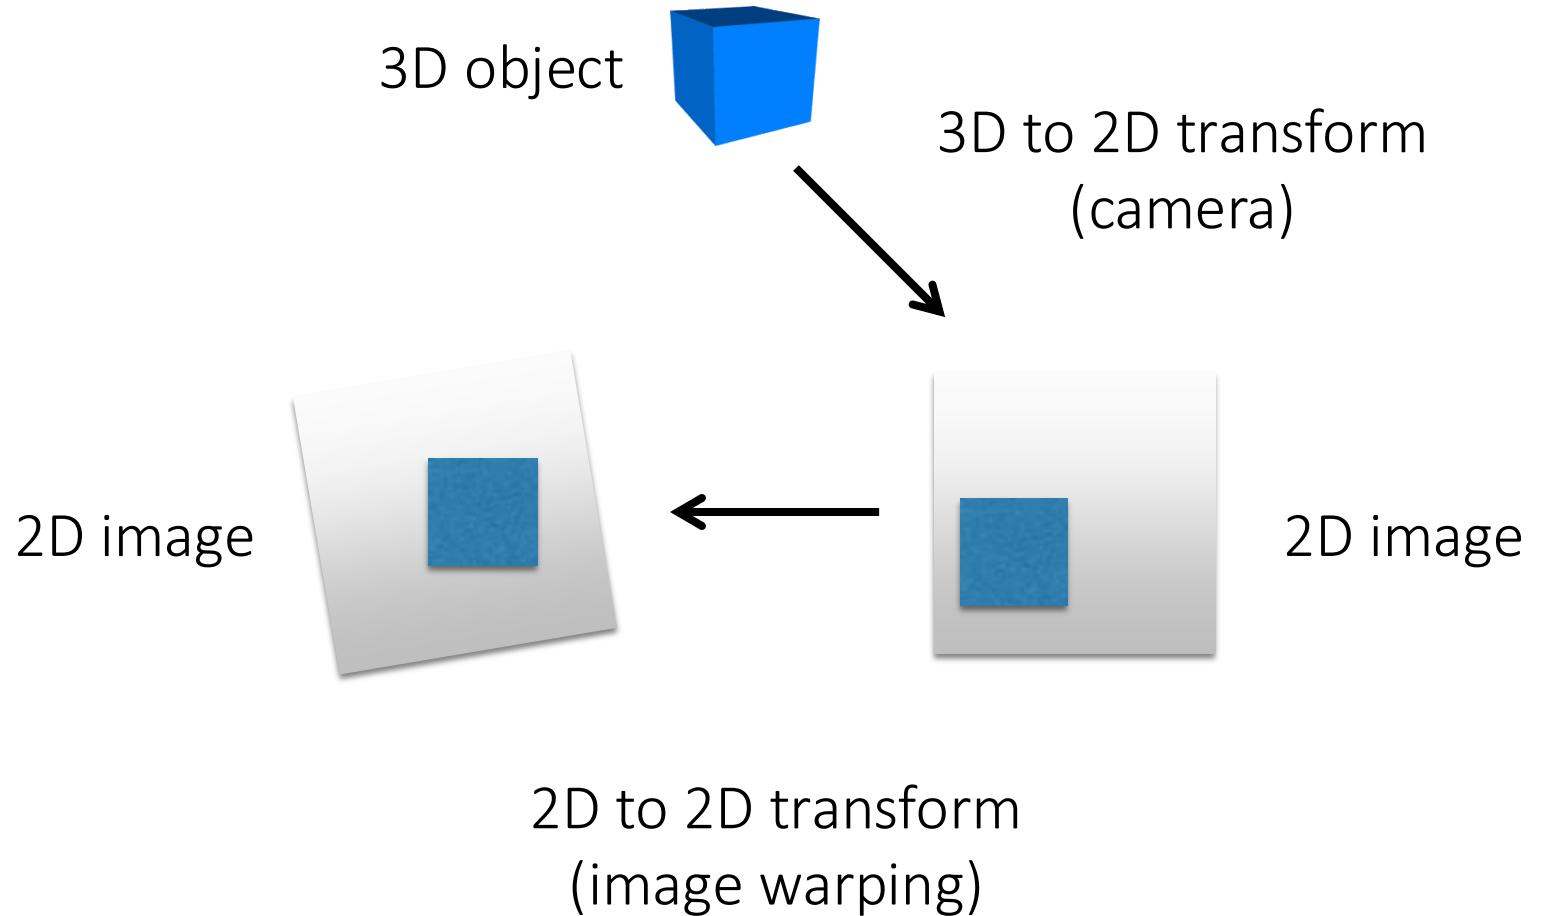


Recap: The camera as a coordinate transformation

A camera is a mapping from:

the 3D world
to:

a 2D image



Recap: The camera as a coordinate transformation

$$\mathbf{x} = \mathbf{P}\mathbf{X}$$

$$\begin{bmatrix} X \\ Y \\ Z \end{bmatrix} = \begin{bmatrix} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

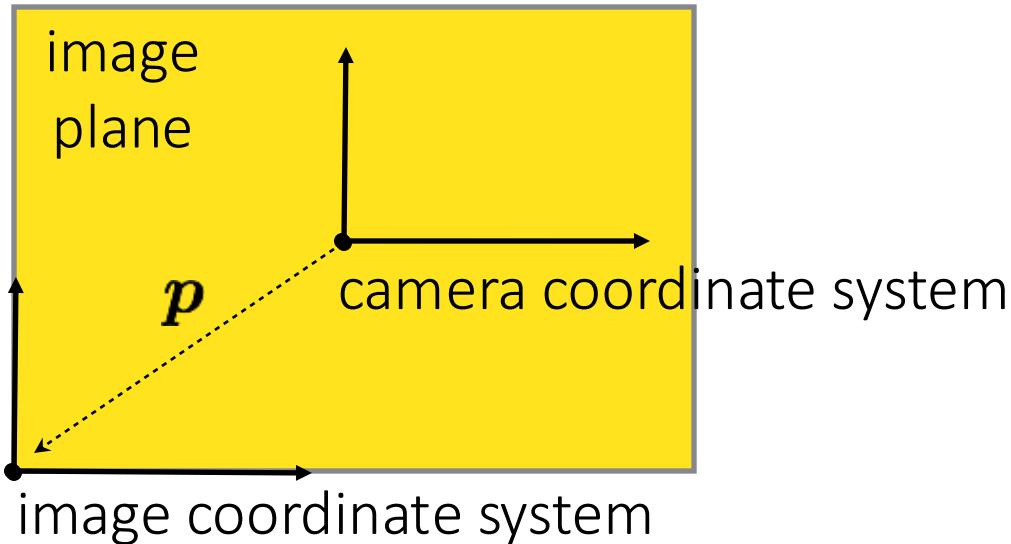
homogeneous
image coordinates
 3×1

camera
matrix
 3×4

homogeneous
world coordinates
 4×1

Recap: Generalizing the camera matrix

In particular, the camera origin and image origin may be different:



How does the camera matrix change?

$$\mathbf{P} = \begin{bmatrix} f & 0 & p_x & 0 \\ 0 & f & p_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

shift vector
transforming
camera origin to
image origin

Recap: Camera matrix decomposition

We can decompose the camera matrix like this:

$$\mathbf{P} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & | & 0 \\ 0 & 1 & 0 & | & 0 \\ 0 & 0 & 1 & | & 0 \end{bmatrix}$$



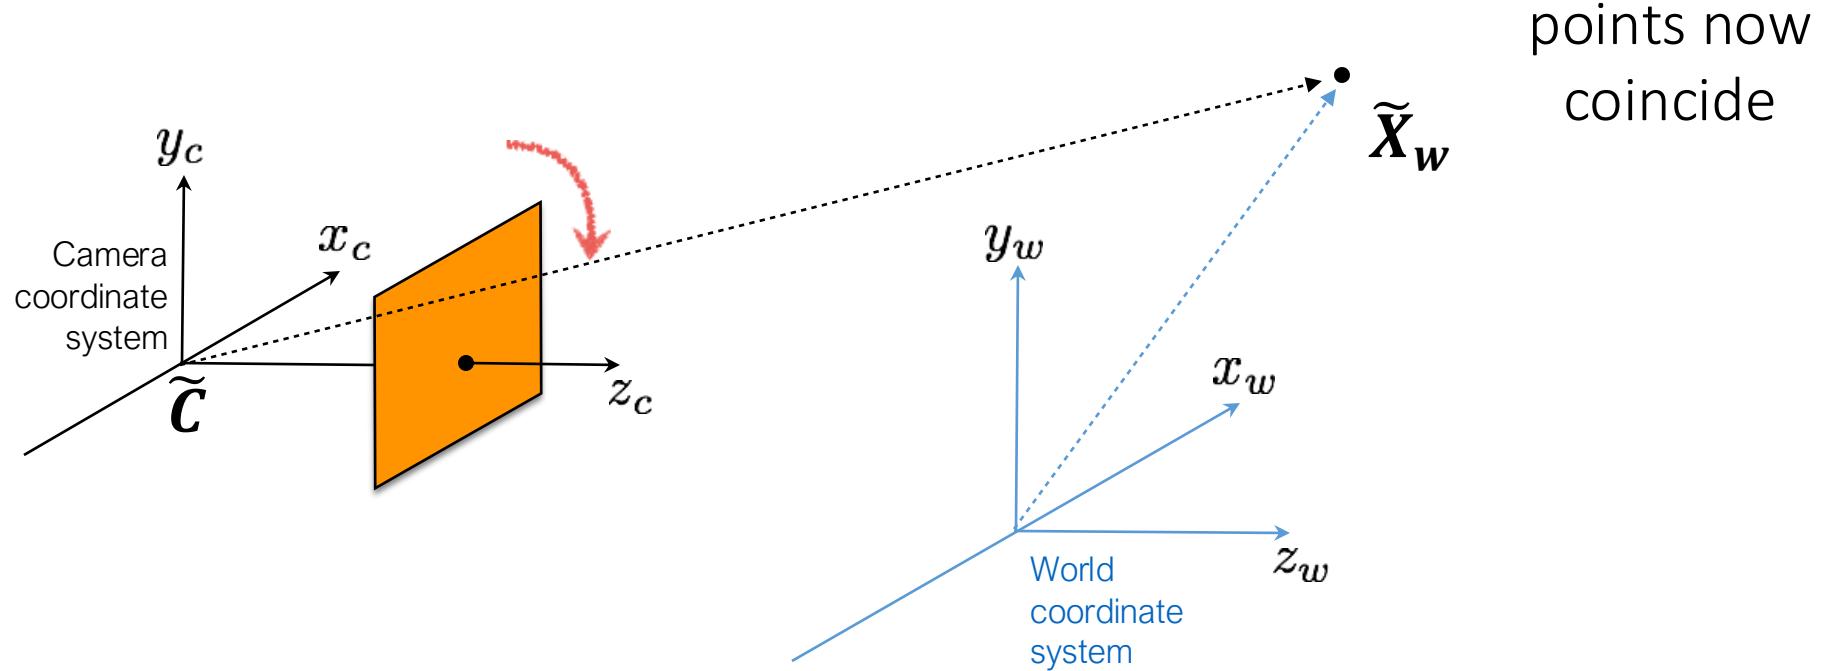
(homogeneous) transformation
from 2D to 2D, accounting for non
unit focal length and origin shift

(homogeneous) perspective projection
from 3D to 2D, assuming image plane at
 $z = 1$ and shared camera/image origin

Also written as: $\mathbf{P} = \mathbf{K}[\mathbf{I}|0]$

where $\mathbf{K} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix}$

Recap: World-to-camera coordinate system



$$R \cdot (\tilde{X}_w - \tilde{C})$$

rotate translate

Recap: Modeling the coordinate system

In heterogeneous coordinates, we have:

$$\tilde{\mathbf{X}}_c = \mathbf{R} \cdot (\tilde{\mathbf{X}}_w - \tilde{\mathbf{C}})$$

How do we write this transformation in homogeneous coordinates?

Recap: Modeling the coordinate system

In heterogeneous coordinates, we have:

$$\tilde{\mathbf{X}}_c = \mathbf{R} \cdot (\tilde{\mathbf{X}}_w - \tilde{\mathbf{C}})$$

In homogeneous coordinates, we have:

$$\begin{bmatrix} X_c \\ Y_c \\ Z_c \\ 1 \end{bmatrix} = \begin{bmatrix} \mathbf{R} & -\mathbf{R}\tilde{\mathbf{C}} \\ \mathbf{0} & 1 \end{bmatrix} \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix} \quad \text{or} \quad \mathbf{x}_c = \begin{bmatrix} \mathbf{R} & -\mathbf{R}\tilde{\mathbf{C}} \\ \mathbf{0} & 1 \end{bmatrix} \mathbf{x}_w$$

Recap: Putting it all together

We can write everything into a single projection:

$$\mathbf{x} = \mathbf{P}\mathbf{X}_w$$

The camera matrix now looks like:

$$\mathbf{P} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} [\mathbf{I} \quad | \quad \mathbf{0}] \begin{bmatrix} \mathbf{R} & -\mathbf{R}\tilde{\mathbf{C}} \\ \mathbf{0} & 1 \end{bmatrix}$$

intrinsic parameters (3×3):
correspond to camera
internals (image-to-image
transformation)

perspective projection (3×4):
maps 3D to 2D points
(camera-to-image
transformation)

extrinsic parameters (4×4):
correspond to camera
externals (world-to-camera
transformation)

Recap: Putting it all together

We can write everything into a single projection:

$$\mathbf{x} = \mathbf{P}\mathbf{X}_w$$

The camera matrix now looks like:

$$\mathbf{P} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \left[\begin{array}{c|c} \mathbf{R} & -\mathbf{RC} \end{array} \right]$$

intrinsic parameters (3×3):
correspond to camera internals
(sensor not at $f = 1$ and origin shift)

extrinsic parameters (3×4):
correspond to camera externals
(world-to-image transformation)

Recap: General pinhole camera matrix

$$\mathbf{P} = \mathbf{K}[\mathbf{R}|\mathbf{t}]$$

$$\mathbf{P} = \left[\begin{array}{ccc} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{array} \right] \left[\begin{array}{ccc|c} r_1 & r_2 & r_3 & t_1 \\ r_4 & r_5 & r_6 & t_2 \\ r_7 & r_8 & r_9 & t_3 \end{array} \right]$$

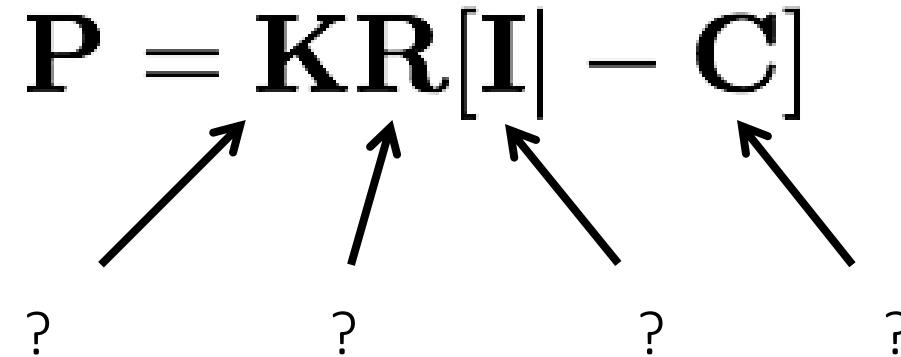
intrinsic extrinsic
parameters parameters

$$\mathbf{R} = \left[\begin{array}{ccc} r_1 & r_2 & r_3 \\ r_4 & r_5 & r_6 \\ r_7 & r_8 & r_9 \end{array} \right] \quad \mathbf{t} = \left[\begin{array}{c} t_1 \\ t_2 \\ t_3 \end{array} \right]$$

3D rotation 3D translation

Recap

What is the size and meaning of each term in the camera matrix?

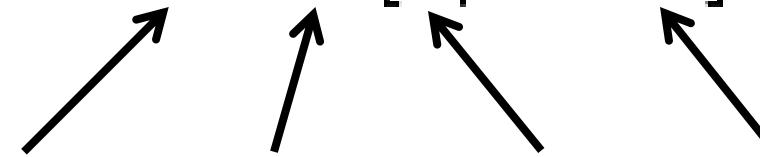
$$\mathbf{P} = \mathbf{K}\mathbf{R}[\mathbf{I}] - \mathbf{C}$$


The diagram consists of a mathematical equation $\mathbf{P} = \mathbf{K}\mathbf{R}[\mathbf{I}] - \mathbf{C}$ centered on the page. Below the equation, there are four question marks positioned under the terms \mathbf{K} , \mathbf{R} , $[\mathbf{I}]$, and \mathbf{C} . Four black arrows originate from these question marks and point diagonally upwards towards their respective terms in the equation.

Recap

What is the size and meaning of each term in the camera matrix?

$$\mathbf{P} = \mathbf{K}\mathbf{R}[\mathbf{I}] - \mathbf{C}$$



3x3

?

?

?

intrinsics

Recap

What is the size and meaning of each term in the camera matrix?

$$\mathbf{P} = \mathbf{K}\mathbf{R}[\mathbf{I}] - \mathbf{C}$$

The diagram illustrates the components of the camera matrix \mathbf{P} . Four arrows point from labels at the bottom to specific terms in the equation:

- An arrow points from "3x3 intrinsics" to the matrix \mathbf{K} .
- An arrow points from "3x3 3D rotation" to the matrix \mathbf{R} .
- An arrow points from "?" to the term $[\mathbf{I}]$.
- An arrow points from "?" to the matrix \mathbf{C} .

Recap

What is the size and meaning of each term in the camera matrix?

$$\mathbf{P} = \mathbf{K}\mathbf{R}[\mathbf{I}] - \mathbf{C}$$

The diagram illustrates the components of the camera matrix \mathbf{P} . The equation is $\mathbf{P} = \mathbf{K}\mathbf{R}[\mathbf{I}] - \mathbf{C}$. Four arrows point from labels below the equation to its terms: the first arrow points to \mathbf{K} with the label "3x3 intrinsics"; the second arrow points to \mathbf{R} with the label "3x3 3D rotation"; the third arrow points to $[\mathbf{I}]$ with the label "3x3 identity"; and the fourth arrow points to $-\mathbf{C}$ with a question mark "?".

Recap

What is the size and meaning of each term in the camera matrix?

$$\mathbf{P} = \mathbf{K}\mathbf{R}[\mathbf{I}] - \mathbf{C}$$

3x3

3x3

3x3

3x1

intrinsics 3D rotation identity 3D translation

Quiz

The camera matrix relates what two quantities?

Quiz

The camera matrix relates what two quantities?

$$\mathbf{x} = \mathbf{P}\mathbf{X}$$

homogeneous 3D points to 2D image points

Quiz

The camera matrix relates what two quantities?

$$\mathbf{x} = \mathbf{P}\mathbf{X}$$

homogeneous 3D points to 2D image points

The camera matrix can be decomposed into?

Quiz

The camera matrix relates what two quantities?

$$\mathbf{x} = \mathbf{P}\mathbf{X}$$

homogeneous 3D points to 2D image points

The camera matrix can be decomposed into?

$$\mathbf{P} = \mathbf{K}[\mathbf{R}|\mathbf{t}]$$

intrinsic and extrinsic parameters

More general camera matrices

The following is the standard camera matrix we saw.

$$\mathbf{P} = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} \quad [\mathbf{R} \quad \vdots \quad -\mathbf{RC}]$$

More general camera matrices

CCD camera: pixels may not be square.

Lens: focal length along x and y may be different

$$\mathbf{P} = \begin{bmatrix} \alpha_x & 0 & p_x \\ 0 & \alpha_y & p_y \\ 0 & 0 & 1 \end{bmatrix} \quad \left[\begin{array}{c|c} \mathbf{R} & -\mathbf{RC} \end{array} \right]$$

How many degrees of freedom?



Spherical



Anamorphic Lens



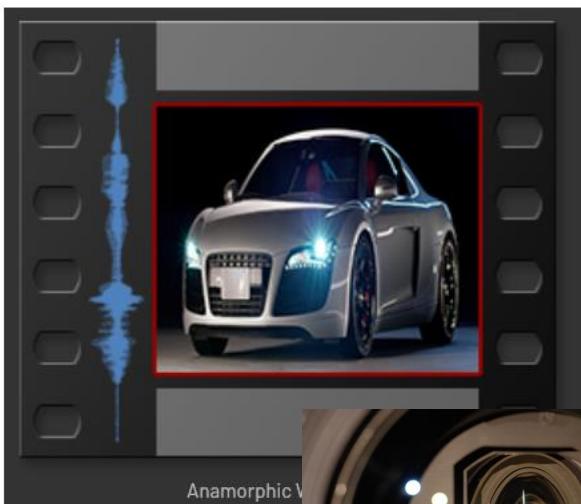
Lens flares are scattered light that you usually try to avoid having in photos. It's usually caused when a bright light shines into the lens. [Paramount Pictures](#)



Abrams loves them. [Paramount Pictures](#)



Widescreen



Anamorphic



You can see them across his rebooted "Star Trek" movies. [Paramount Pictures](#)

More general camera matrices

CCD camera: pixels may not be square.

Lens: focal length along x and y may be different

$$\mathbf{P} = \begin{bmatrix} \alpha_x & 0 & p_x \\ 0 & \alpha_y & p_y \\ 0 & 0 & 1 \end{bmatrix} \quad \left[\begin{array}{c|c} \mathbf{R} & -\mathbf{RC} \end{array} \right]$$

How many degrees of freedom?

10 DOF

More general camera matrices

Finite projective camera: sensor be skewed.

$$\mathbf{P} = \begin{bmatrix} \alpha_x & s & p_x \\ 0 & \alpha_y & p_y \\ 0 & 0 & 1 \end{bmatrix} \left[\begin{array}{c|c} \mathbf{R} & -\mathbf{RC} \end{array} \right]$$

How many degrees of freedom?

More general camera matrices

Finite projective camera: sensor be skewed.

$$\mathbf{P} = \begin{bmatrix} \alpha_x & s & p_x \\ 0 & \alpha_y & p_y \\ 0 & 0 & 1 \end{bmatrix} \left[\begin{array}{c|c} \mathbf{R} & -\mathbf{RC} \end{array} \right]$$

How many degrees of freedom?

11 DOF

Projection matrix

$$\Pi = \mathbf{K} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} \mathbf{R} & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \mathbf{I}_{3 \times 3} & -\mathbf{c} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

intrinsics projection rotation translation

The **K** matrix converts 3D rays in the camera's coordinate system to 2D image points in image (pixel) coordinates.

This part converts 3D points in world coordinates to 3D rays in the camera's coordinate system. There are 6 parameters represented (3 for position/translation, 3 for rotation).

Projection matrix

$$\Pi = \mathbf{K} \underbrace{\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}}_{\text{intrinsics}} \underbrace{\begin{bmatrix} \mathbf{R} & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}}_{\text{rotation}} \underbrace{\begin{bmatrix} \mathbf{I}_{3 \times 3} & -\mathbf{c} \\ 0 & 0 & 0 & 1 \end{bmatrix}}_{\text{translation}}$$

$$\left[\begin{array}{c|c} \mathbf{R} & -\mathbf{R}\mathbf{c} \end{array} \right]$$



(\mathbf{t} in book's
notation)

$$\Pi = \mathbf{K} \left[\begin{array}{c|c} \mathbf{R} & -\mathbf{R}\mathbf{c} \end{array} \right]$$

Perspective projection

$$\underbrace{\begin{bmatrix} f & 0 & c_x \\ 0 & f & c_y \\ 0 & 0 & 1 \end{bmatrix}}_{\mathbf{K} \text{ (intrinsics)}} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

\mathbf{K} (converts from 3D rays in camera coordinate system to pixel coordinates)

in general, $\mathbf{K} = \begin{bmatrix} f & s & c_x \\ 0 & \alpha f & c_y \\ 0 & 0 & 1 \end{bmatrix}$ (upper triangular matrix)

α : **aspect ratio** (1 unless pixels are not square)

s : **skew** (0 unless pixels are shaped like rhombi/parallelograms)

(c_x, c_y) : **principal point** ((w/2,h/2) unless optical axis doesn't intersect projection plane at image center)

Focal length

- Can think of as “zoom”



24mm



50mm



200mm

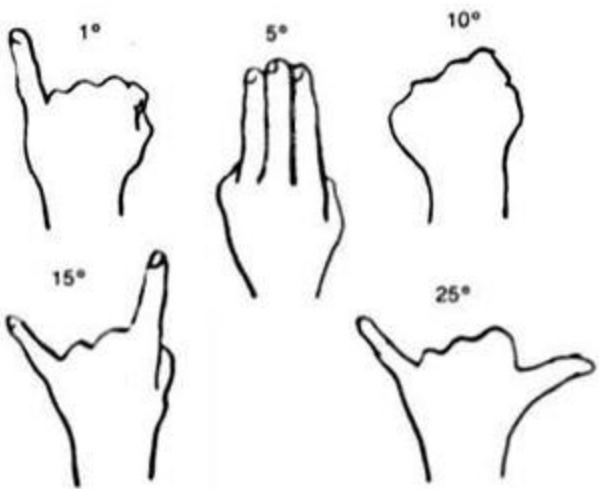


800mm



- Also related to *field of view*

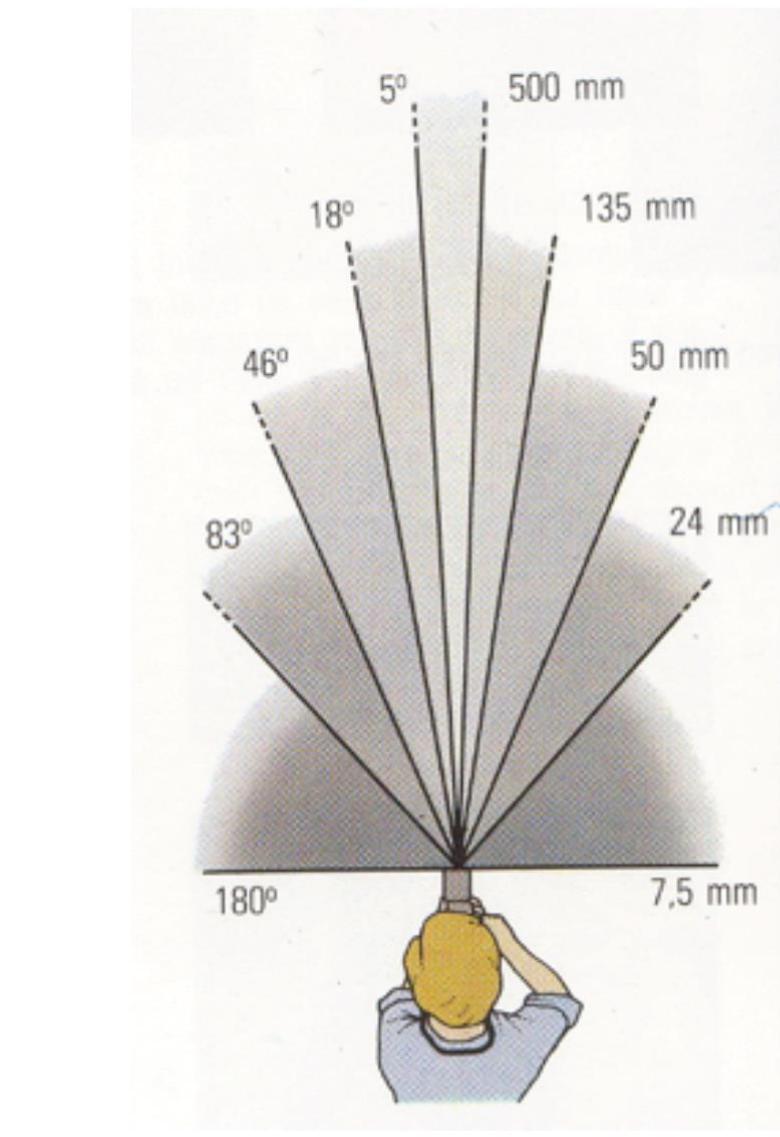
Field of view



APS-C Crop Body Measurement Table

Lens	After 1.62 Multiplier	APS-C Sensor (1.62 lens multiplier) Canon 60D, 7D, 70D, T3i, T4i	Hand Positions
18mm	29.16mm	Three hands wide at full arms length.	18mm Lens on Canon APS-C Sensor = 29.16mm Frame Width = Three Hands Wide
28mm	45.36mm	Slightly less than two hands wide at full arms length.	28mm Lens on Canon APS-C Sensor = 45.36mm Frame Width = Slightly less than Two Hands Wide
35mm	56.7mm	One hand + width of one fist at full arms length.	35mm Lens on Canon APS-C Sensor = 56.7mm Frame Width = One Hand + Width of One Fist
50mm	81mm	One hand wide + width of thumb at full arms length.	50mm Lens on Canon APS-C Sensor = 81mm Frame Width = One Hand wide + Width of Thumb
55mm	89.1mm	Slightly less than one hand wide at full arms length.	55mm Lens on Canon APS-C Sensor = 89.1mm Frame Width = Slightly less than One Hand wide
85mm	137.7mm	Inside edge of thumb to tip of forefinger wide with hand in "L" shape, thumb up.	85mm Lens on Canon APS-C Sensor = 137.7mm Frame Width = Inside Edge of Thumb to Tip of Forefinger Wide

Focal length in practice



24mm



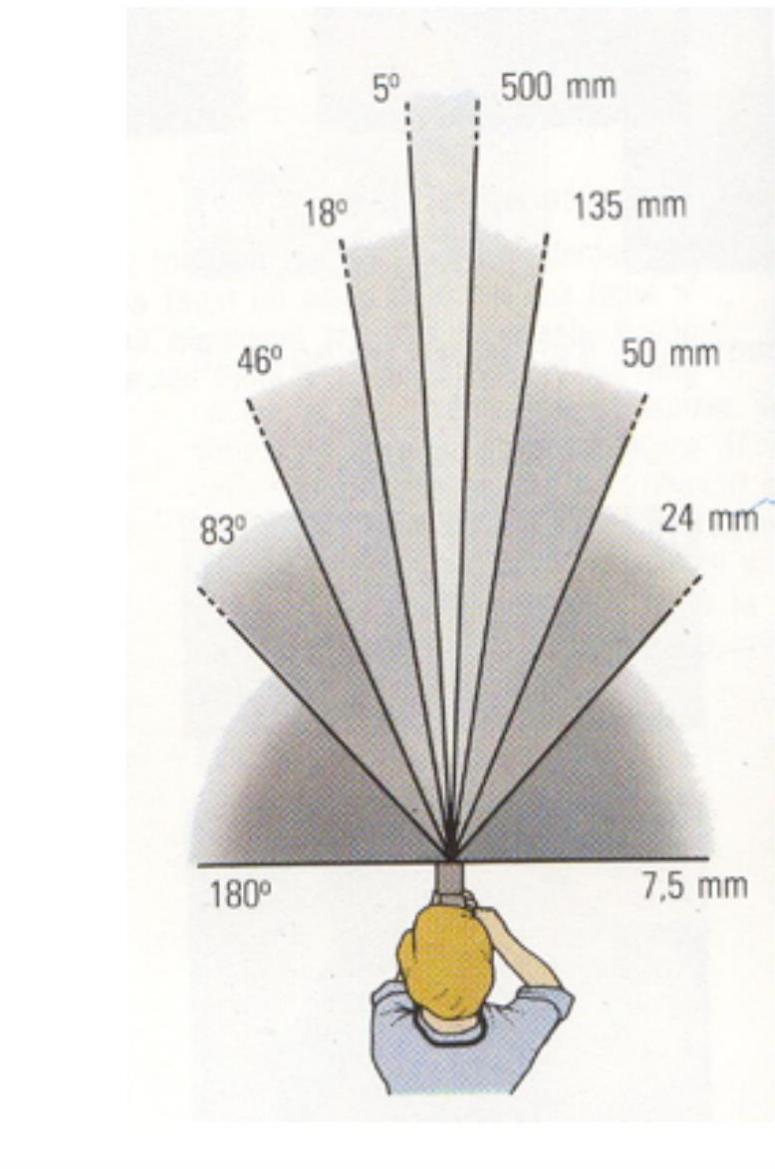
50mm



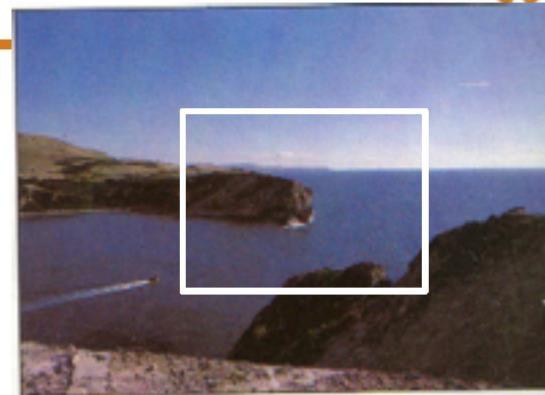
135mm



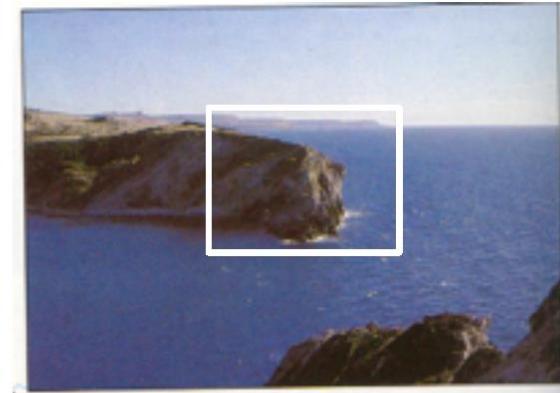
Focal length = cropping



24mm



50mm

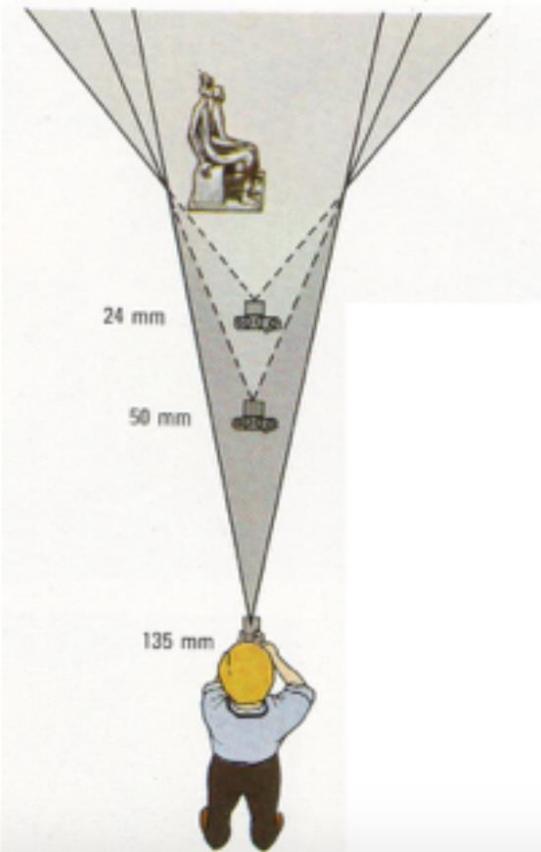


135mm



Focal length vs. viewpoint

- Telephoto makes it easier to select background (a small change in viewpoint is a big change in background).



Fredo Durand

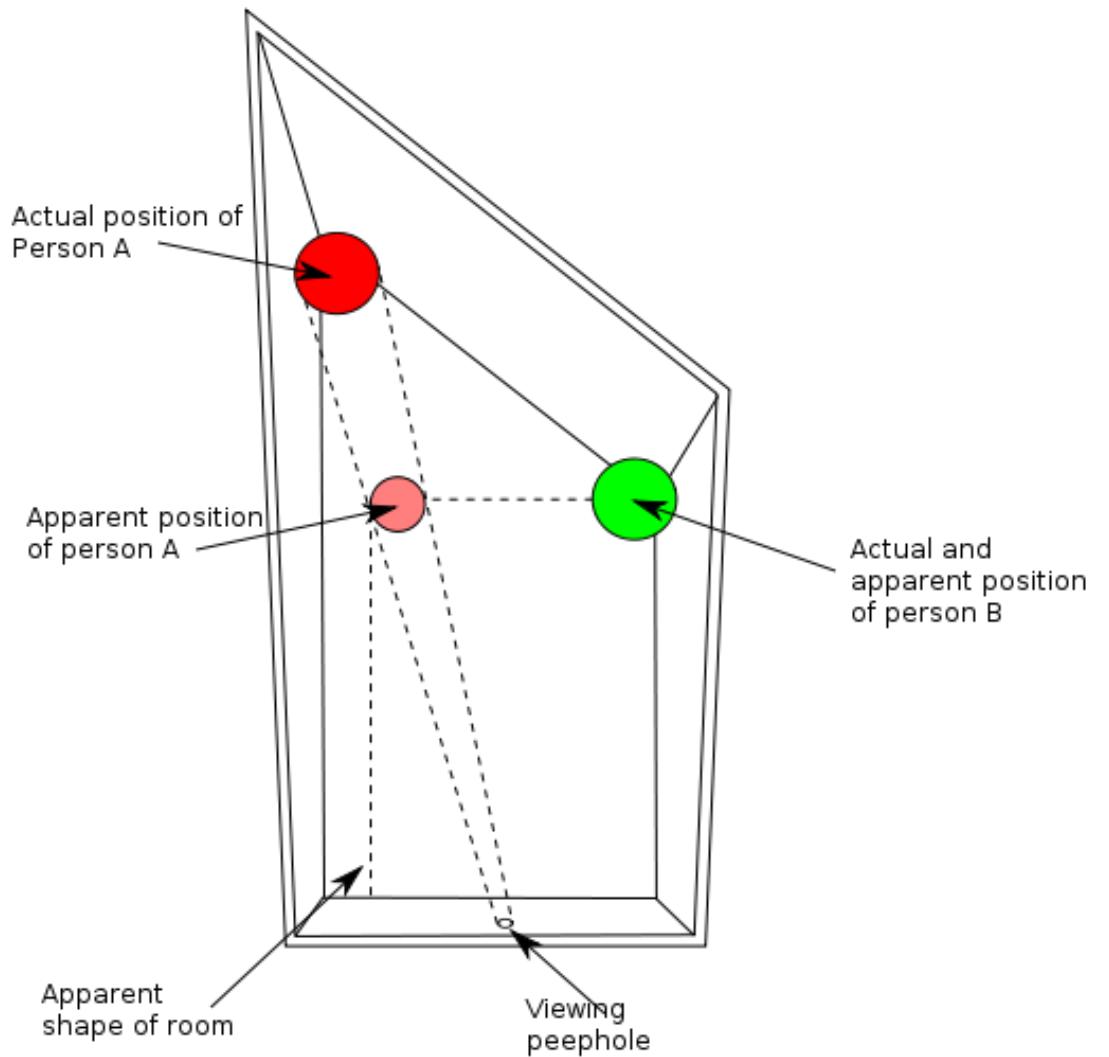
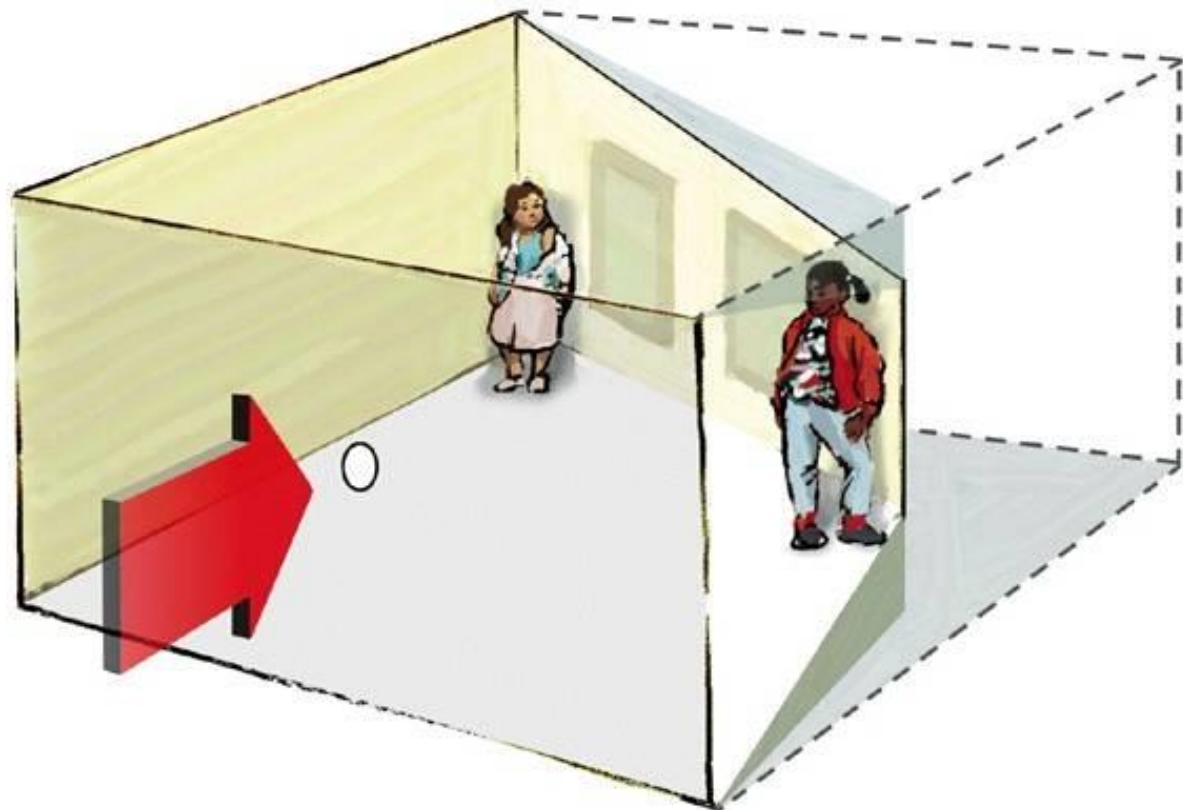
Forced perspective



The Ames room illusion



The Ames room illusion



The arrow illusion



Forced Perspective Display

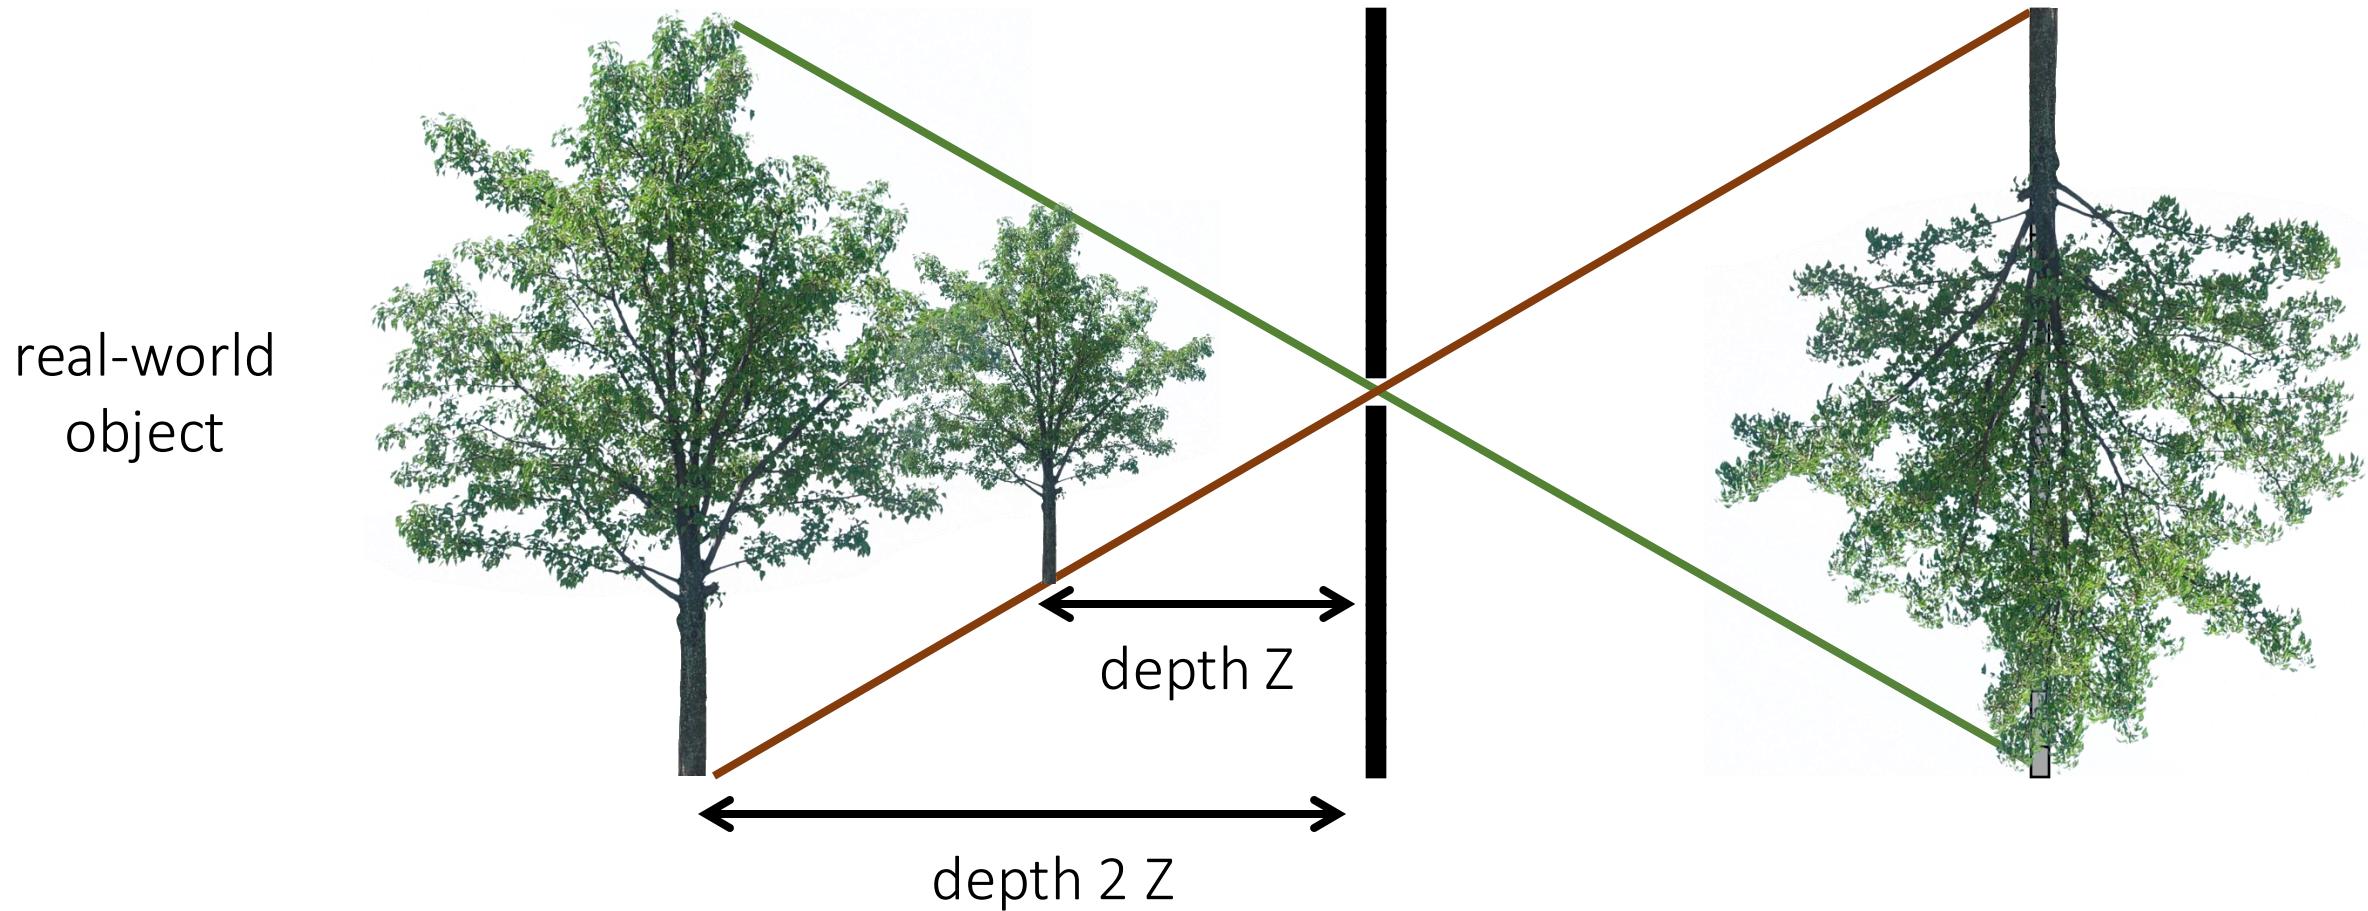


Forced Perspective Display

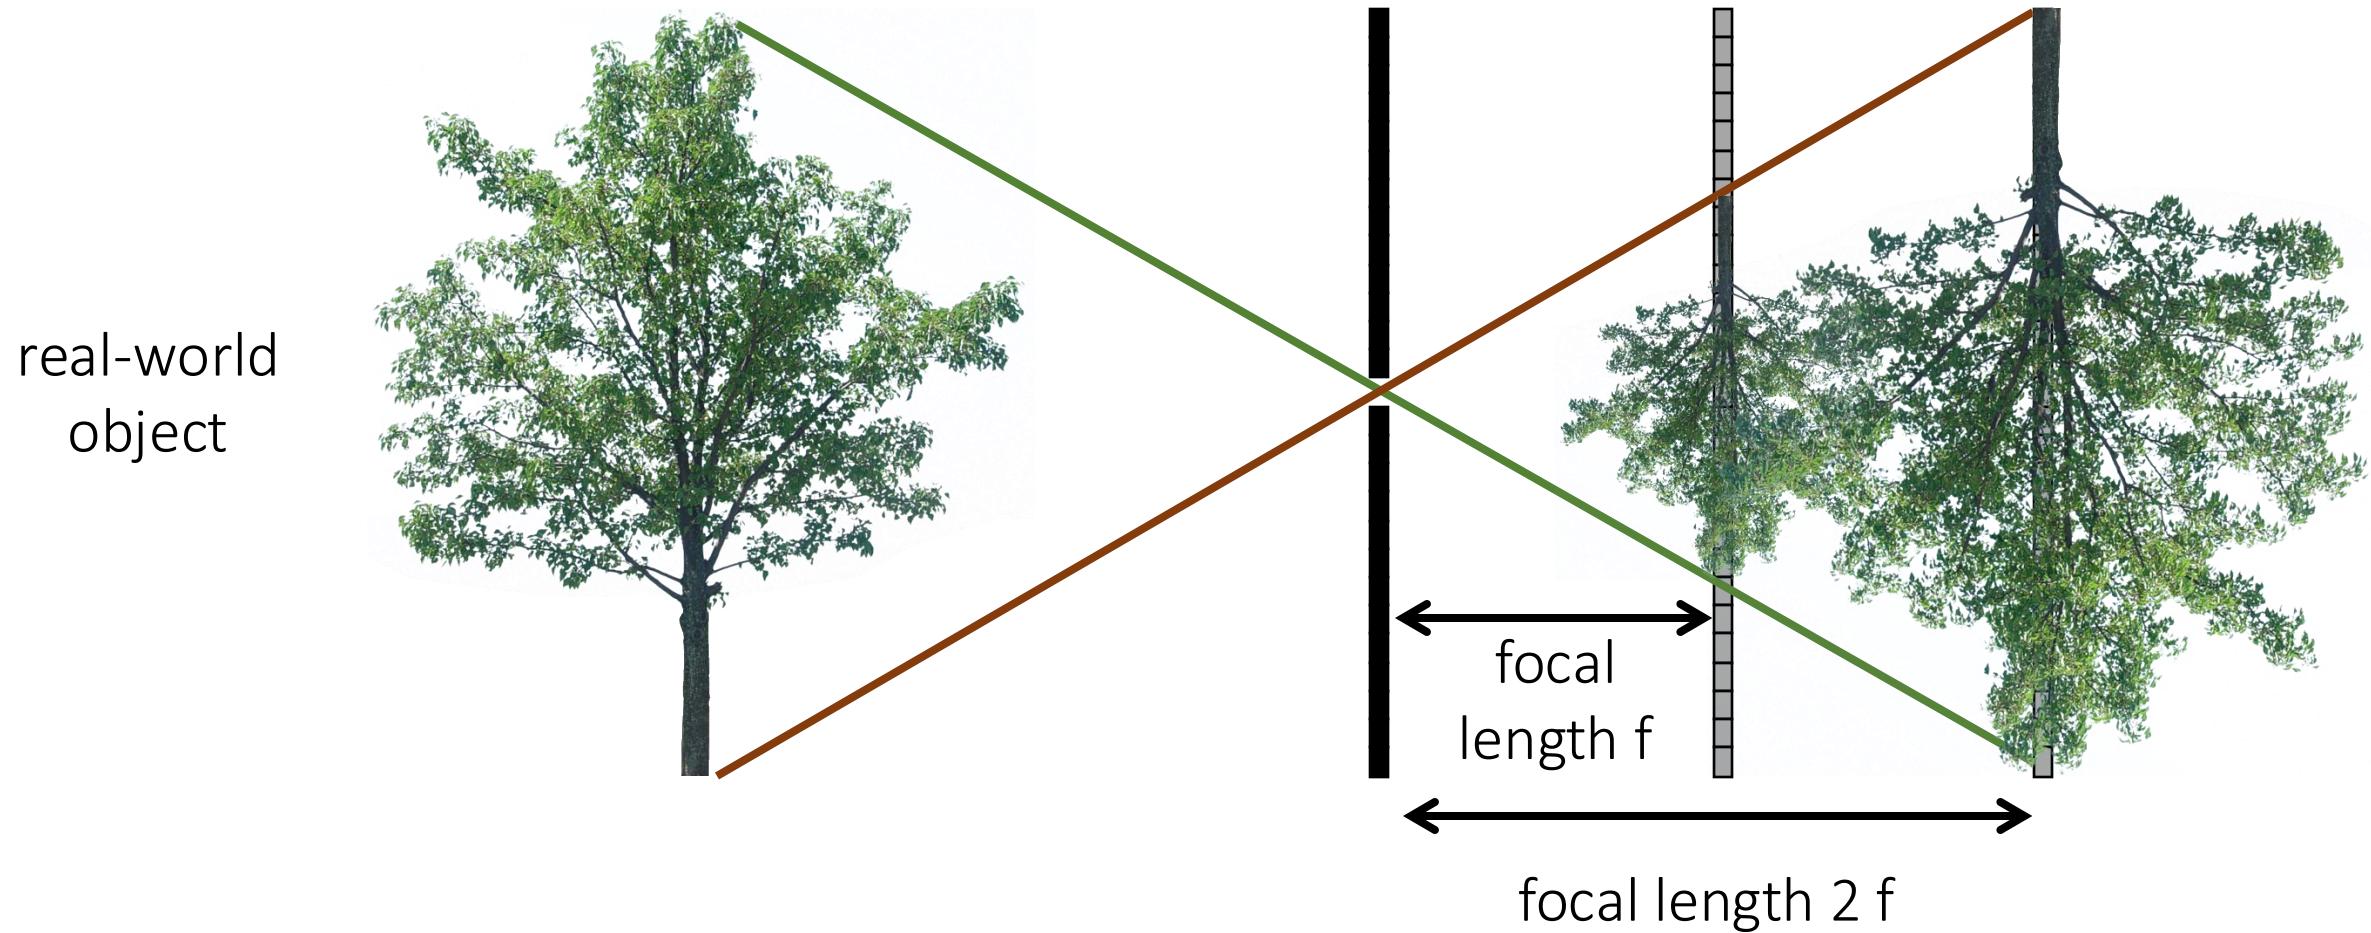


Magnification depends on depth

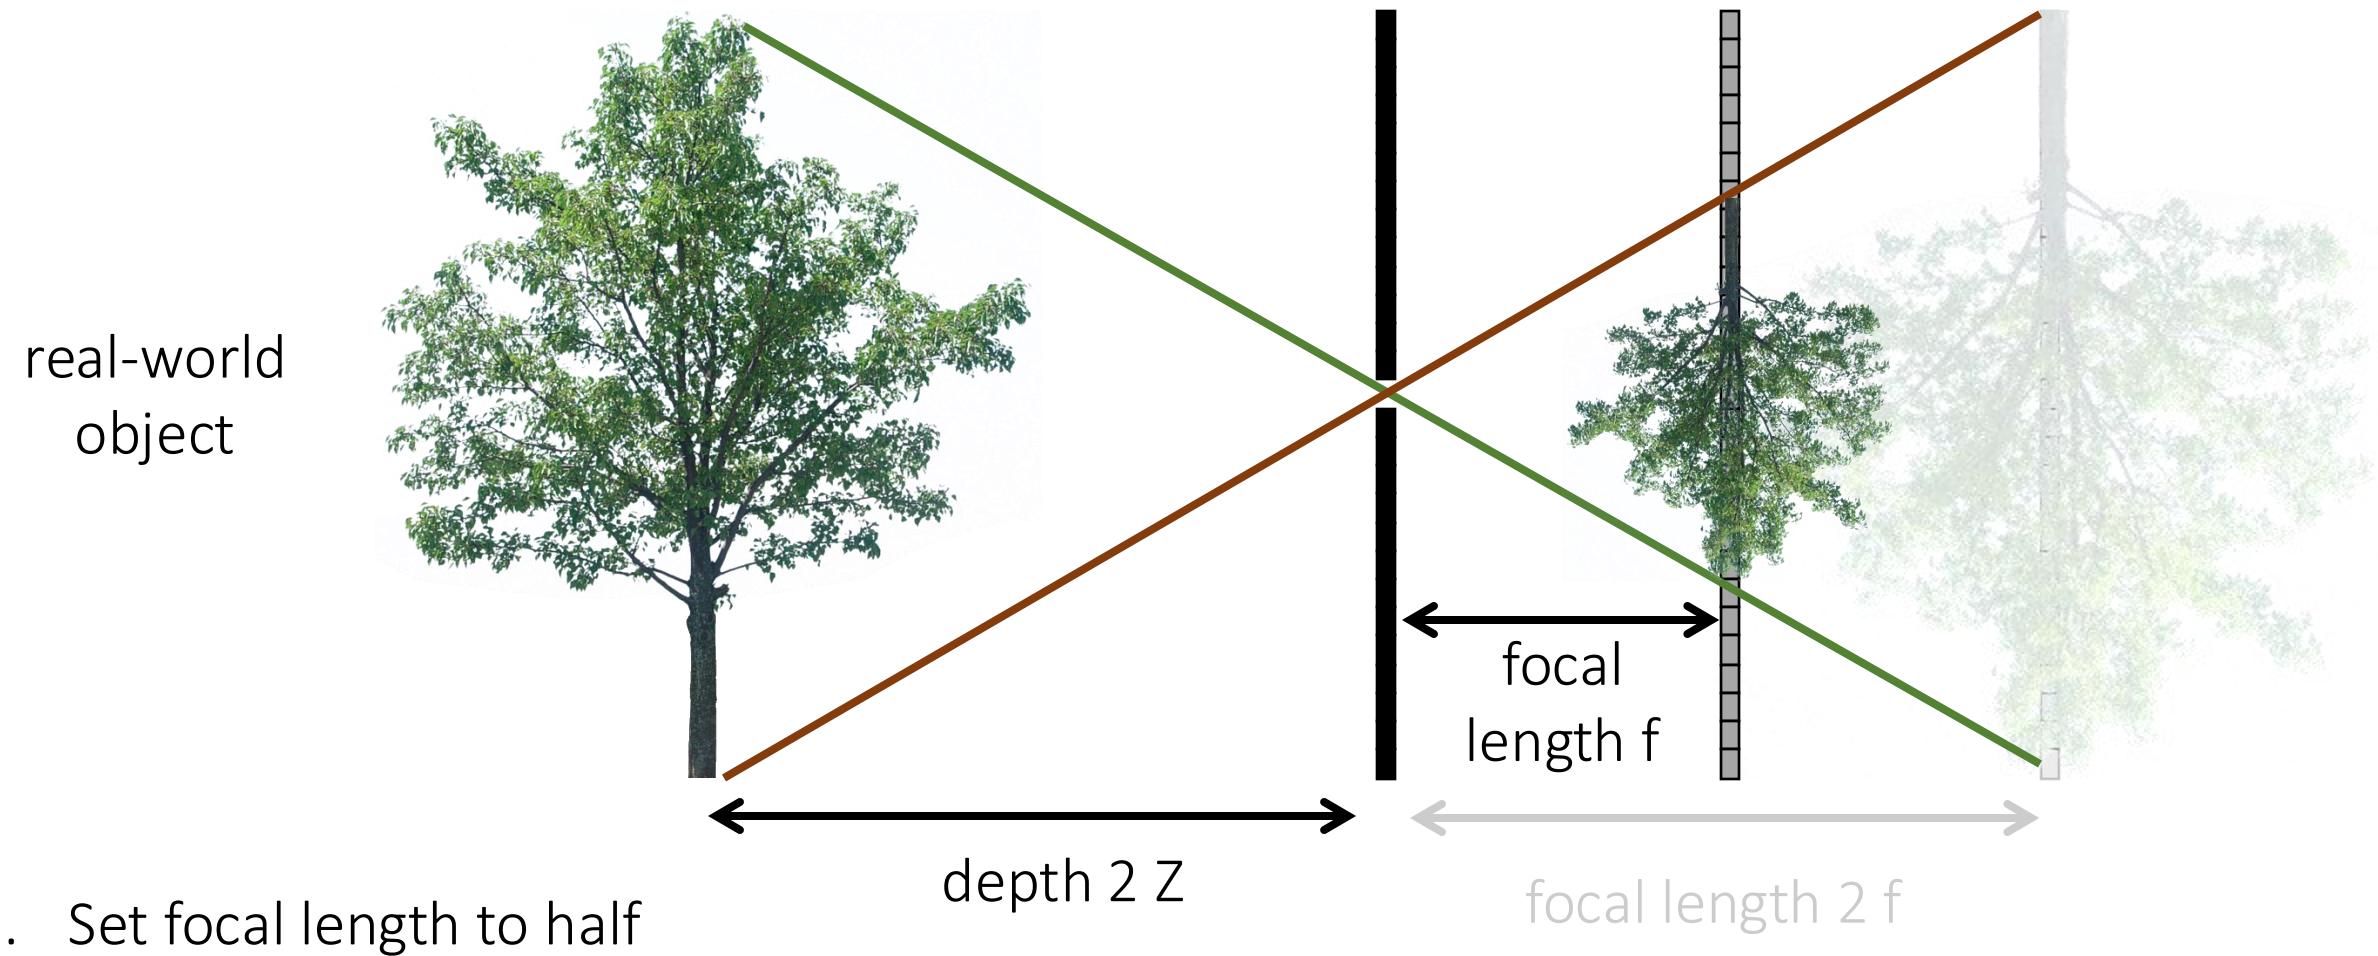
What happens as we change the focal length?



Magnification depends on focal length

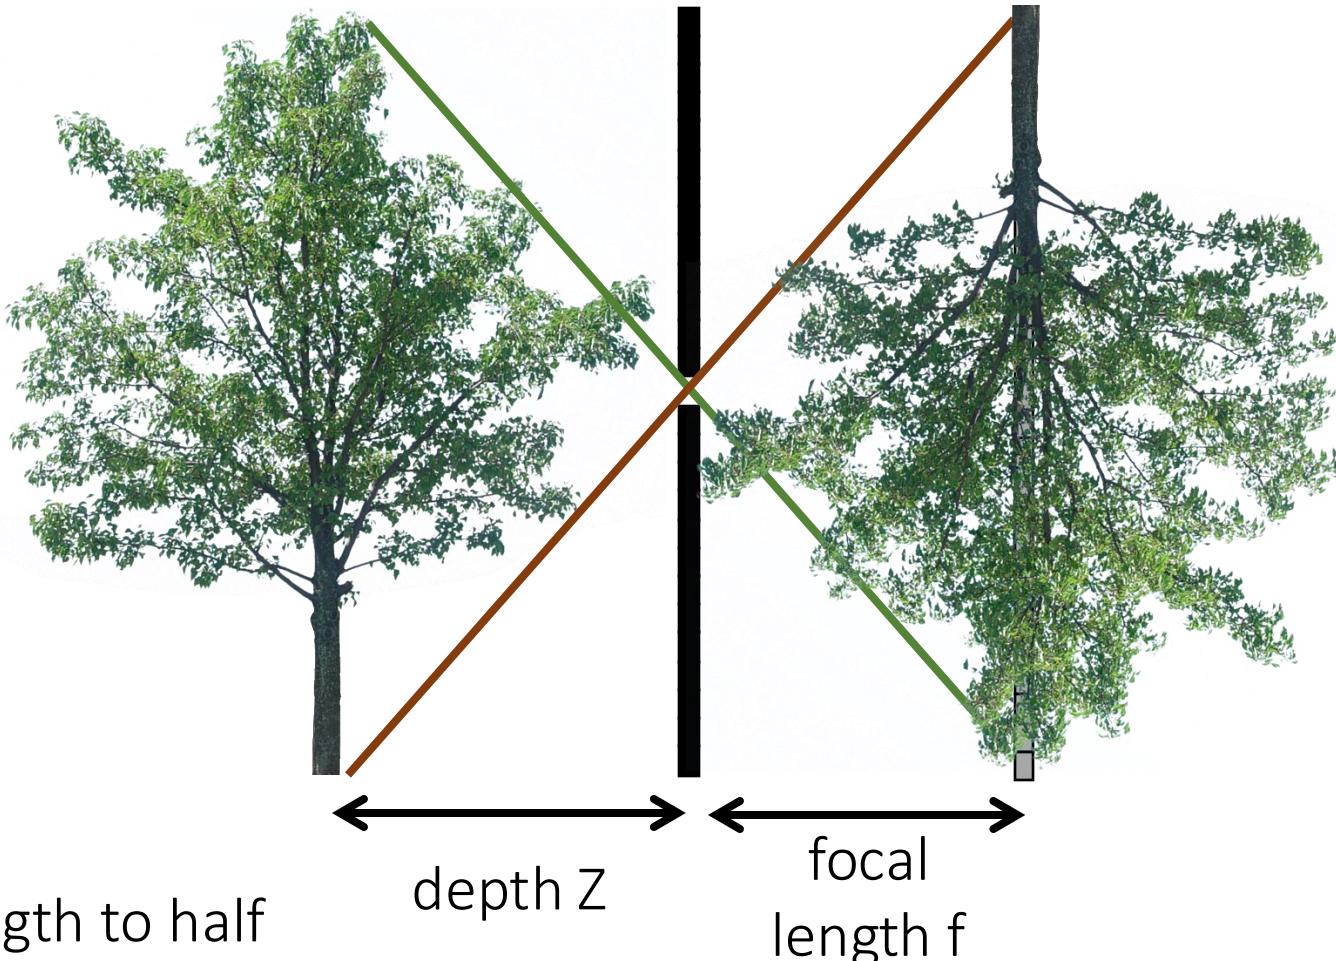


What if...



What if...

real-world
object



Is this the same image as
the one I had at focal
length $2f$ and distance $2Z$?

1. Set focal length to half
2. Set depth to half

Vertigo effect

Named after Alfred Hitchcock's movie

- also known as “dolly zoom”



Vertigo effect



How would you
create this effect?

Vertigo effect



Perspective distortion



long focal length

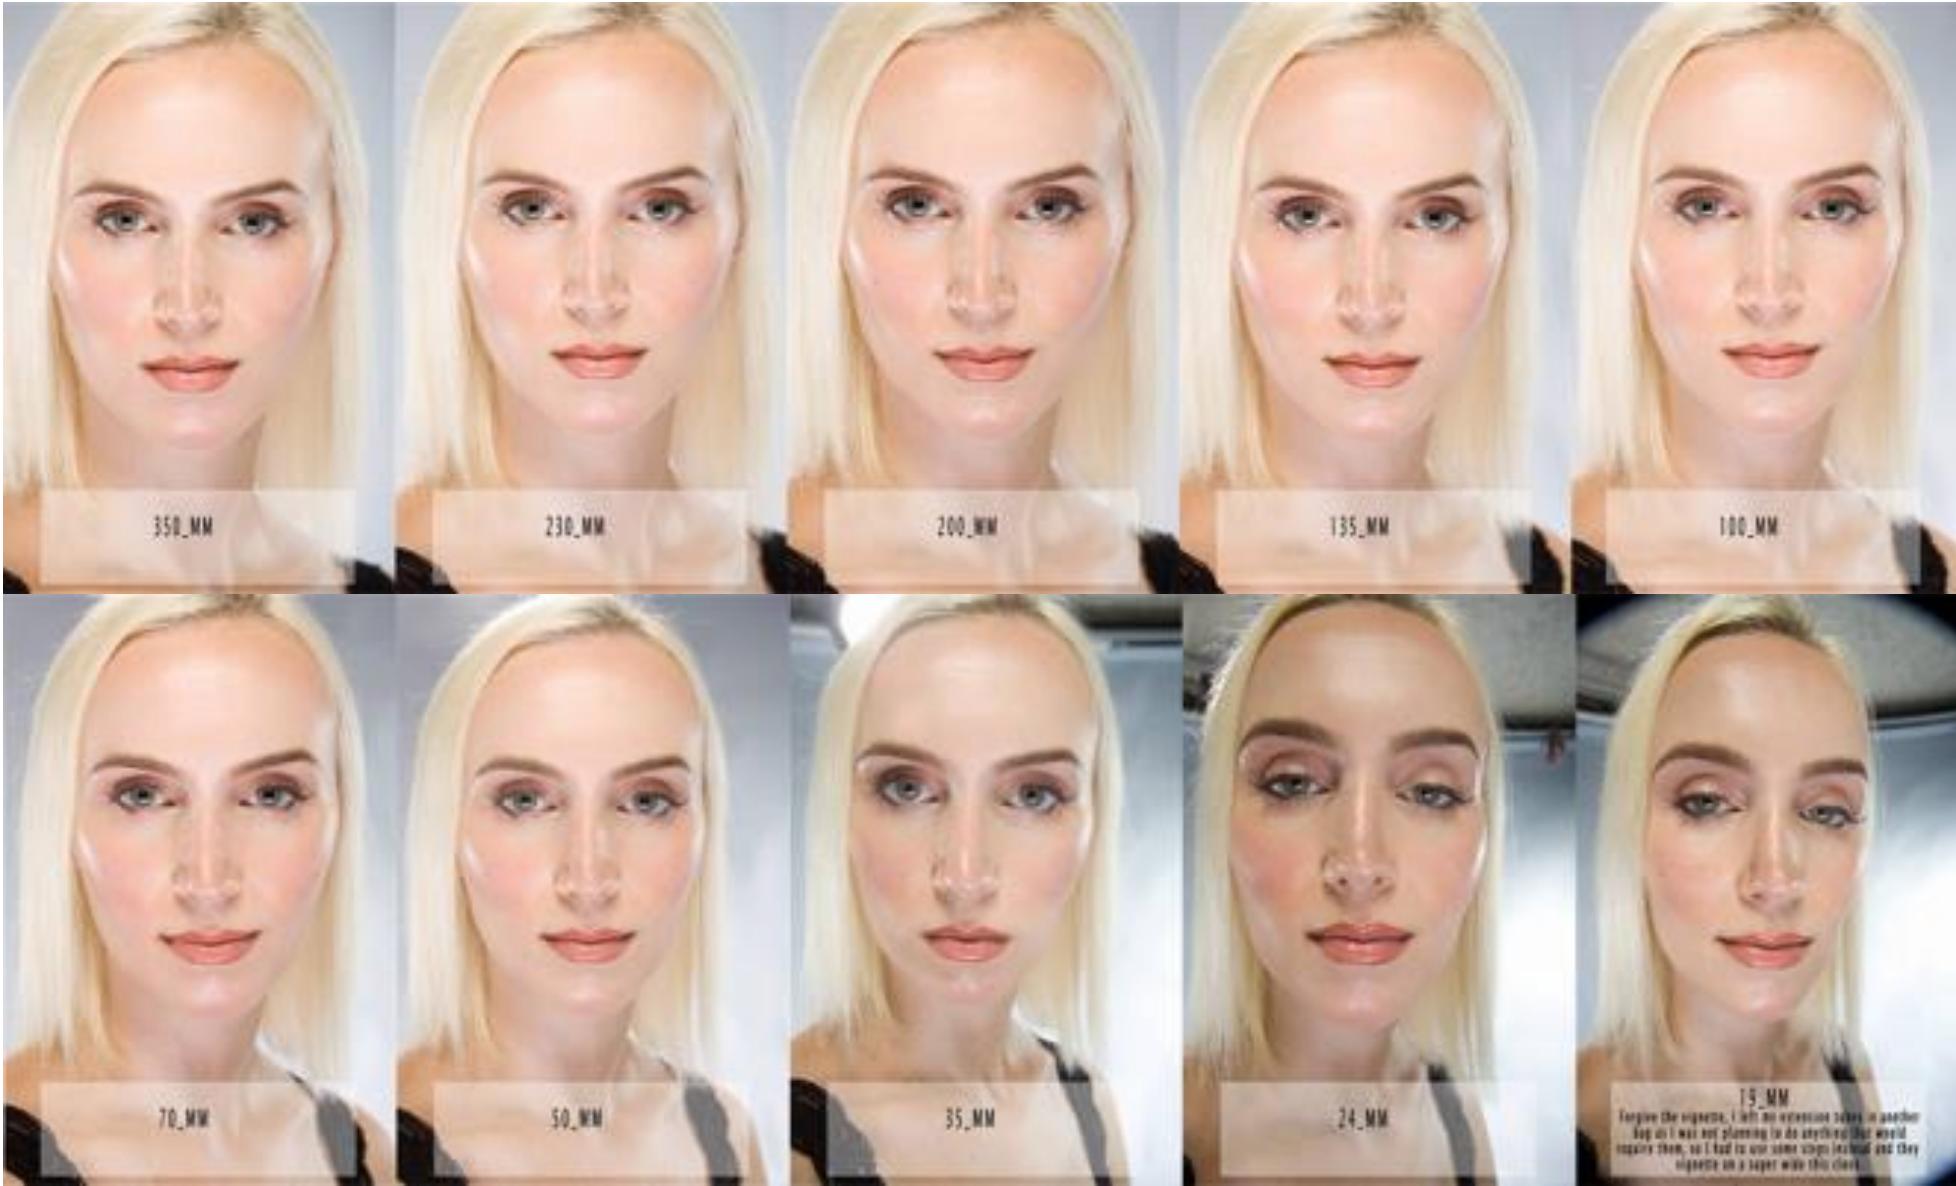


mid focal length



short focal length

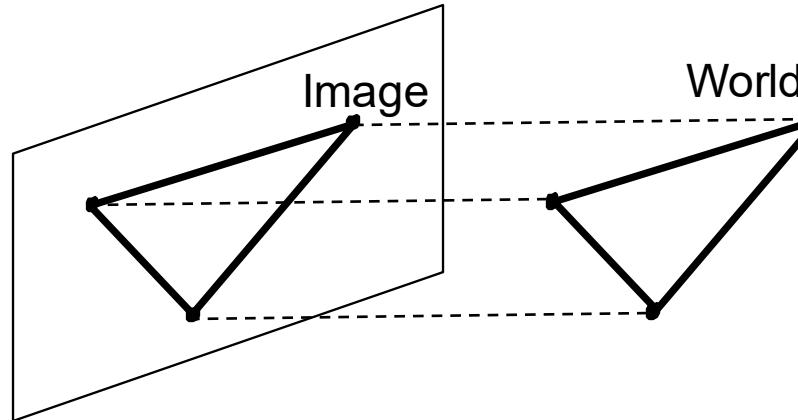
Perspective distortion



Other camera models

Orthographic projection

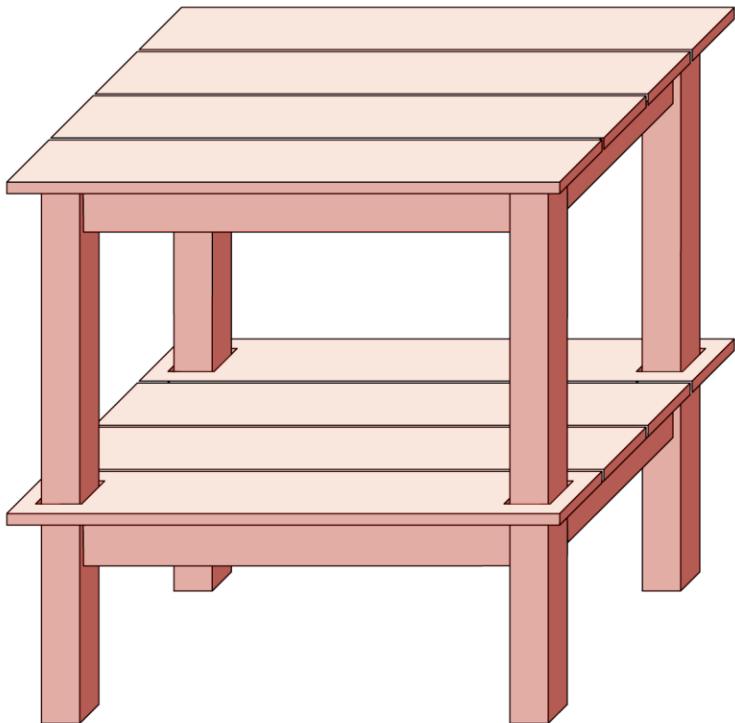
- Special case of perspective projection
 - Distance from the COP to the PP is infinite



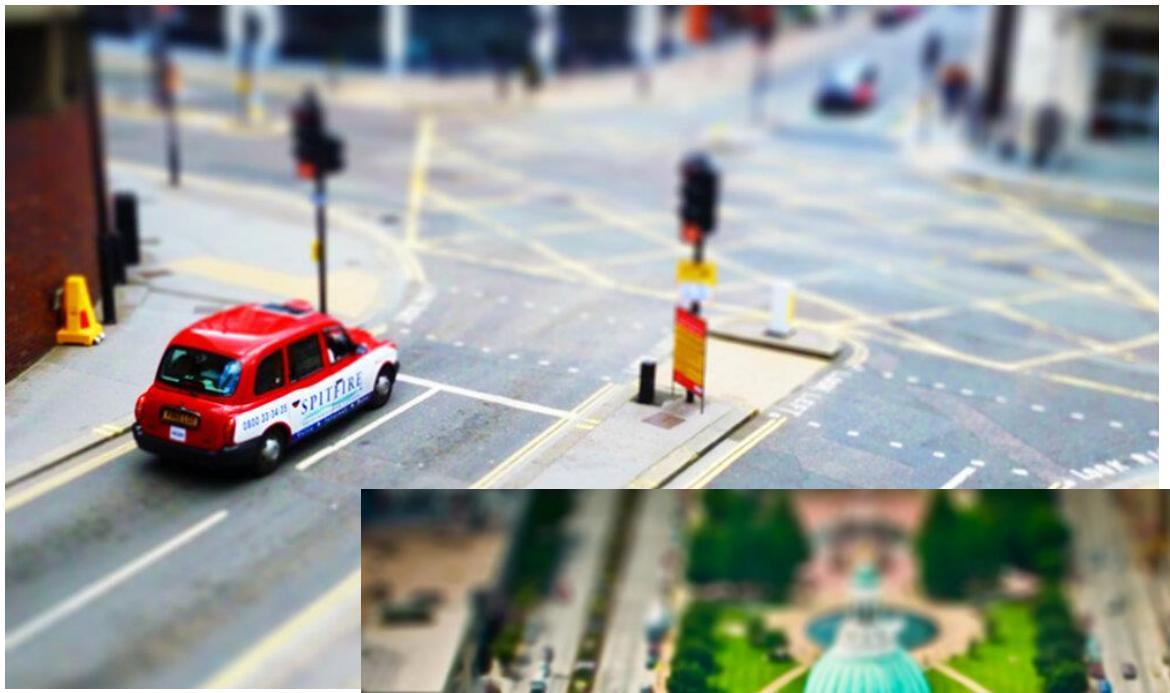
- Good approximation for telephone optics

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \Rightarrow (x, y)$$

Orthographic projection



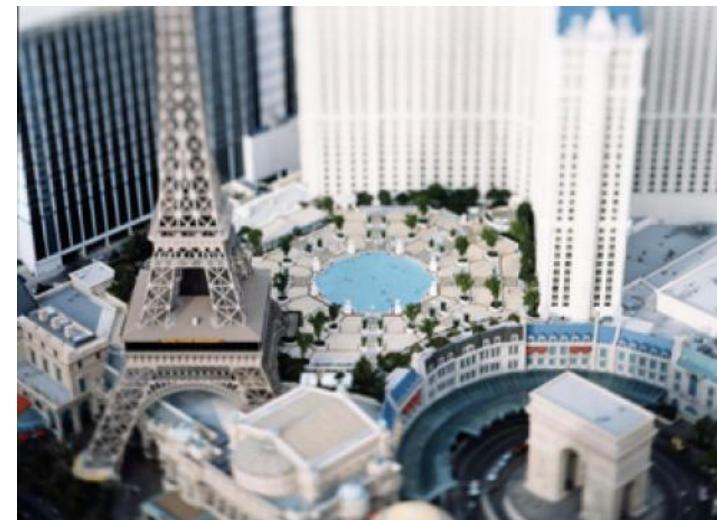
How's this related to
"miniature" effect by
tilt-shift photography?



Tilt-shift

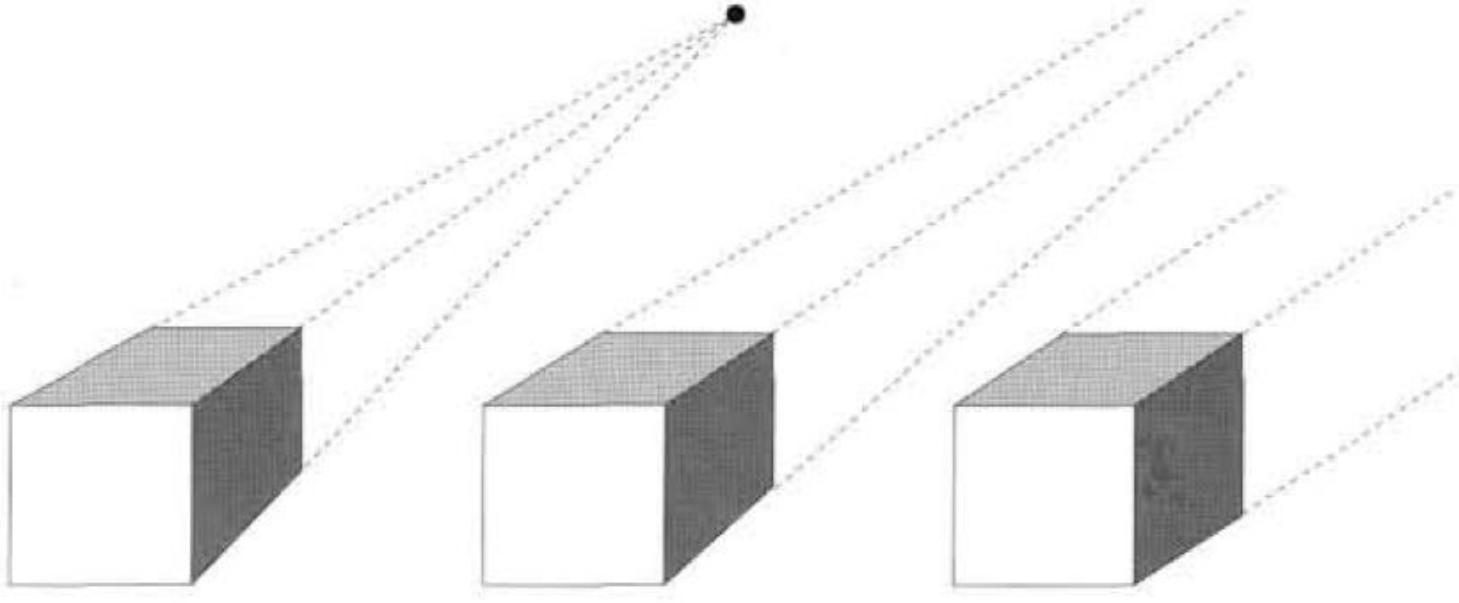


http://www.northlight-images.co.uk/article_pages/tilt_and_shift_ts-e.html



Tilt-shift images from [Olivo Barbieri](#)
and Photoshop [imitations](#)

camera is *close*
to object and has
small focal length



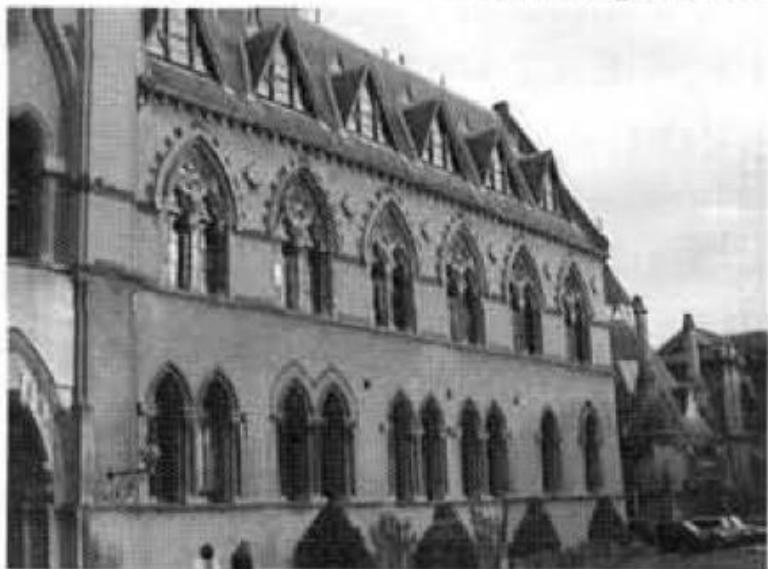
perspective

camera is *far* from
object and has
large focal length

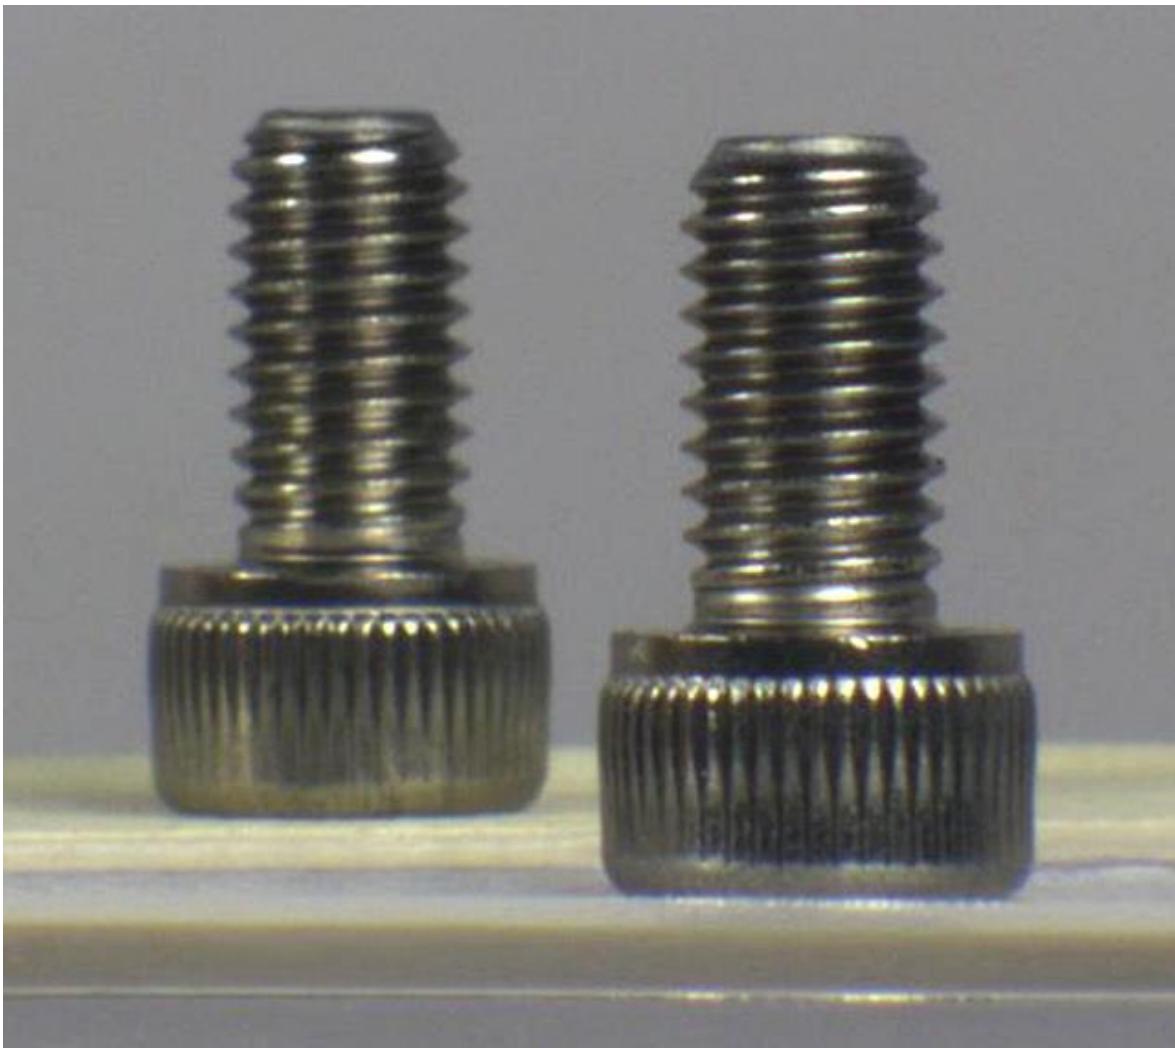
weak perspective

increasing focal length →

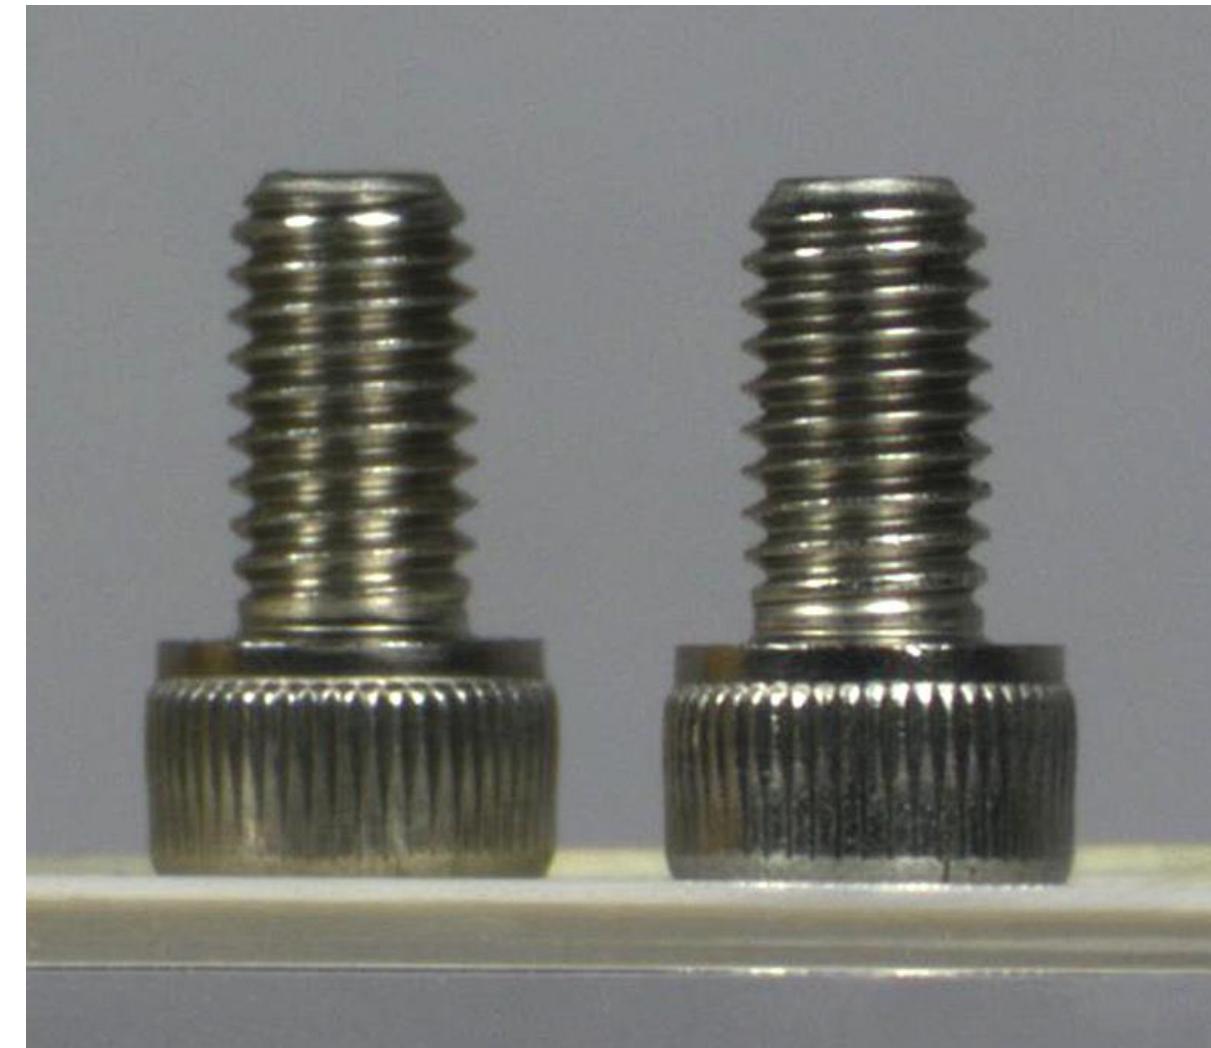
increasing distance from camera →



Different cameras



perspective camera



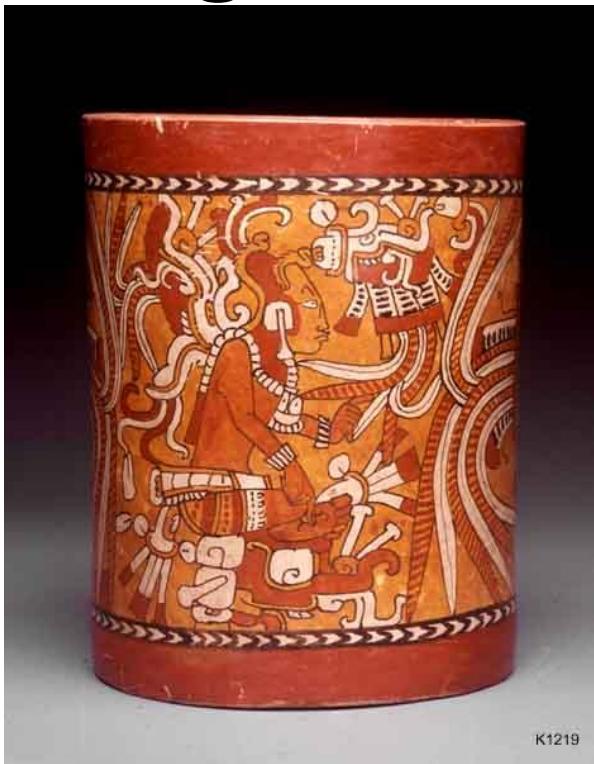
weak perspective camera

360 degree field of view...



- Basic approach
 - Take a photo of a parabolic mirror with an orthographic lens (Nayar)
 - Or buy one a lens from a variety of omnicam manufacturers...
 - See <http://www.cis.upenn.edu/~kostas/omni.html>

Rotating sensor (or object)

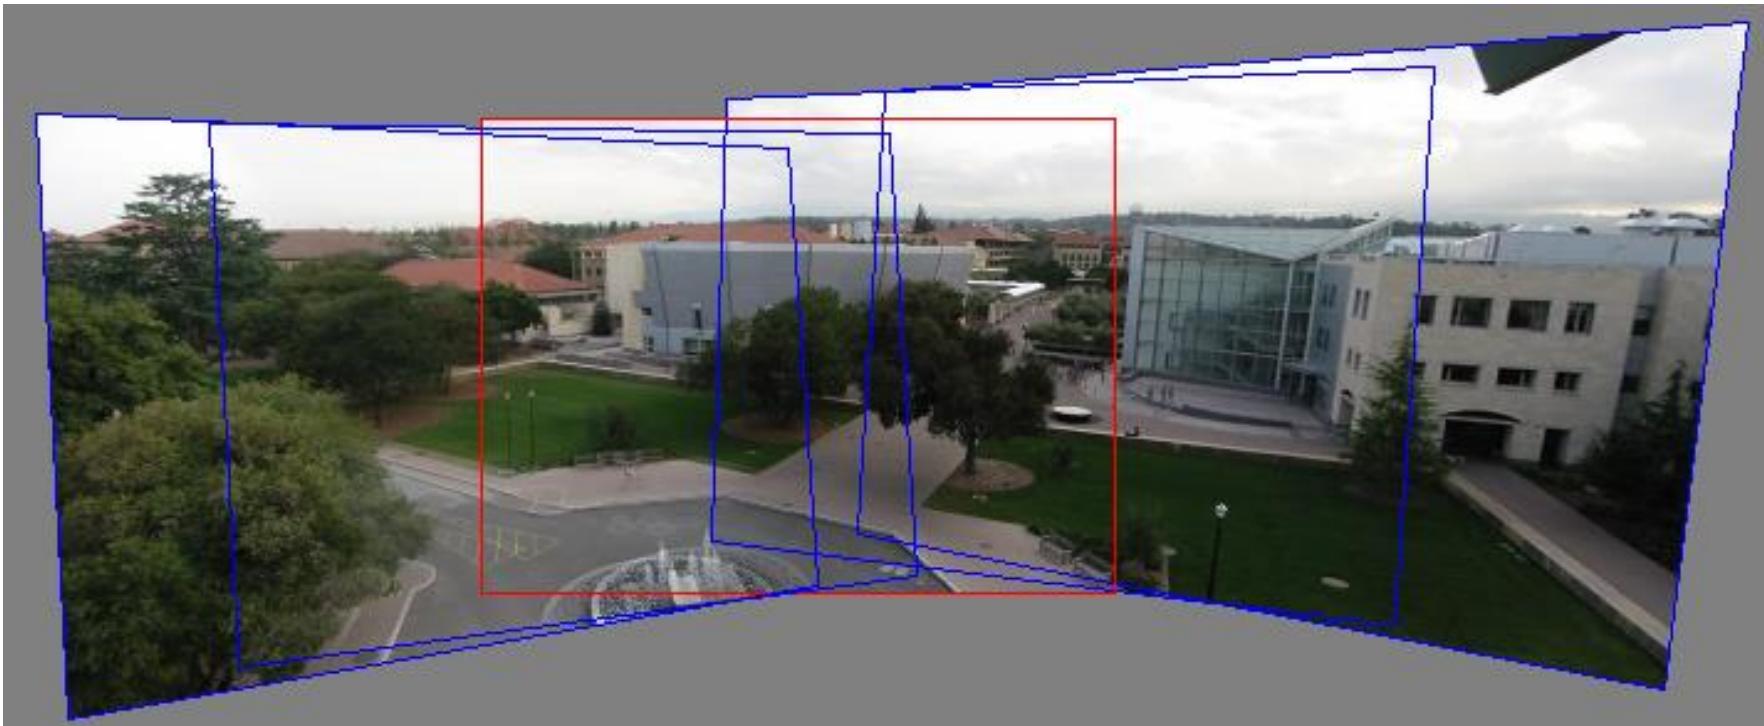


Rollout Photographs © Justin Kerr

<http://research.famsi.org/kerrmaya.html>

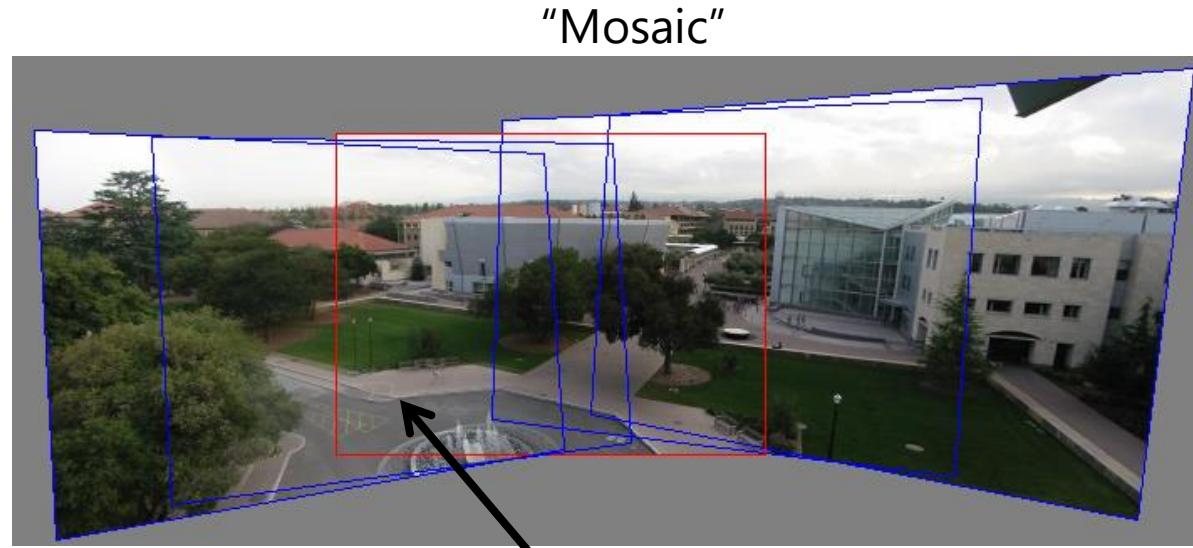
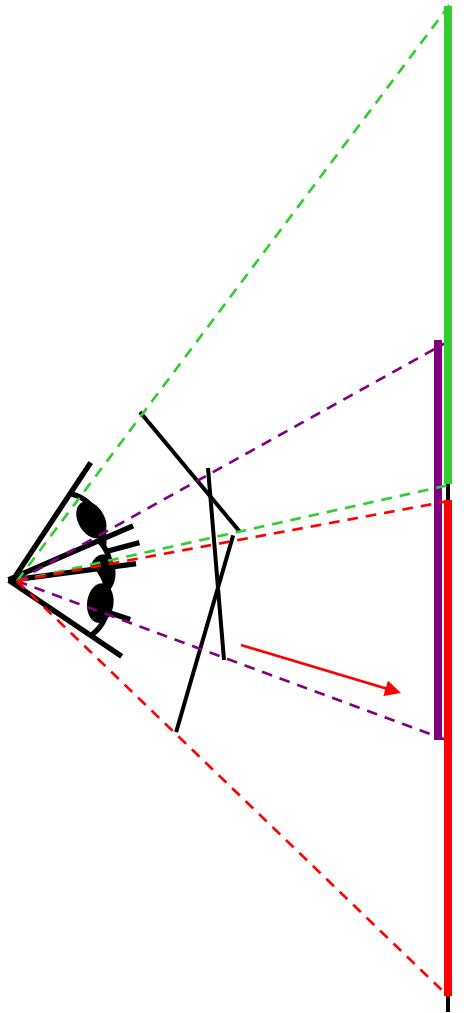
Also known as “cyclographs”, “peripheral images”

Back to panoramas



Can we use homographies to create a 360 degree panorama?

Idea: project images onto a common plane



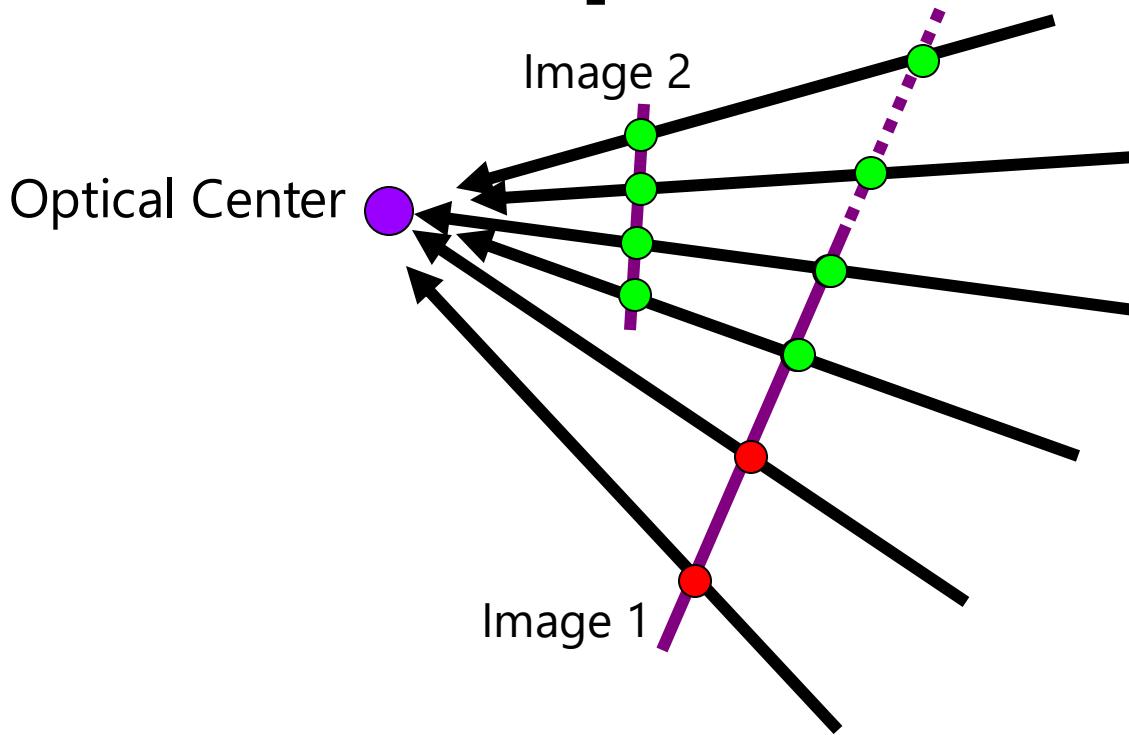
each image is warped
with a homography \mathbf{H}

We'll see what this homography means next
Can't create a 360 panorama this way... we'll fix this shortly

Creating a panorama

- Basic Procedure
 - Take a sequence of images from the same position
 - Rotate the camera about its optical center
 - Compute transformation between second image and first
 - Transform the second image to overlap with the first
 - Blend the two together to create a mosaic
 - If there are more images, repeat

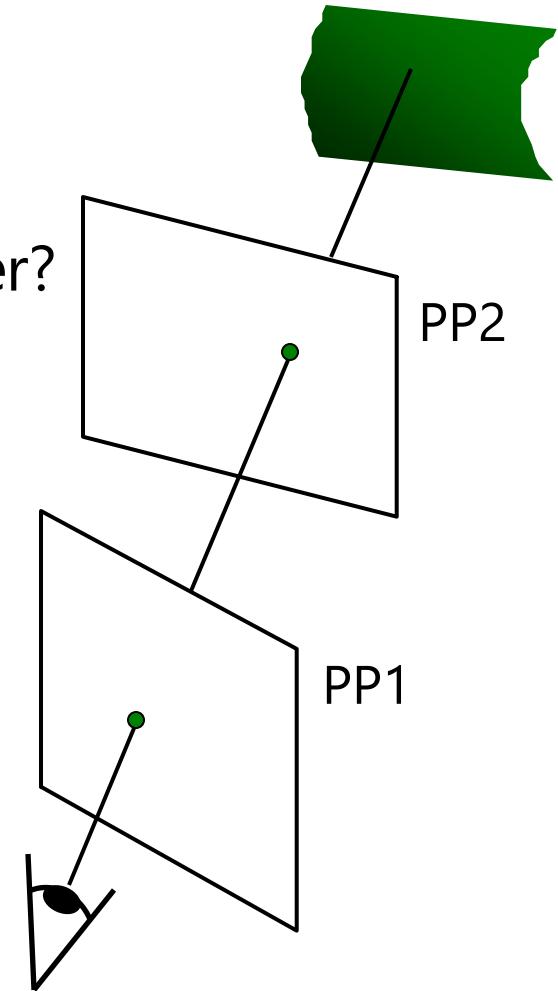
Geometric interpretation of mosaics



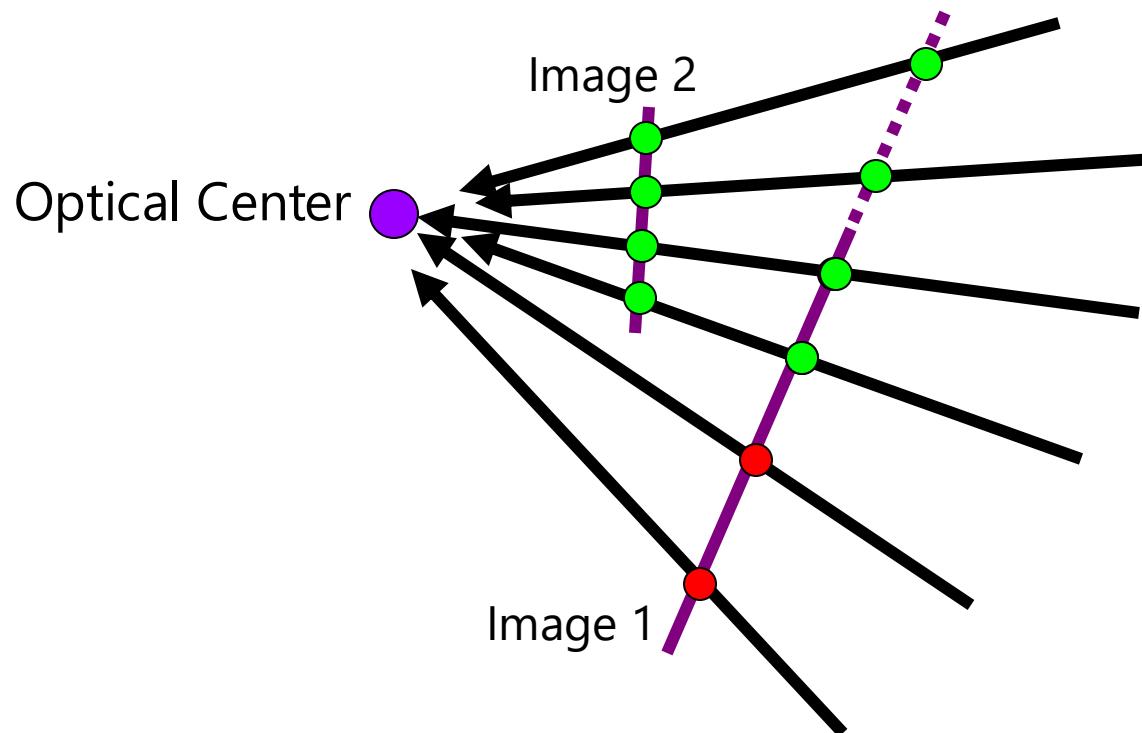
- If we capture all 360° of rays, we can create a 360° panorama
- The basic operation is *projecting* an image from one plane to another
- The projective transformation is **scene-INDEPENDENT**
 - This depends on all the images having the same optical center

Image reprojection

- Basic question
 - How to relate two images from the same camera center?
 - how to map a pixel from PP1 to PP2
- Answer
 - Cast a ray through each pixel in PP1
 - Draw the pixel where that ray intersects PP2



What is the transformation?



How do we map points in image 2 into image 1?

intrinsics	image 1 \mathbf{K}_1
extrinsics (rotation only)	$\mathbf{R}_1 = \mathbf{I}_{3 \times 3}$

image 2 \mathbf{K}_2	
\mathbf{R}_2	

Step 1: Convert pixels in image 2 to rays in camera 2's coordinate system.

$$\begin{bmatrix} X_2 \\ Y_2 \\ Z_2 \end{bmatrix} = \mathbf{K}_2^{-1} \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix}$$

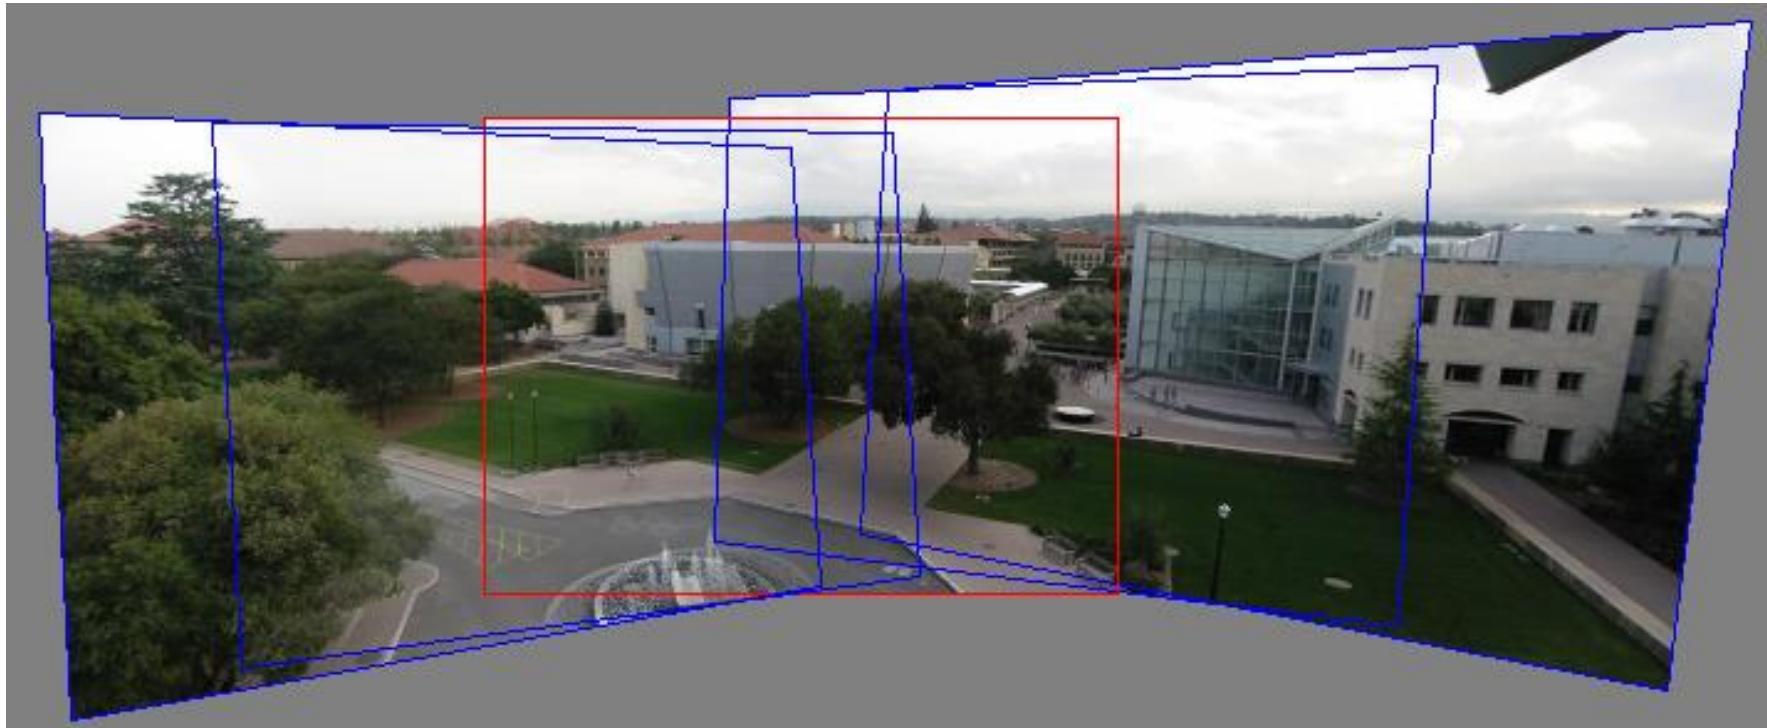
Step 2: Convert rays in camera 2's coordinates to rays in camera 1's coordinates.

$$\begin{bmatrix} X_1 \\ Y_1 \\ Z_1 \end{bmatrix} = \mathbf{R}_2^T \mathbf{K}_2^{-1} \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix}$$

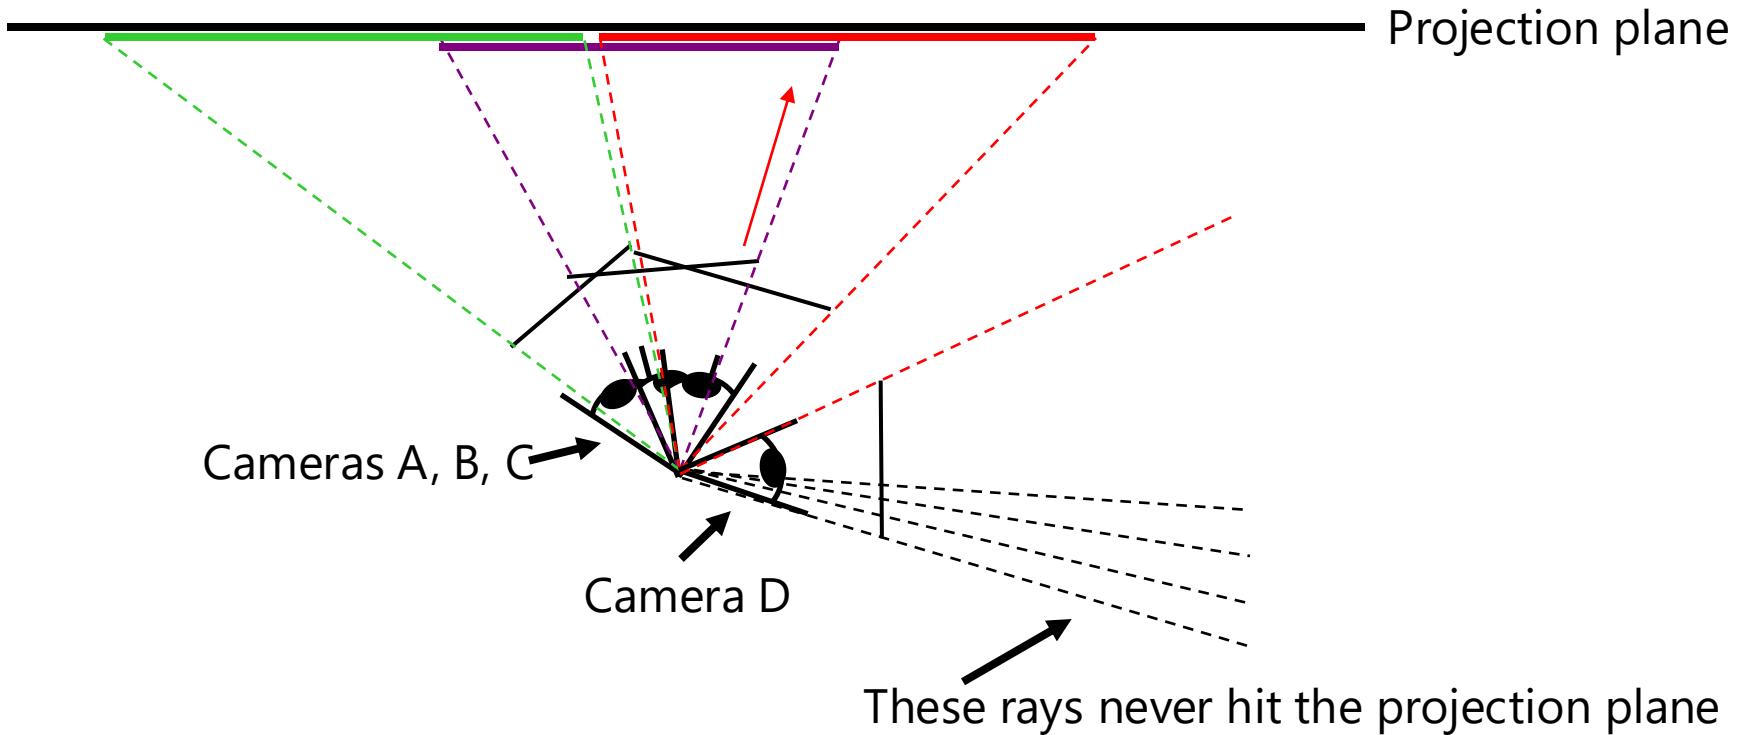
Step 3: Convert rays in camera 1's coordinates to pixels in image 1's coordinates.

$$\begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} \sim \underbrace{\mathbf{K}_1 \mathbf{R}_2^T \mathbf{K}_2^{-1}}_{\text{3x3 homography}} \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix}$$

Can we use homography to create a 360 panorama?

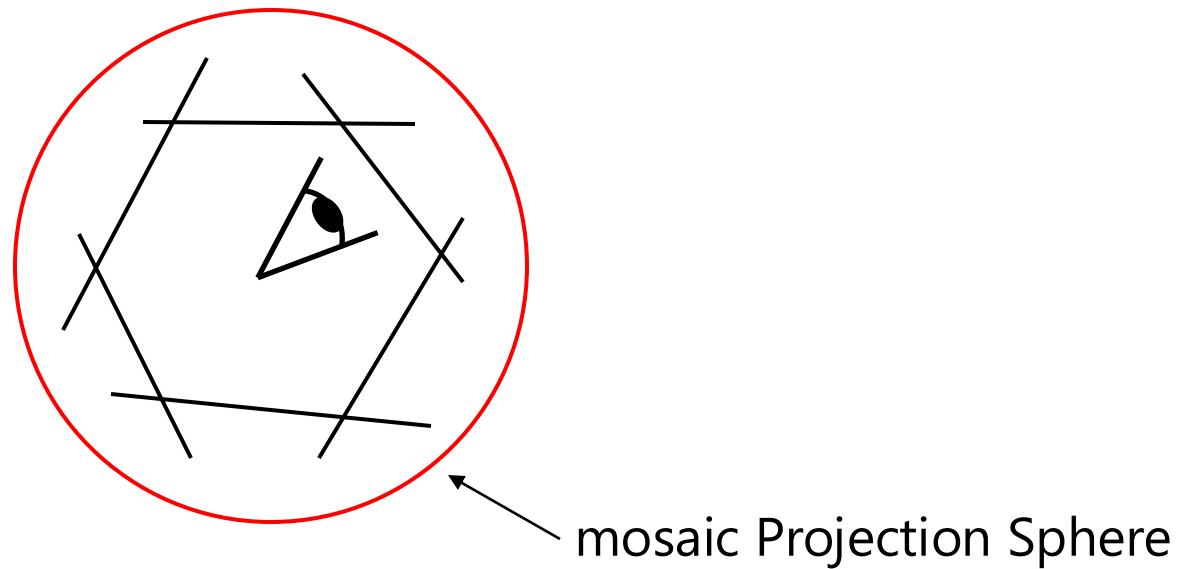


Answer: No

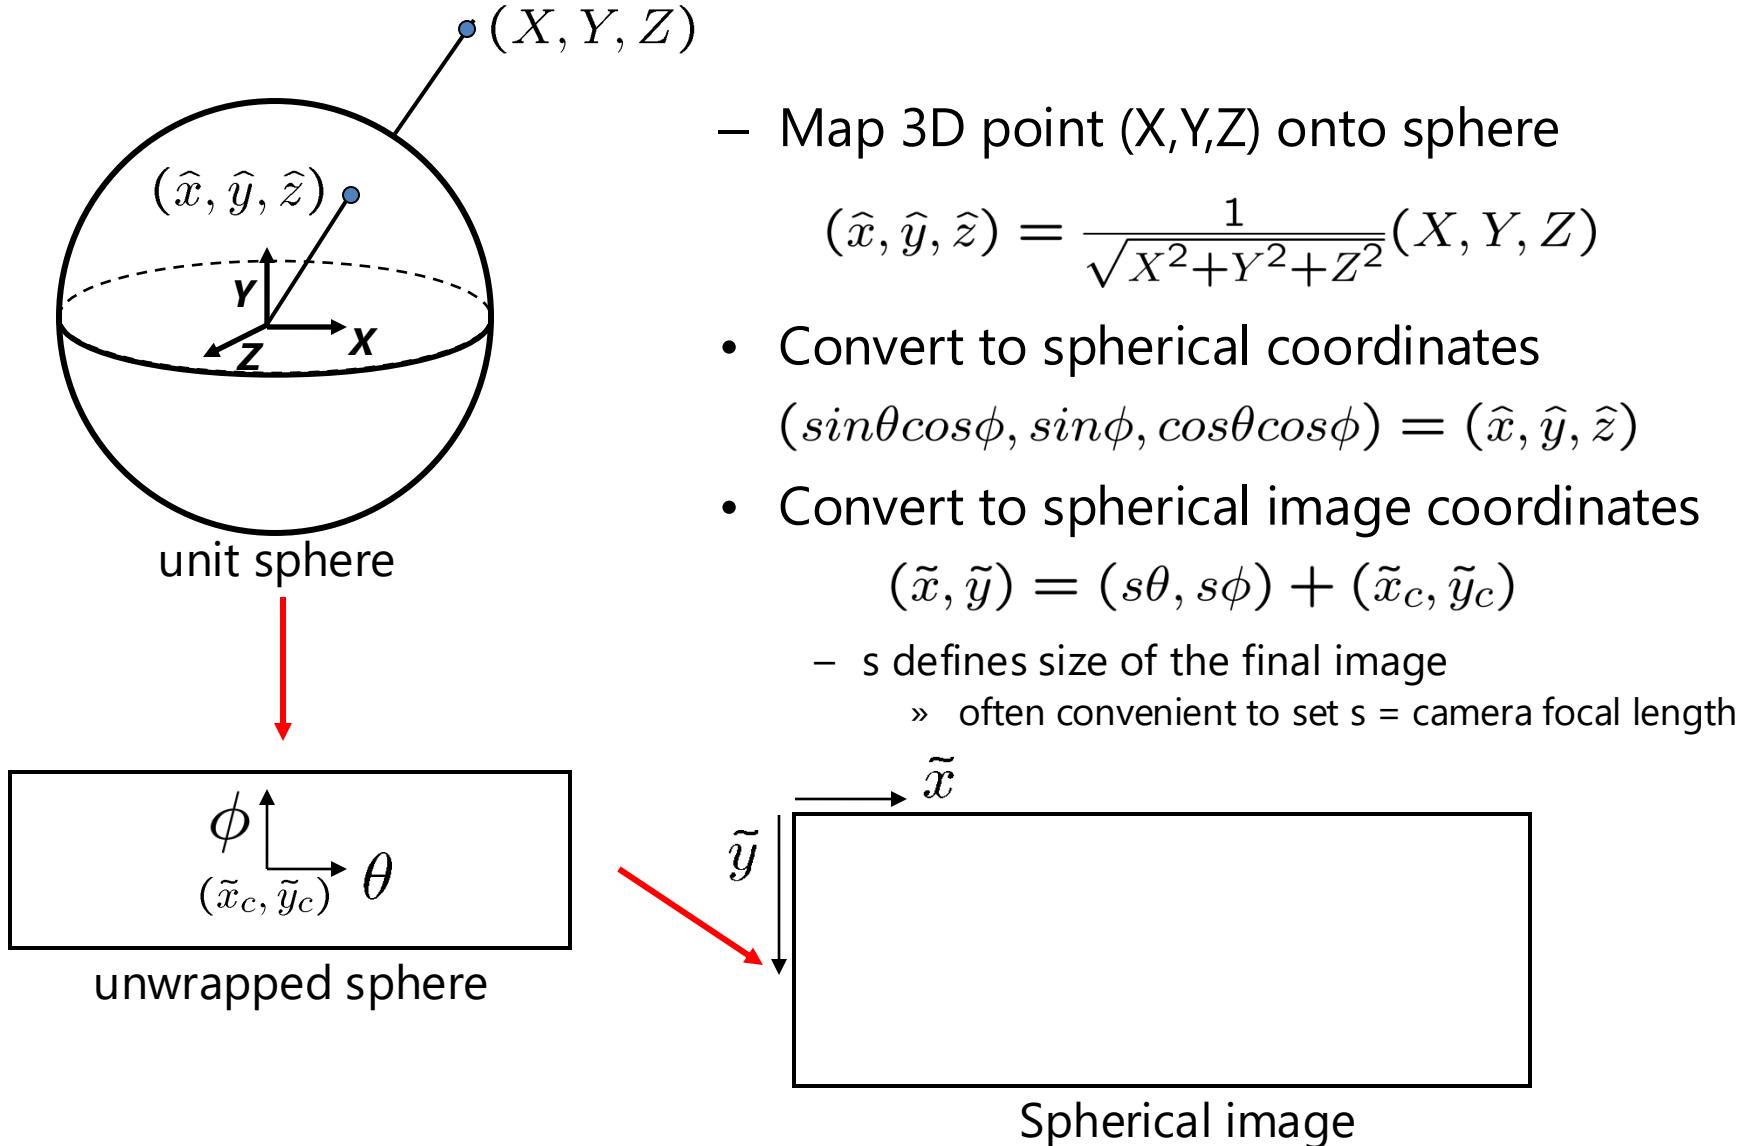


Panoramas

- What if you want a 360° field of view?



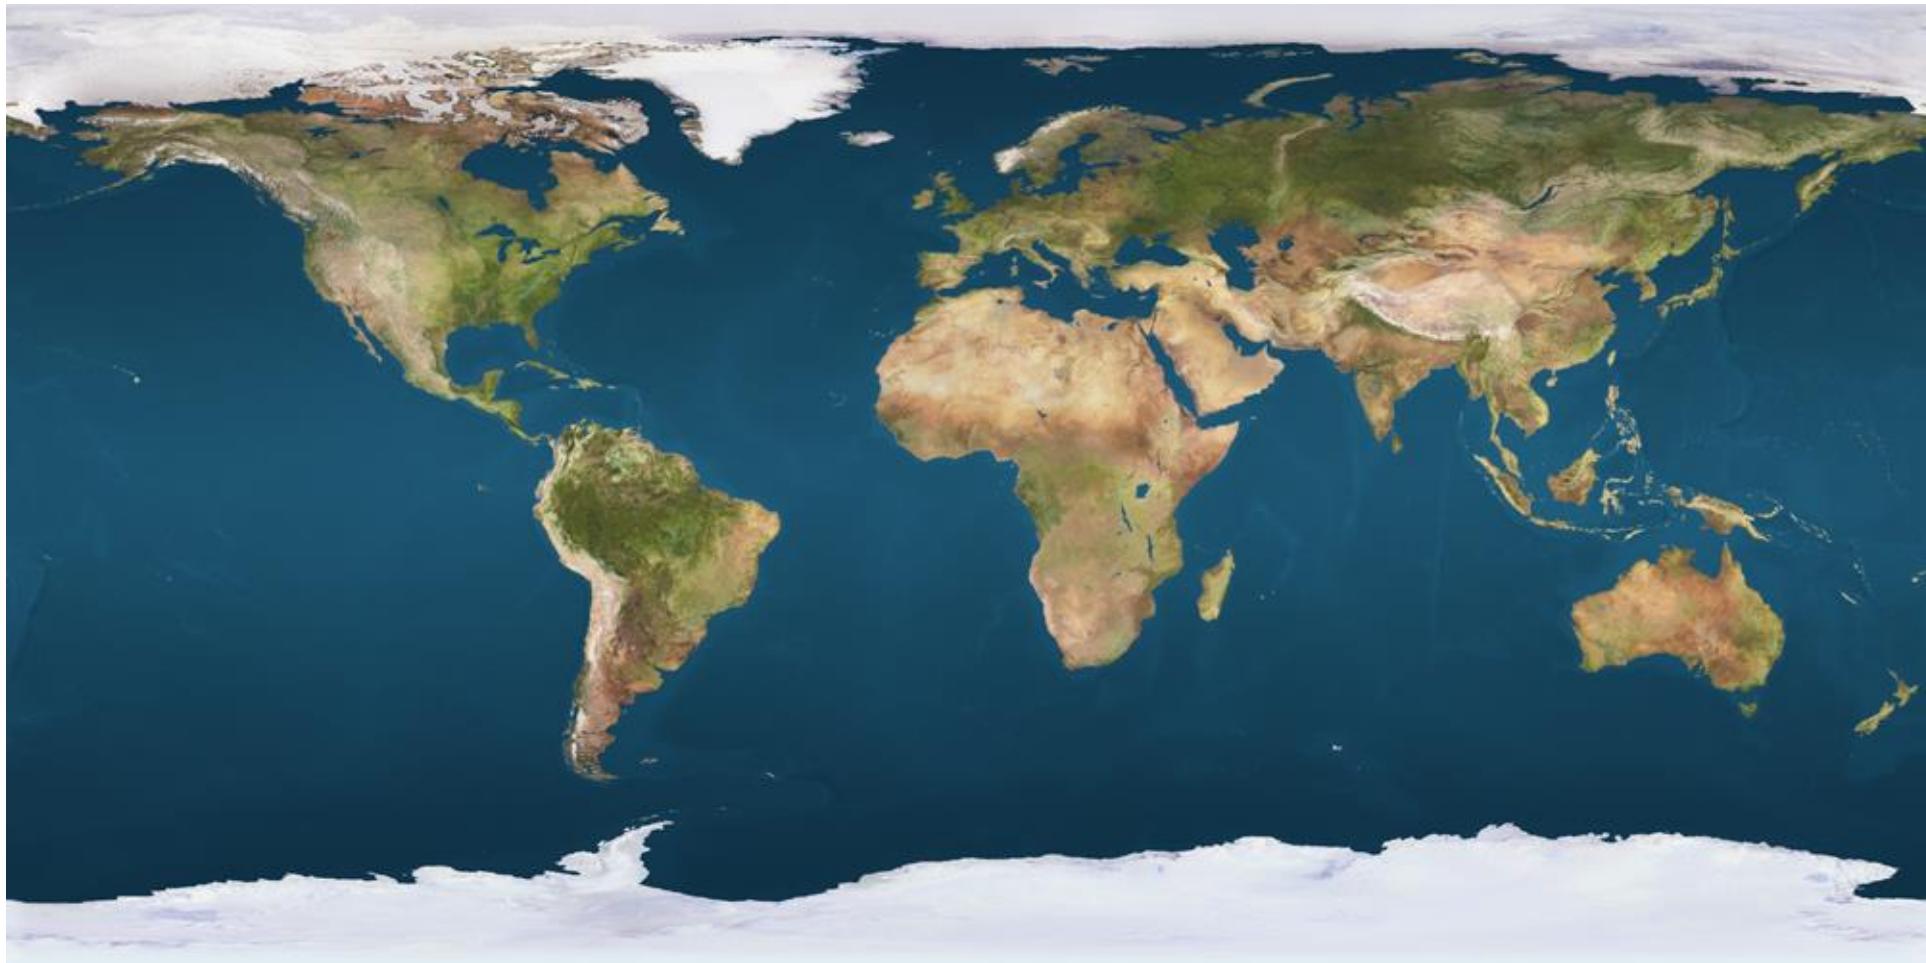
Spherical projection



Unwrapping a sphere



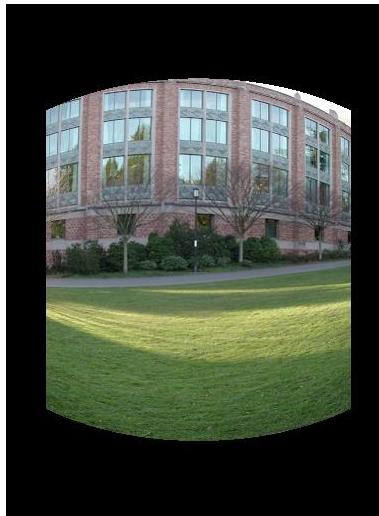
Credit: JHT's Planetary Pixel Emporium



Spherical reprojection



input



$f = 200$ (pixels)



$f = 400$



$f = 800$

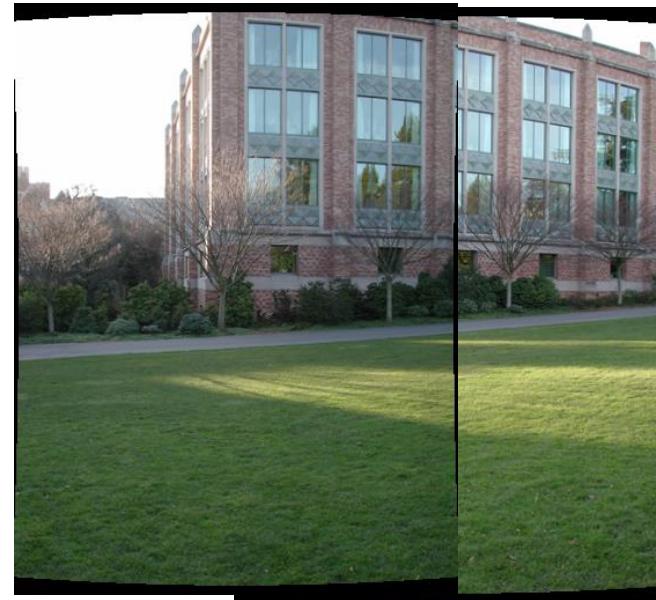
- Map image to spherical coordinates
 - need to know the focal length

Aligning spherical images



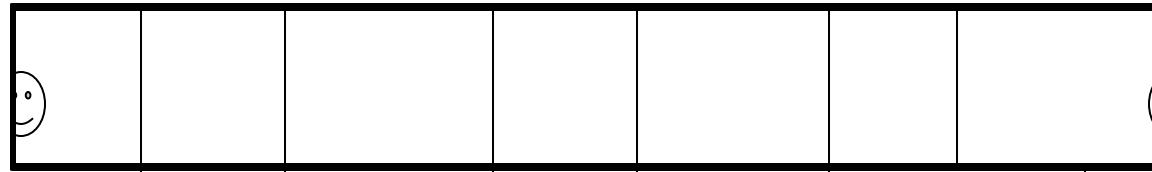
- Suppose we rotate the camera by θ about the vertical axis
 - How does this change the spherical image?

Aligning spherical images



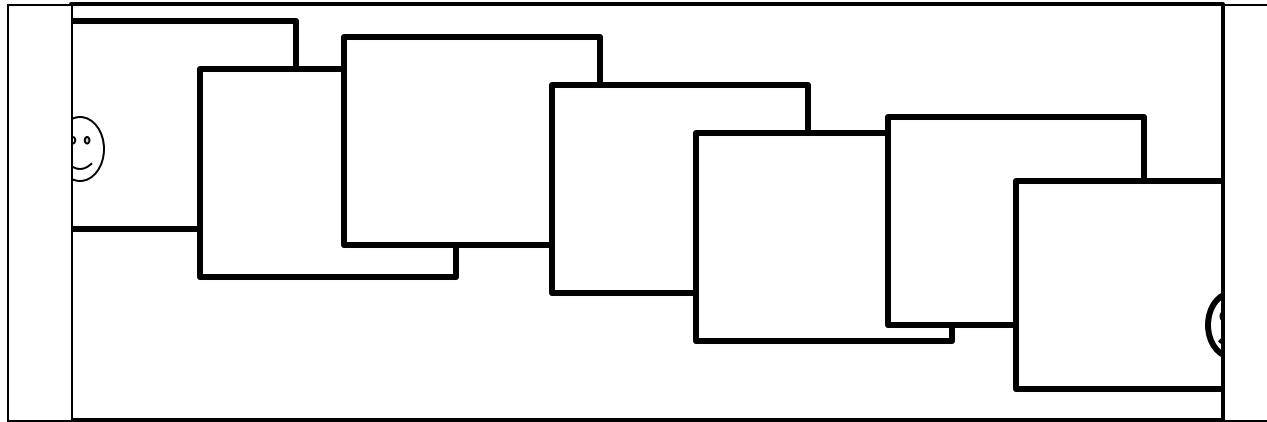
- Suppose we rotate the camera by θ about the vertical axis
 - How does this change the spherical image?
 - Translation by θ
 - This means that we can align spherical images by translation

Assembling the panorama



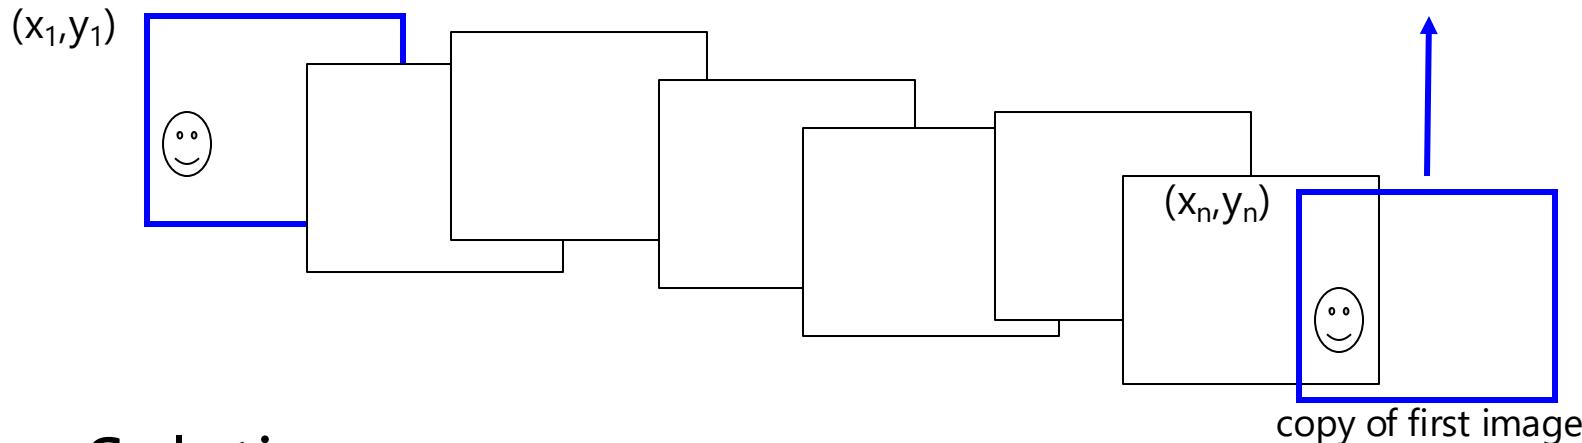
- Stitch pairs together, blend, then crop

Problem: Drift



- Error accumulation
 - small errors accumulate over time

Problem: Drift



- Solution
 - add another copy of first image at the end
 - this gives a constraint: $y_n = y_1$
 - there are a bunch of ways to solve this problem
 - add displacement of $(y_1 - y_n)/(n - 1)$ to each image after the first
 - **apply an affine warp:** $y' = y + ax$ [**you will implement this for P3**]
 - run a big optimization problem, incorporating this constraint
 - best solution, but more complicated
 - known as “bundle adjustment”

Project 3

1. Take pictures on a tripod (or handheld)
 2. Warp to spherical coordinates (not needed if using homographies to align images)
 3. Extract features
 4. Align neighboring pairs using feature matching + RANSAC
 5. Write out list of neighboring translations
 6. Correct for drift
 7. Read in warped images and blend them
 8. Crop the result and import into a viewer
- Roughly based on **Autostitch**
 - By Matthew Brown and David Lowe
 - <http://www.cs.ubc.ca/~mbrown/autostitch/autostitch.html>

Spherical panoramas

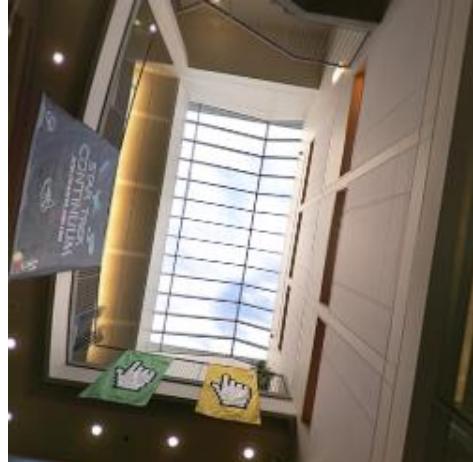


Microsoft Lobby: <http://www.acm.org/pubs/citations/proceedings/graph/258734/p251-szeliski>

Different projections are possible



Cube-map



Blending

- We've aligned the images – now what?

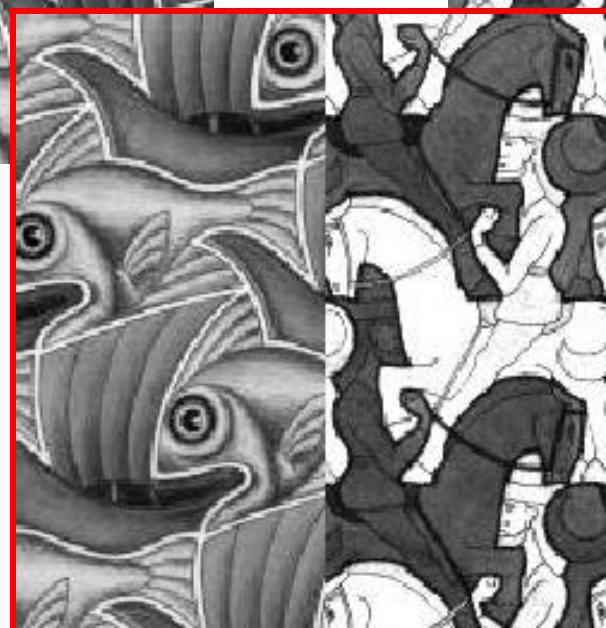
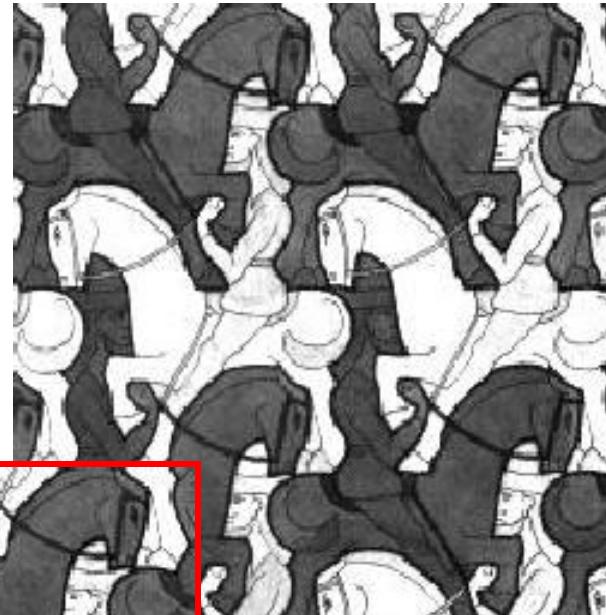
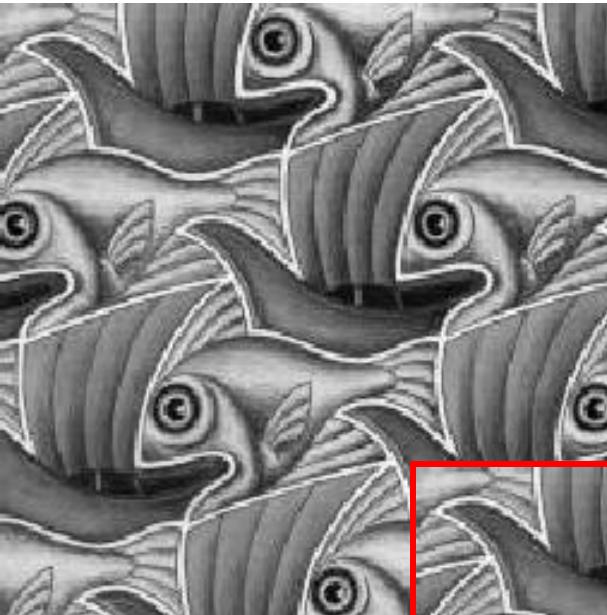


Blending

- Want to seamlessly blend them together



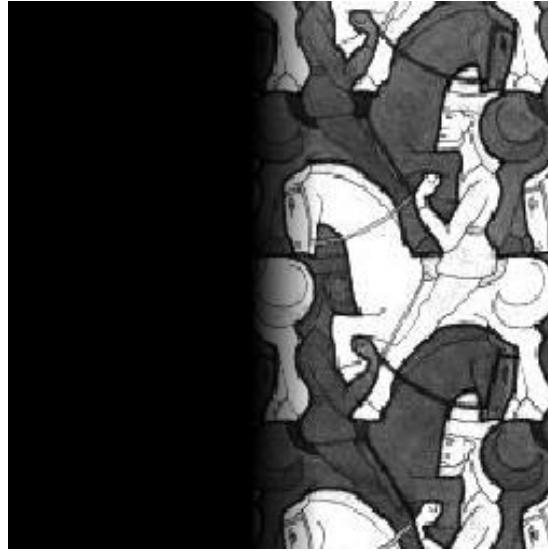
Image Blending



Feathering



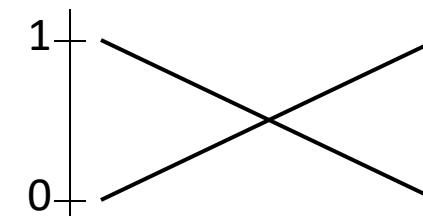
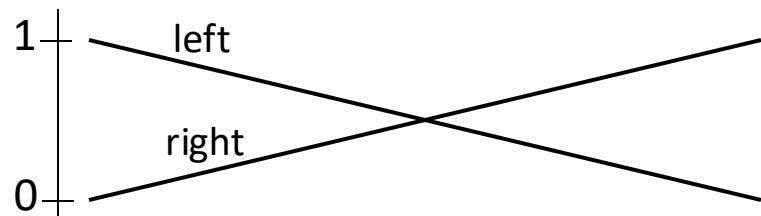
+



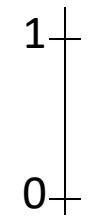
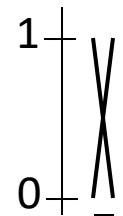
=



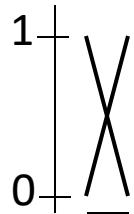
Effect of window size



Effect of window size



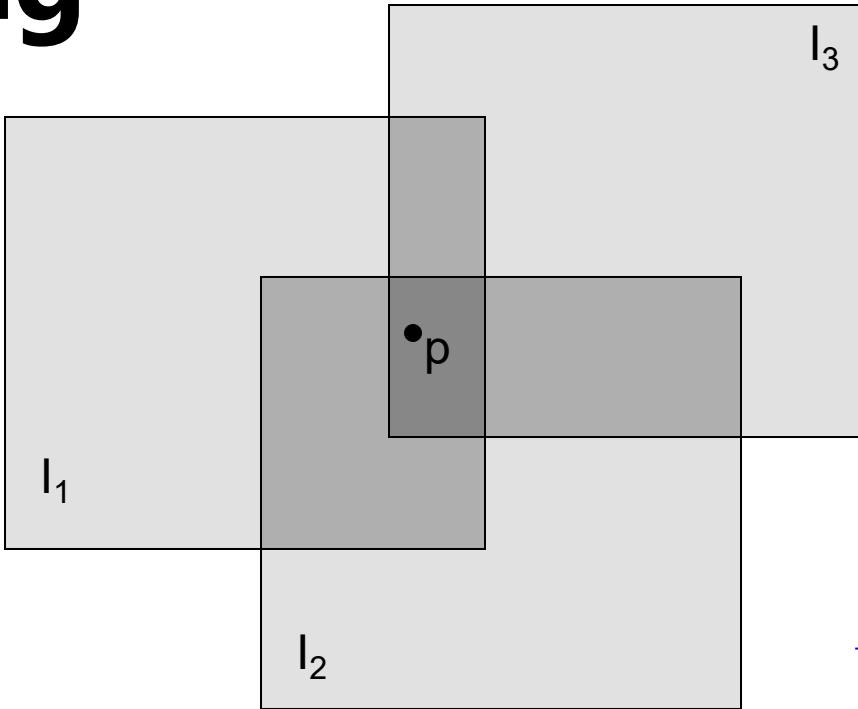
Good window size



“Optimal” window: smooth but not ghosted

- Doesn’t always work...

Alpha Blending



see Blinn (CGA, 1994) for details:

[Compositing, Part 1: Theory](#)

Encoding blend weights: $I(x,y) = (\alpha R, \alpha G, \alpha B, \alpha)$

color at $p = \frac{(\alpha_1 R_1, \alpha_1 G_1, \alpha_1 B_1) + (\alpha_2 R_2, \alpha_2 G_2, \alpha_2 B_2) + (\alpha_3 R_3, \alpha_3 G_3, \alpha_3 B_3)}{\alpha_1 + \alpha_2 + \alpha_3}$

Implement this in two steps:

1. accumulate: add up the (α premultiplied) RGB α values at each pixel
2. normalize: divide each pixel's accumulated RGB by its α value

Q: what if $\alpha = 0$?

Poisson Image Editing



For more info: [Perez et al, SIGGRAPH 2003](#)

Some panorama examples



"Before SIGGRAPH Deadline" Photo credit: Doug Zongker

Some panorama examples

- Every image on Google Streetview



Magic: ghost removal



M. Uyttendaele, A. Eden, and R. Szeliski.
Eliminating ghosting and exposure artifacts in image mosaics.
ICCV 2001

Magic: ghost removal

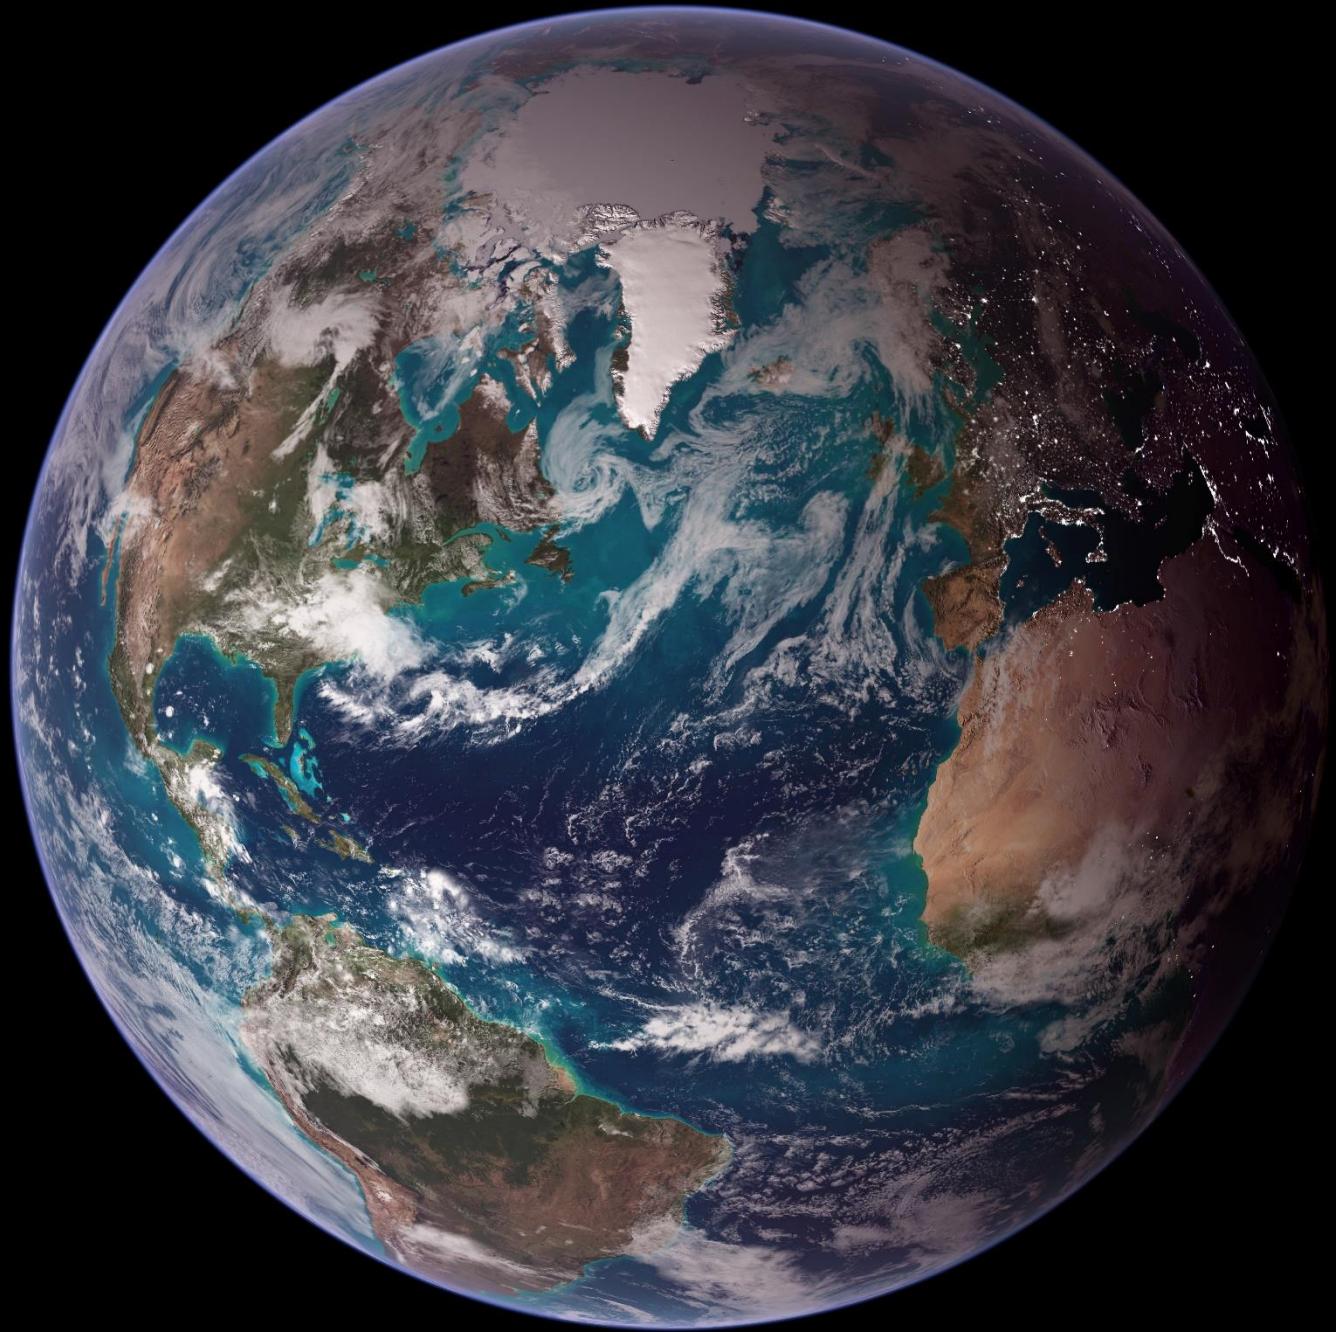


M. Uyttendaele, A. Eden, and R. Szeliski.
Eliminating ghosting and exposure artifacts in image mosaics.
ICCV 2001

Other types of mosaics



- Can mosaic onto *any* surface if you know the geometry
 - See NASA's [Visible Earth project](#) for some stunning earth mosaics





https://www.nasa.gov/centers/wallops/news/frozen_sos.html

Science on a Sphere

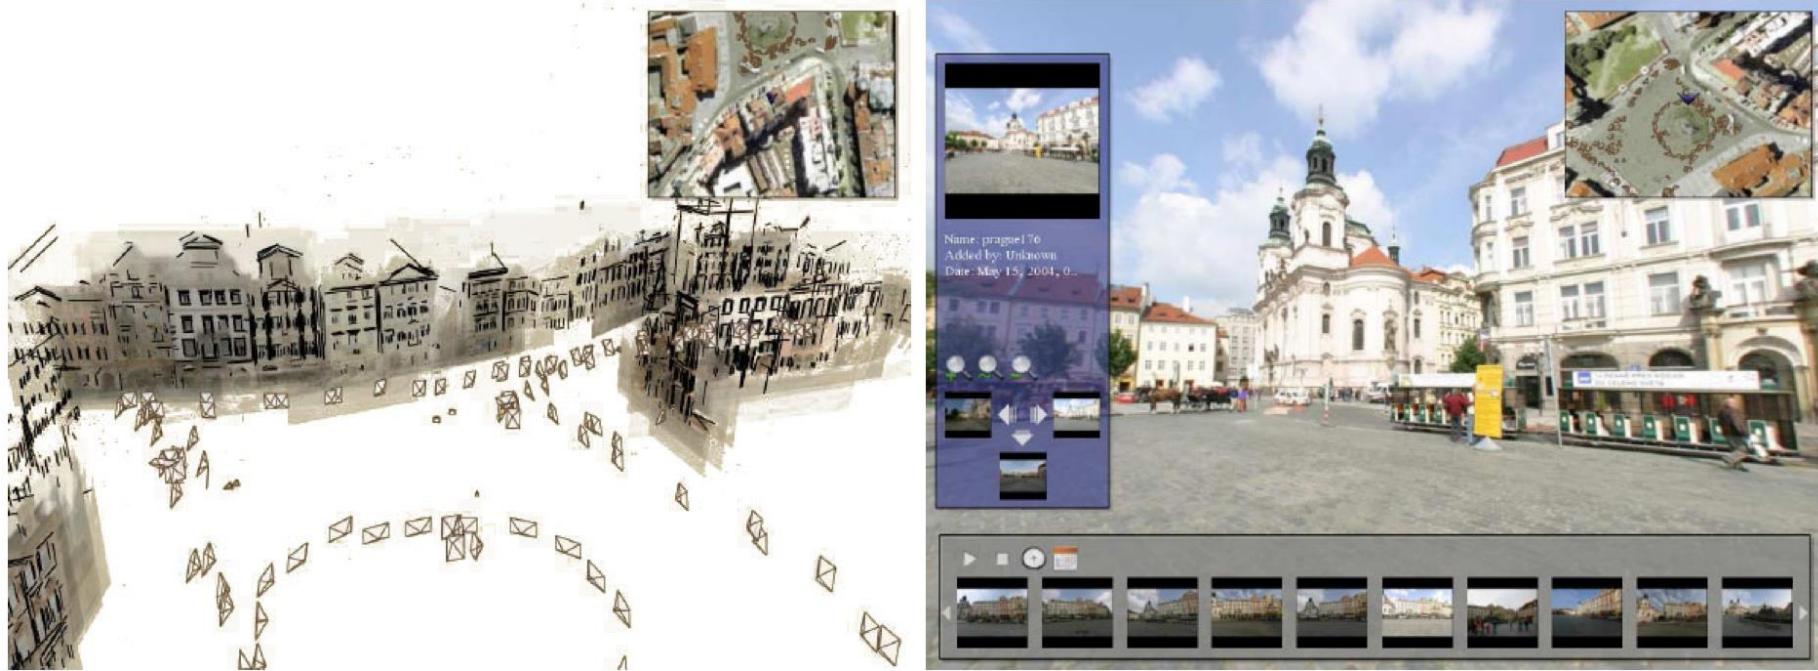


https://news.vcu.edu/article/Science_On_a_Sphere_now_at_VCU_offers_a_world_of_possibilities

Geometric camera calibration

	Structure (scene geometry)	Motion (camera geometry)	Measurements
Camera Calibration (a.k.a. Pose Estimation)	known	estimate	3D to 2D correspondences
Triangulation	estimate	known	2D to 2D coorespondences
Reconstruction	estimate	estimate	2D to 2D coorespondences

Pose Estimation



Given a single image,
estimate the exact position of the photographer

Geometric camera calibration

Given a set of matched points

$$\{\mathbf{X}_i, \mathbf{x}_i\}$$

point in 3D
space point in the
image

and camera model

$$\mathbf{x} = f(\mathbf{X}; \mathbf{p}) = \mathbf{P}\mathbf{X}$$

projection
model parameters Camera
matrix

Find the (pose) estimate of

P

We'll use a **perspective** camera
model for pose estimation

Same setup as homography estimation
(slightly different derivation here)

Mapping between 3D point and image points

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

What are the unknowns?

Mapping between 3D point and image points

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} \text{---} & \mathbf{p}_1^\top & \text{---} \\ \text{---} & \mathbf{p}_2^\top & \text{---} \\ \text{---} & \mathbf{p}_3^\top & \text{---} \end{bmatrix} \begin{bmatrix} | \\ \mathbf{X} \\ | \end{bmatrix}$$

Heterogeneous coordinates

$$x' = \frac{\mathbf{p}_1^\top \mathbf{X}}{\mathbf{p}_3^\top \mathbf{X}} \quad y' = \frac{\mathbf{p}_2^\top \mathbf{X}}{\mathbf{p}_3^\top \mathbf{X}}$$

(non-linear relation between coordinates)

How can we make these relations linear?

How can we make these relations linear?

$$x' = \frac{\mathbf{p}_1^\top \mathbf{X}}{\mathbf{p}_3^\top \mathbf{X}} \quad y' = \frac{\mathbf{p}_2^\top \mathbf{X}}{\mathbf{p}_3^\top \mathbf{X}}$$

Make them linear with algebraic manipulation...

$$\mathbf{p}_2^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} y' = 0$$

$$\mathbf{p}_1^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} x' = 0$$

Now we can setup a system of linear equations with multiple point correspondences

$$\mathbf{p}_2^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} y' = 0$$

$$\mathbf{p}_1^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} x' = 0$$

How do we proceed?

$$\mathbf{p}_2^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} y' = 0$$

$$\mathbf{p}_1^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} x' = 0$$

In matrix form ...

$$\begin{bmatrix} \mathbf{X}^\top & \mathbf{0} & -x' \mathbf{X}^\top \\ \mathbf{0} & \mathbf{X}^\top & -y' \mathbf{X}^\top \end{bmatrix} \begin{bmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{bmatrix} = \mathbf{0}$$

How do we proceed?

$$\mathbf{p}_2^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} y' = 0$$

$$\mathbf{p}_1^\top \mathbf{X} - \mathbf{p}_3^\top \mathbf{X} x' = 0$$

In matrix form ...

$$\begin{bmatrix} \mathbf{X}^\top & \mathbf{0} & -x' \mathbf{X}^\top \\ \mathbf{0} & \mathbf{X}^\top & -y' \mathbf{X}^\top \end{bmatrix} \begin{bmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{bmatrix} = \mathbf{0}$$

For N points ...

$$\begin{bmatrix} \mathbf{X}_1^\top & \mathbf{0} & -x' \mathbf{X}_1^\top \\ \mathbf{0} & \mathbf{X}_1^\top & -y' \mathbf{X}_1^\top \\ \vdots & \vdots & \vdots \\ \mathbf{X}_N^\top & \mathbf{0} & -x' \mathbf{X}_N^\top \\ \mathbf{0} & \mathbf{X}_N^\top & -y' \mathbf{X}_N^\top \end{bmatrix} \begin{bmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{bmatrix} = \mathbf{0}$$

*How do we solve
this system?*

Solve for camera matrix by

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \|\mathbf{A}\mathbf{x}\|^2 \text{ subject to } \|\mathbf{x}\|^2 = 1$$

$$\mathbf{A} = \begin{bmatrix} \mathbf{X}_1^\top & \mathbf{0} & -\mathbf{x}' \mathbf{X}_1^\top \\ \mathbf{0} & \mathbf{X}_1^\top & -\mathbf{y}' \mathbf{X}_1^\top \\ \vdots & \vdots & \vdots \\ \mathbf{X}_N^\top & \mathbf{0} & -\mathbf{x}' \mathbf{X}_N^\top \\ \mathbf{0} & \mathbf{X}_N^\top & -\mathbf{y}' \mathbf{X}_N^\top \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{bmatrix}$$

SVD!

Solve for camera matrix by

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \|\mathbf{A}\mathbf{x}\|^2 \text{ subject to } \|\mathbf{x}\|^2 = 1$$

$$\mathbf{A} = \begin{bmatrix} \mathbf{X}_1^\top & \mathbf{0} & -\mathbf{x}' \mathbf{X}_1^\top \\ \mathbf{0} & \mathbf{X}_1^\top & -\mathbf{y}' \mathbf{X}_1^\top \\ \vdots & \vdots & \vdots \\ \mathbf{X}_N^\top & \mathbf{0} & -\mathbf{x}' \mathbf{X}_N^\top \\ \mathbf{0} & \mathbf{X}_N^\top & -\mathbf{y}' \mathbf{X}_N^\top \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{bmatrix}$$

Solution \mathbf{x} is the column of \mathbf{V}
corresponding to smallest singular
value of

$$\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^\top$$

Solve for camera matrix by

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \|\mathbf{A}\mathbf{x}\|^2 \text{ subject to } \|\mathbf{x}\|^2 = 1$$

$$\mathbf{A} = \begin{bmatrix} \mathbf{X}_1^\top & \mathbf{0} & -\mathbf{x}' \mathbf{X}_1^\top \\ \mathbf{0} & \mathbf{X}_1^\top & -\mathbf{y}' \mathbf{X}_1^\top \\ \vdots & \vdots & \vdots \\ \mathbf{X}_N^\top & \mathbf{0} & -\mathbf{x}' \mathbf{X}_N^\top \\ \mathbf{0} & \mathbf{X}_N^\top & -\mathbf{y}' \mathbf{X}_N^\top \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} \mathbf{p}_1 \\ \mathbf{p}_2 \\ \mathbf{p}_3 \end{bmatrix}$$

Equivalently, solution \mathbf{x} is the
Eigenvector corresponding to
smallest Eigenvalue of

$$\mathbf{A}^\top \mathbf{A}$$

Now we have:

$$\mathbf{P} = \begin{bmatrix} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{bmatrix}$$

Are we done?

Almost there ...

$$\mathbf{P} = \begin{bmatrix} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{bmatrix}$$

How do you get the intrinsic and extrinsic parameters from the projection matrix?

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\mathbf{P} = \mathbf{K}[\mathbf{R}|\mathbf{t}]$$

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\begin{aligned}\mathbf{P} &= \mathbf{K}[\mathbf{R}|\mathbf{t}] \\ &= \mathbf{K}[\mathbf{R}] - \mathbf{R}\mathbf{c} \\ &= [\mathbf{M}] - \mathbf{M}\mathbf{c}\end{aligned}$$

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\begin{aligned}\mathbf{P} &= \mathbf{K}[\mathbf{R}|\mathbf{t}] \\ &= \mathbf{K}[\mathbf{R}] - \mathbf{R}\mathbf{c} \\ &= [\mathbf{M}] - \mathbf{M}\mathbf{c}\end{aligned}$$

Find the camera center **C**

What is the projection of the camera center?

Find intrinsic **K** and rotation **R**

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\begin{aligned}\mathbf{P} &= \mathbf{K}[\mathbf{R}|\mathbf{t}] \\ &= \mathbf{K}[\mathbf{R}] - \mathbf{R}\mathbf{c} \\ &= [\mathbf{M}] - \mathbf{M}\mathbf{c}\end{aligned}$$

Find the camera center \mathbf{c}

$$\mathbf{P}\mathbf{c} = \mathbf{0}$$

How do we compute the camera center from this?

Find intrinsic \mathbf{K} and rotation \mathbf{R}

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\begin{aligned}\mathbf{P} &= \mathbf{K}[\mathbf{R}|\mathbf{t}] \\ &= \mathbf{K}[\mathbf{R}] - \mathbf{R}\mathbf{c} \\ &= [\mathbf{M}] - \mathbf{M}\mathbf{c}\end{aligned}$$

Find the camera center \mathbf{c}

$$\mathbf{P}\mathbf{c} = \mathbf{0}$$

SVD of P!

c is the singular vector corresponding
to the smallest singular value

Find intrinsic \mathbf{K} and rotation \mathbf{R}

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\begin{aligned}\mathbf{P} &= \mathbf{K}[\mathbf{R}|\mathbf{t}] \\ &= \mathbf{K}[\mathbf{R}] - \mathbf{R}\mathbf{c} \\ &= [\mathbf{M}] - \mathbf{M}\mathbf{c}\end{aligned}$$

Find the camera center \mathbf{c}

$$\mathbf{P}\mathbf{c} = \mathbf{0}$$

SVD of \mathbf{P} !

c is the singular vector corresponding
to the smallest singular value

Find intrinsic \mathbf{K} and rotation \mathbf{R}

$$\mathbf{M} = \mathbf{K}\mathbf{R}$$

*Any useful properties of K
and R we can use?*

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\begin{aligned}\mathbf{P} &= \mathbf{K}[\mathbf{R}|\mathbf{t}] \\ &= \mathbf{K}[\mathbf{R}] - \mathbf{R}\mathbf{c} \\ &= [\mathbf{M}] - \mathbf{M}\mathbf{c}\end{aligned}$$

Find the camera center \mathbf{c}

$$\mathbf{P}\mathbf{c} = \mathbf{0}$$

SVD of \mathbf{P} !

c is the singular vector corresponding
to the smallest singular value

Find intrinsic \mathbf{K} and rotation \mathbf{R}

$$\mathbf{M} = \mathbf{K}\mathbf{R}$$

right upper triangle orthogonal

*How do we find K
and R?*

Decomposition of the Camera Matrix

$$\mathbf{P} = \left[\begin{array}{ccc|c} p_1 & p_2 & p_3 & p_4 \\ p_5 & p_6 & p_7 & p_8 \\ p_9 & p_{10} & p_{11} & p_{12} \end{array} \right]$$

$$\begin{aligned}\mathbf{P} &= \mathbf{K}[\mathbf{R}|\mathbf{t}] \\ &= \mathbf{K}[\mathbf{R}] - \mathbf{R}\mathbf{c} \\ &= [\mathbf{M}] - \mathbf{M}\mathbf{c}\end{aligned}$$

Find the camera center \mathbf{c}

$$\mathbf{P}\mathbf{c} = \mathbf{0}$$

SVD of \mathbf{P} !

c is the singular vector corresponding
to the smallest singular value

Find intrinsic \mathbf{K} and rotation \mathbf{R}

$$\mathbf{M} = \mathbf{K}\mathbf{R}$$

QR decomposition

Geometric camera calibration

Given a set of matched points

$$\{\mathbf{X}_i, \mathbf{x}_i\}$$

point in 3D
space point in the
image

*Where do we get these
matched points from?*

and camera model

$$\mathbf{x} = f(\mathbf{X}; \mathbf{p}) = \mathbf{P}\mathbf{X}$$

projection
model parameters Camera
matrix

Find the (pose) estimate of

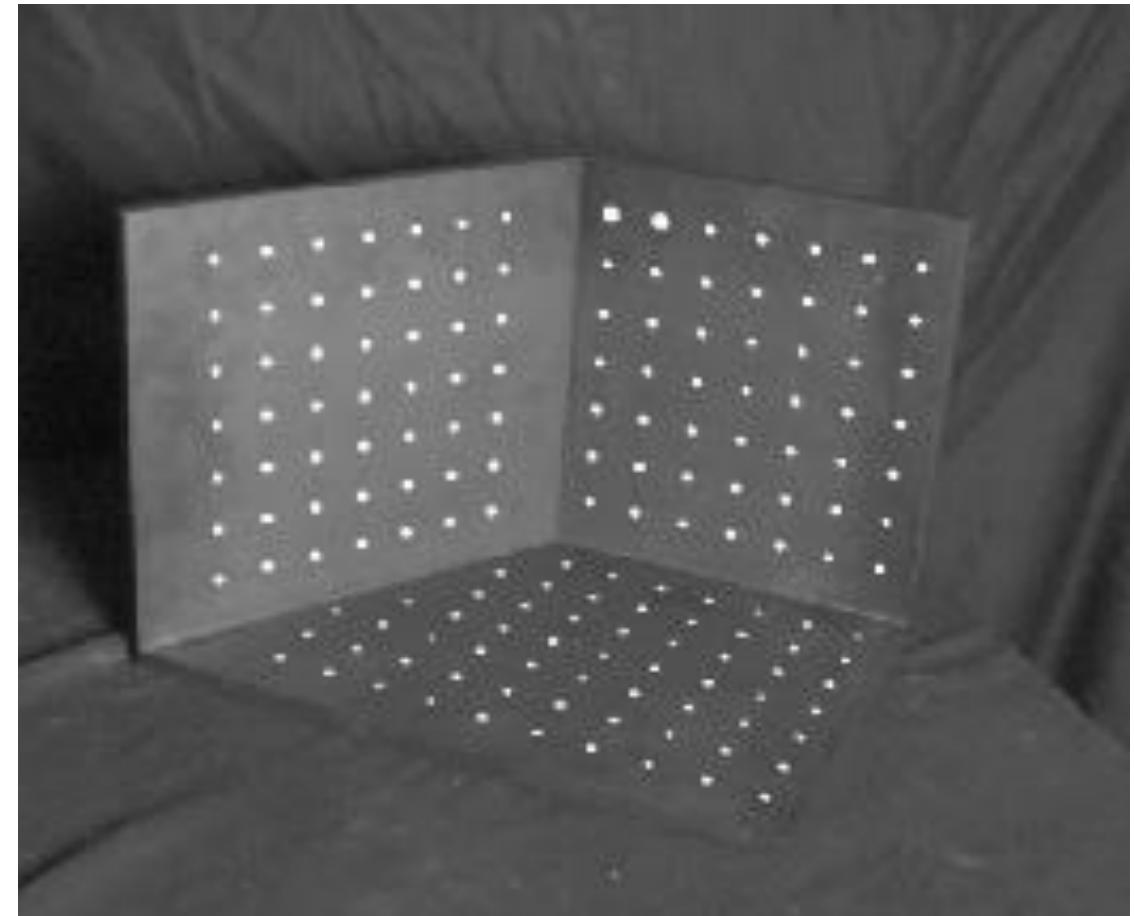
P

We'll use a **perspective** camera
model for pose estimation

Calibration using a reference object

Place a known object in the scene:

- identify correspondences between image and scene
- compute mapping from scene to image



Issues:

- must know geometry very accurately
- must know 3D->2D correspondence

Geometric camera calibration

Advantages:

- Very simple to formulate.
- Analytical solution.

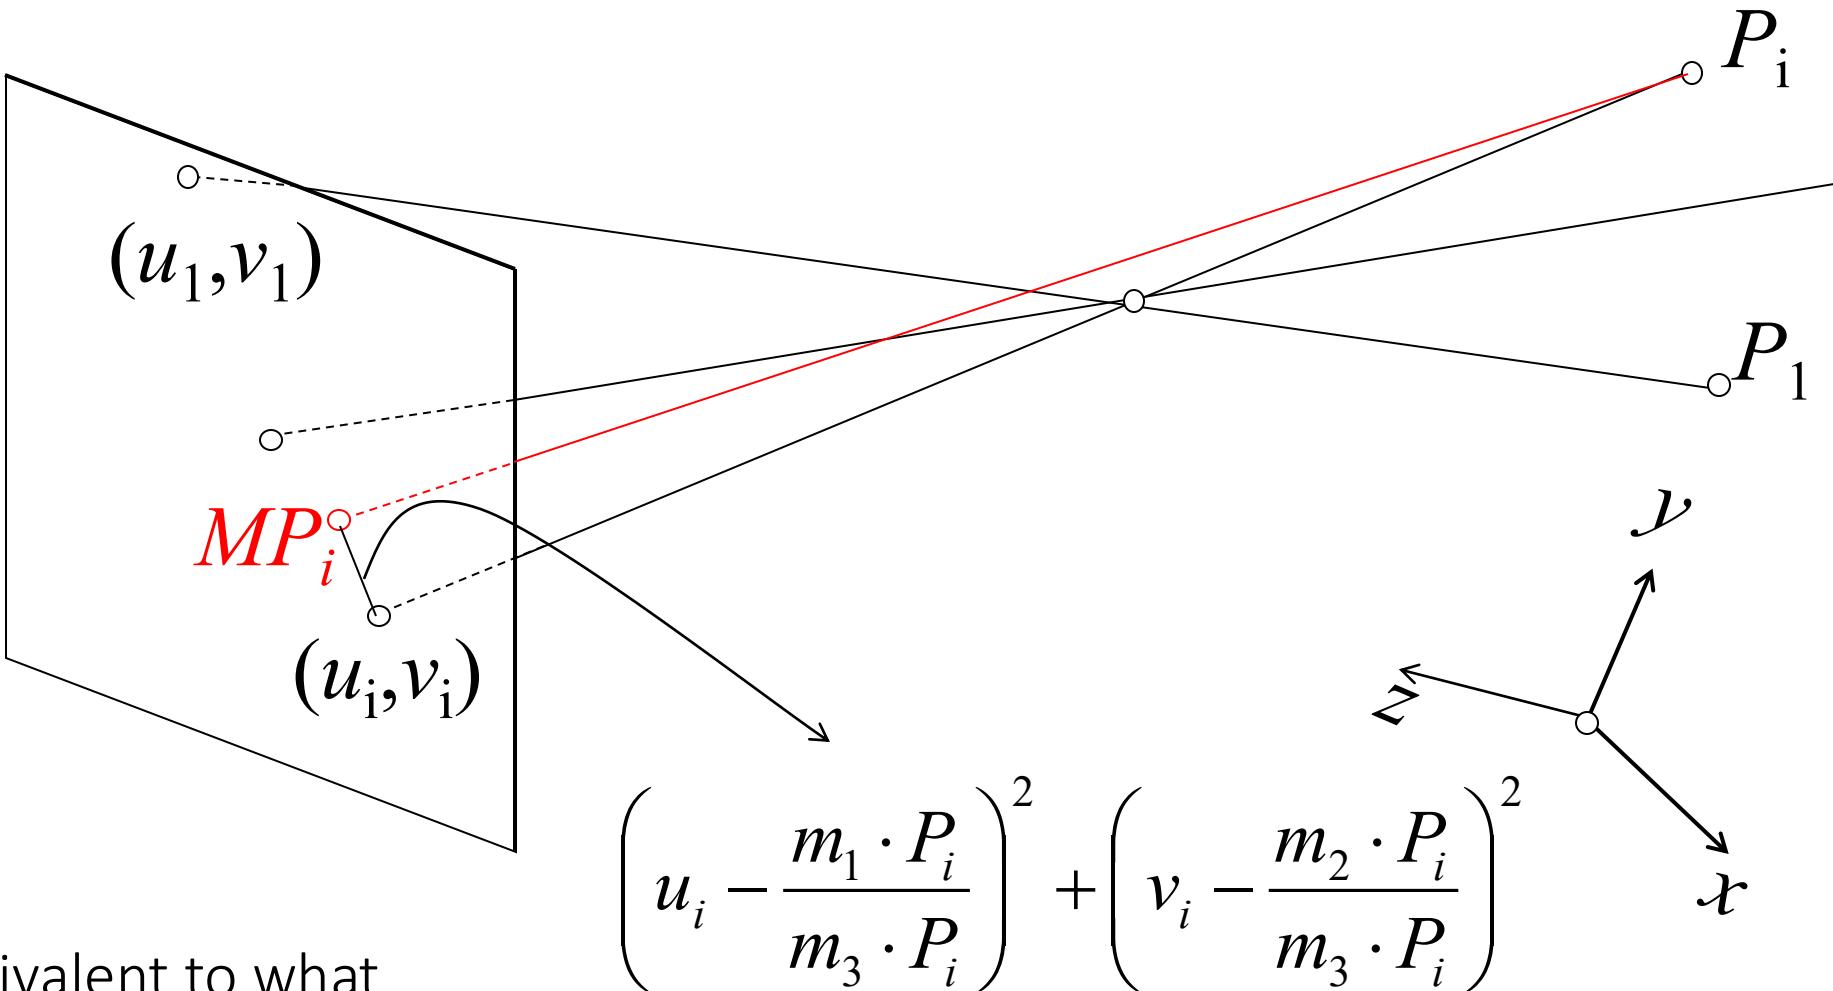
Disadvantages:

- Doesn't model radial distortion.
- Hard to impose constraints (e.g., known f).
- Doesn't minimize the correct error function.

For these reasons, *nonlinear methods* are preferred

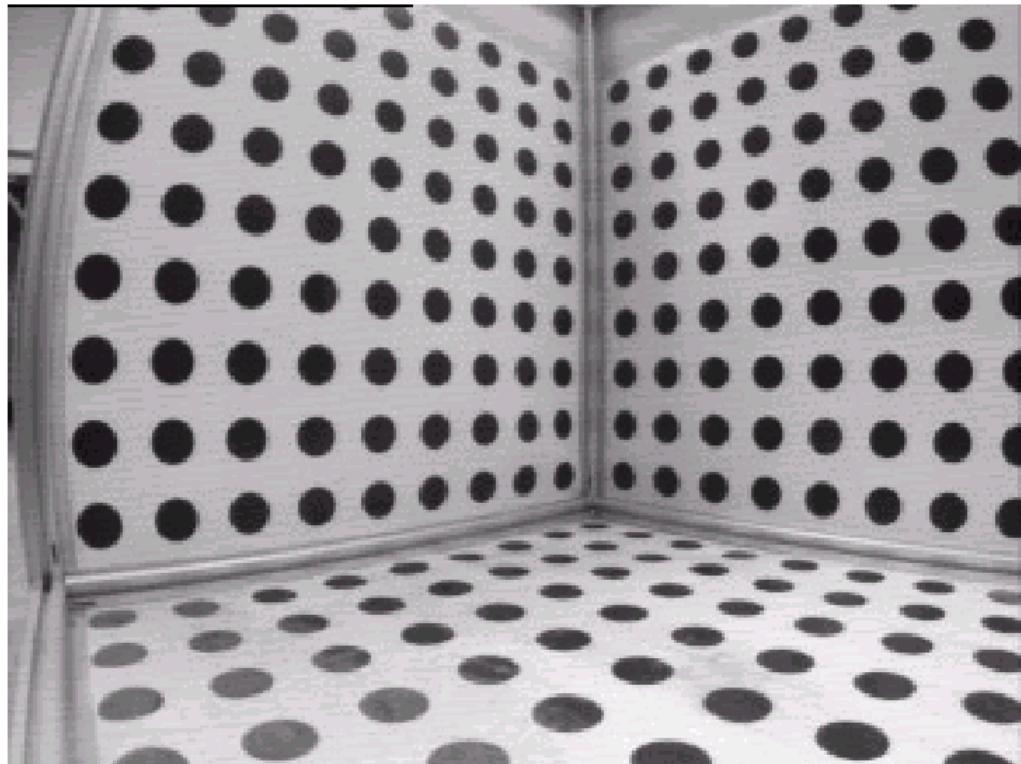
- Define error function E between projected 3D points and image positions
 - E is nonlinear function of intrinsics, extrinsics, radial distortion
- Minimize E using nonlinear optimization techniques

Minimizing reprojection error

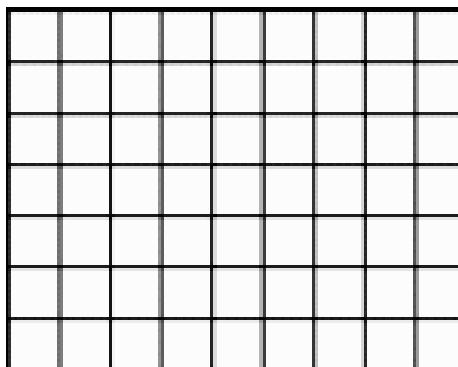


Is this equivalent to what
we were doing previously?

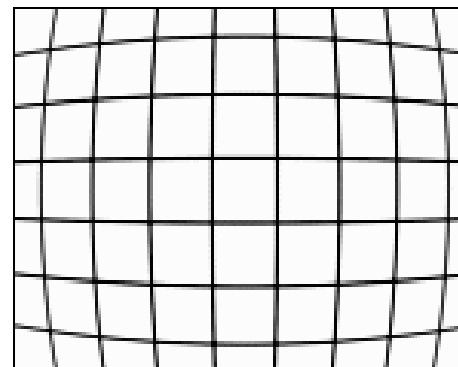
Radial distortion



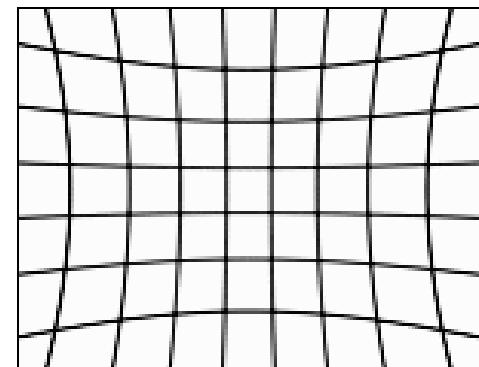
What causes this distortion?



no distortion

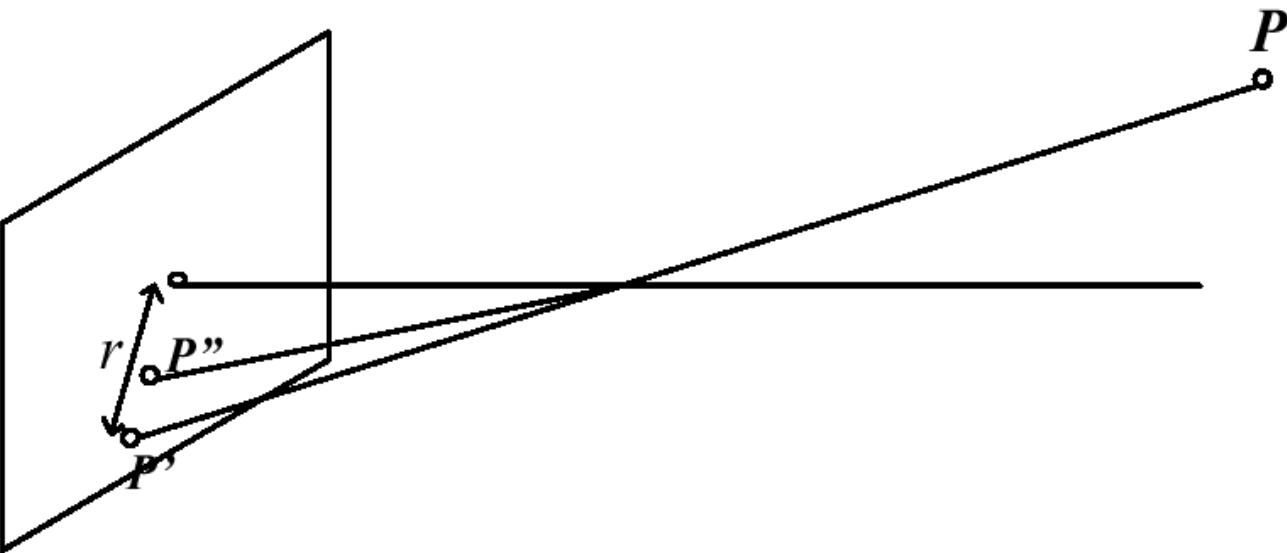


barrel distortion



pincushion distortion

Radial distortion model



Ideal:

$$x' = f \frac{x}{z}$$

$$y' = f \frac{y}{z}$$

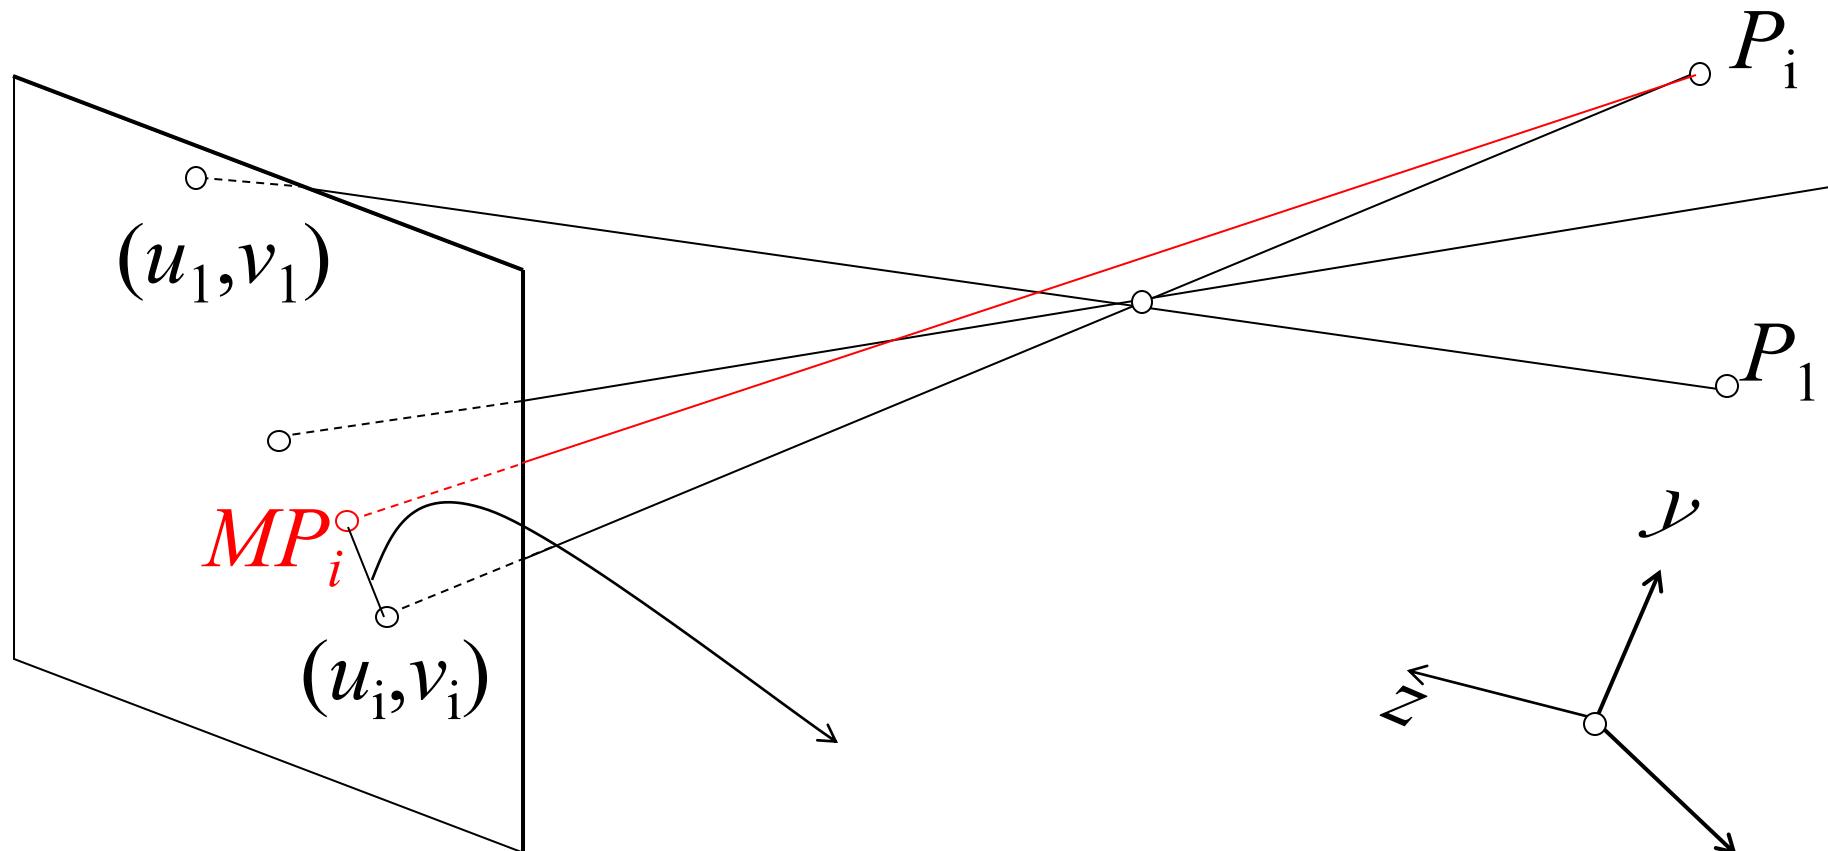
Distorted:

$$x'' = \frac{1}{\lambda} x'$$

$$y'' = \frac{1}{\lambda} y'$$

$$\lambda = 1 + k_1 r^2 + k_2 r^4 + \dots$$

Minimizing reprojection error with radial distortion



Add distortions to
reprojection error:

$$\left(u_i - \frac{1}{\lambda} \frac{m_1 \cdot P_i}{m_3 \cdot P_i} \right)^2 + \left(v_i - \frac{1}{\lambda} \frac{m_2 \cdot P_i}{m_3 \cdot P_i} \right)^2$$

Correcting radial distortion

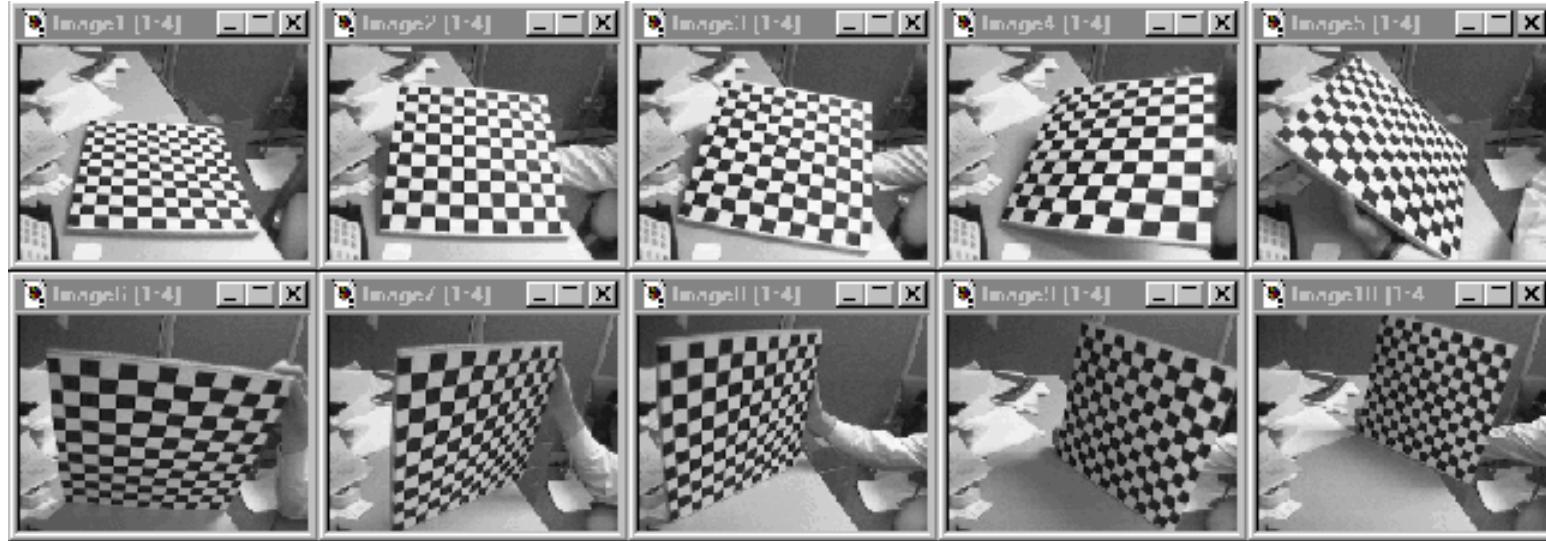


before



after

Alternative: Multi-plane calibration

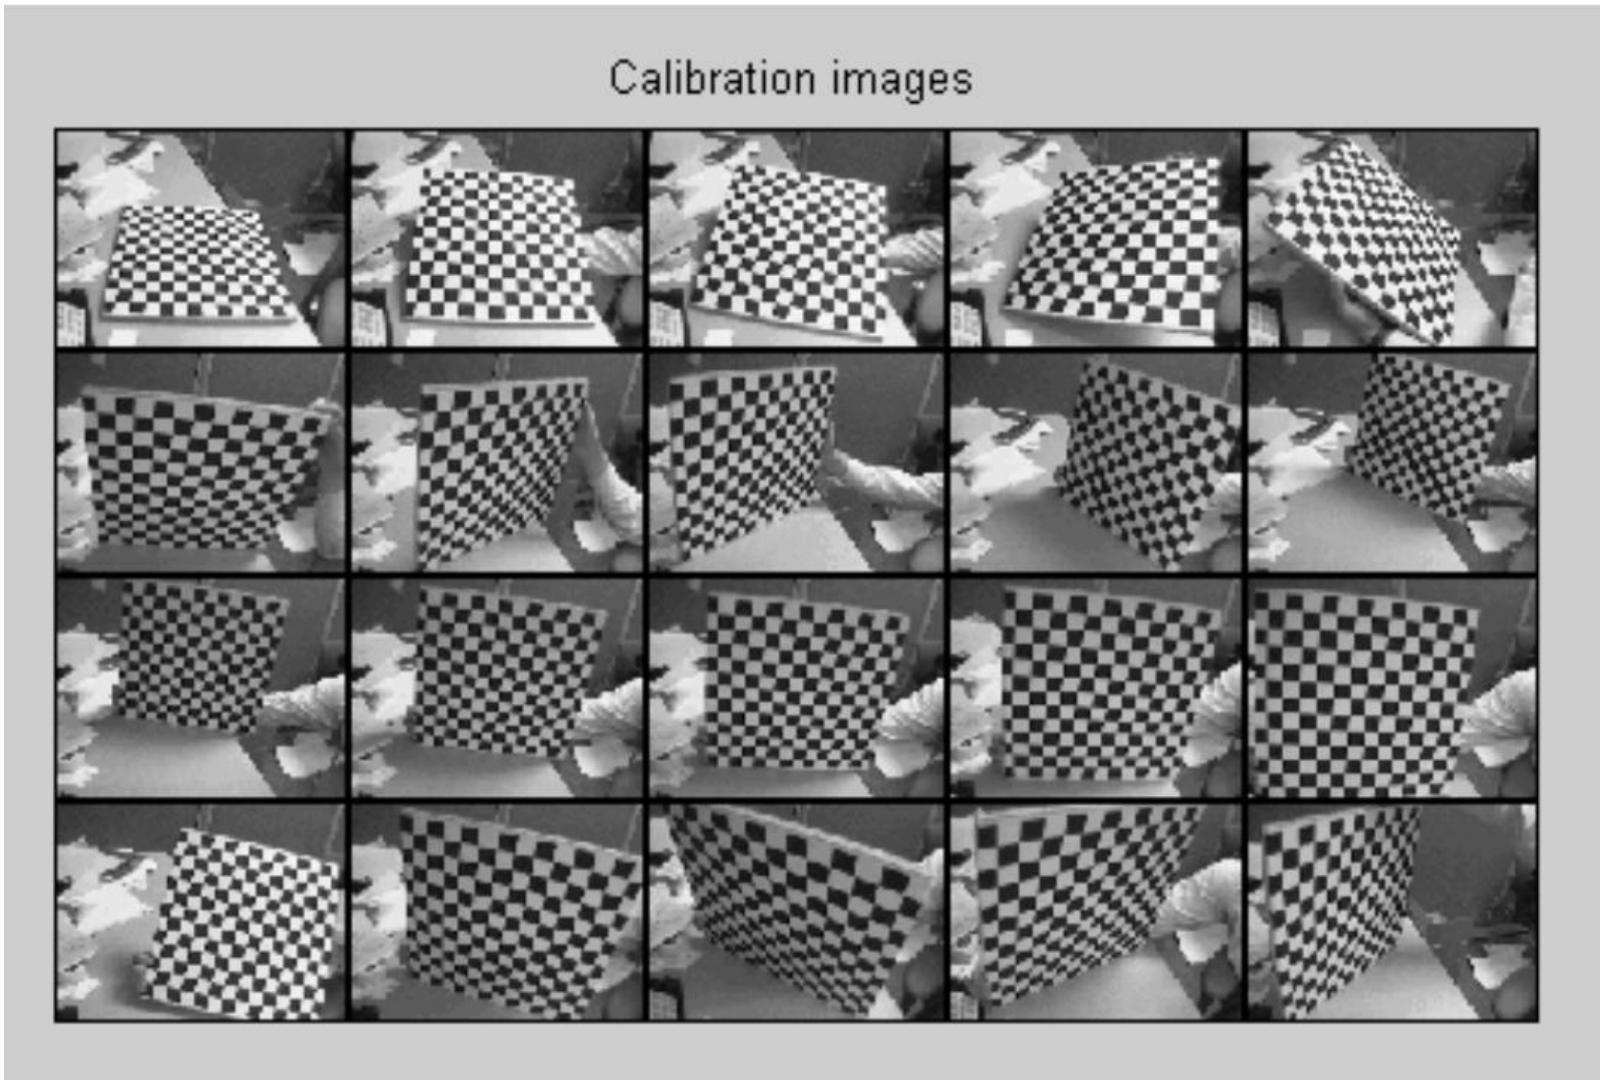


Advantages:

- Only requires a plane
- Don't have to know positions/orientations
- Great code available online!
 - Matlab version: http://www.vision.caltech.edu/bouguetj/calib_doc/index.html
 - Also available on OpenCV.

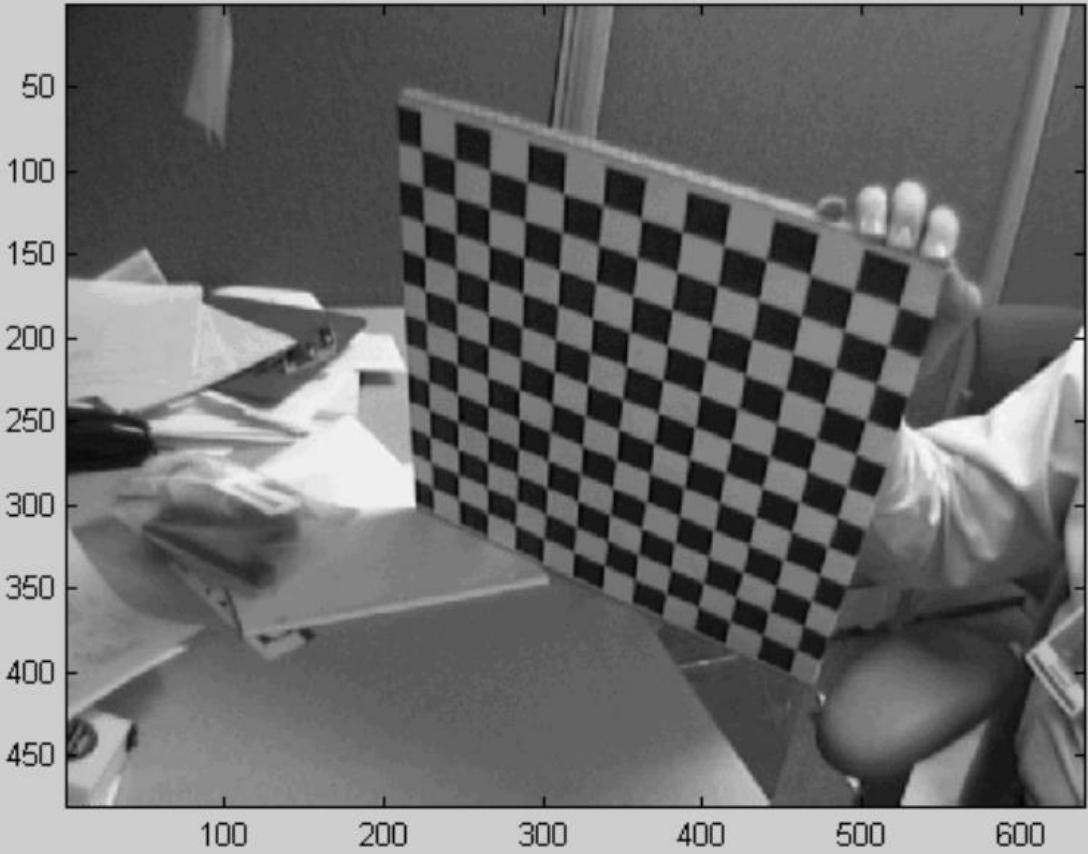
Disadvantage: Need to solve non-linear optimization problem.

Step-by-step demonstration

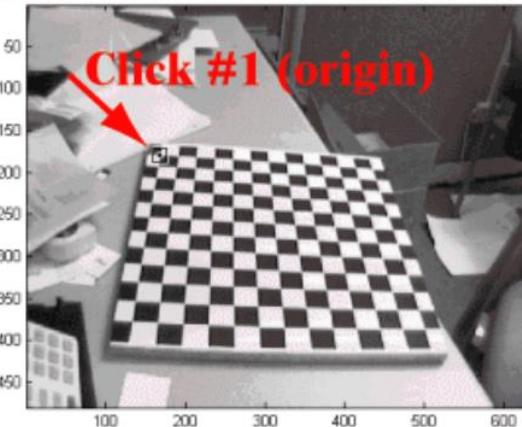


Step-by-step demonstration

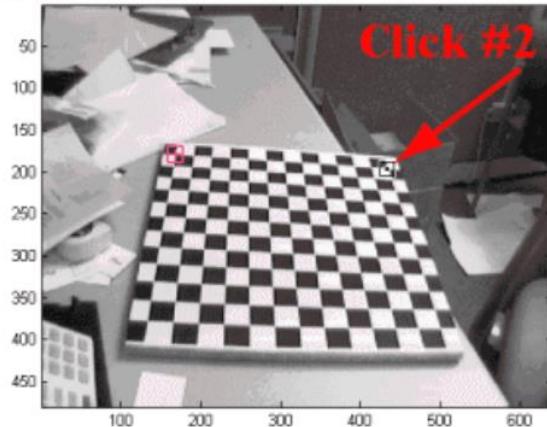
Click on the four extreme corners of the rectangular pattern...



Click on the four extreme corners of the rectangular pattern (first corner = origin)... Image 1



Click on the four extreme corners of the rectangular pattern (first corner = origin)... Image 1



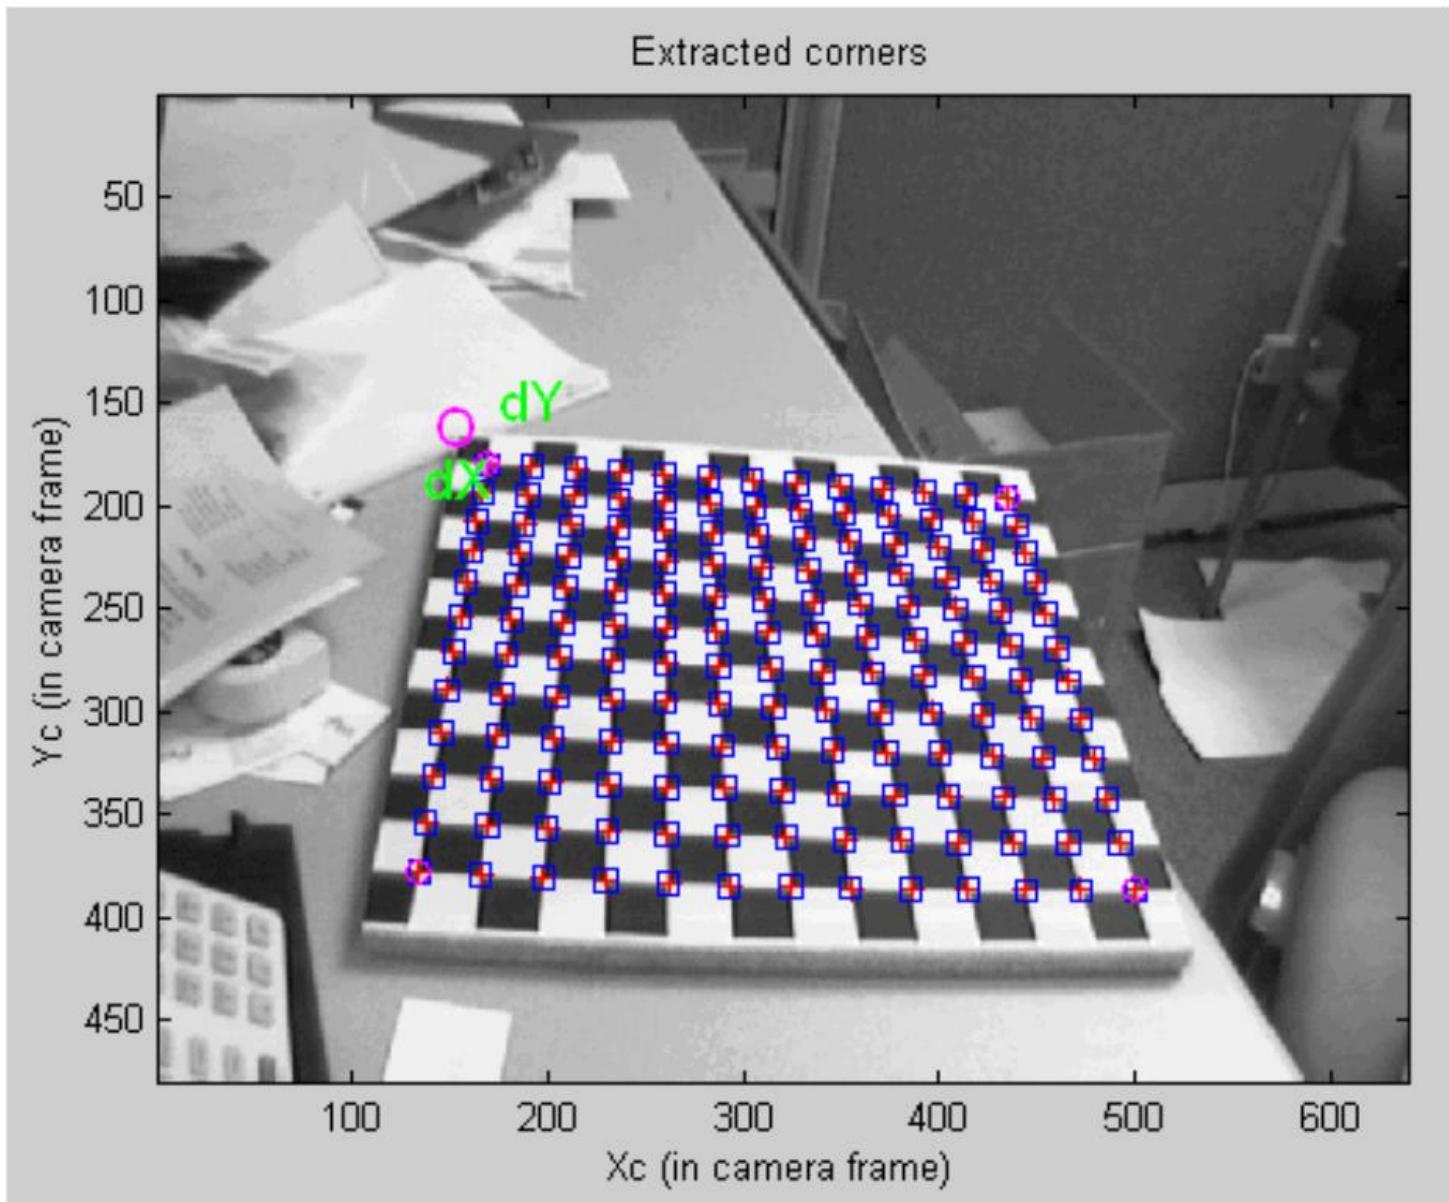
Click on the four extreme corners of the rectangular pattern (first corner = origin)... Image 1



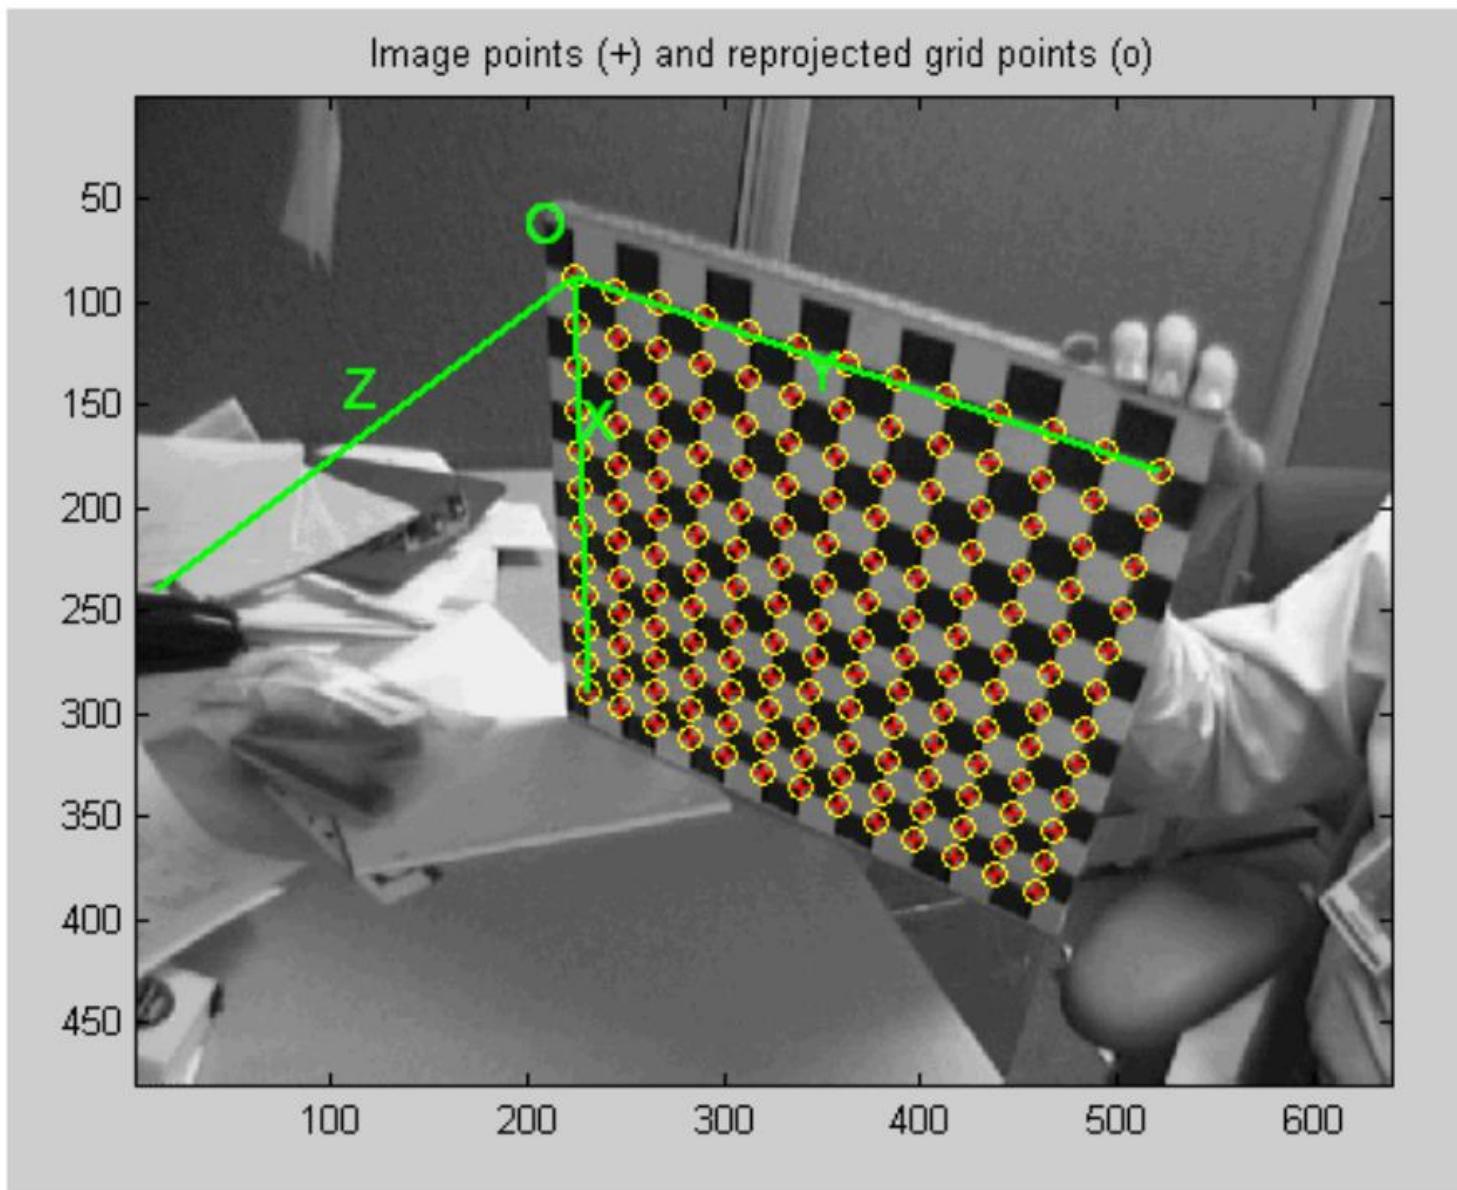
Click on the four extreme corners of the rectangular pattern (first corner = origin)... Image 1



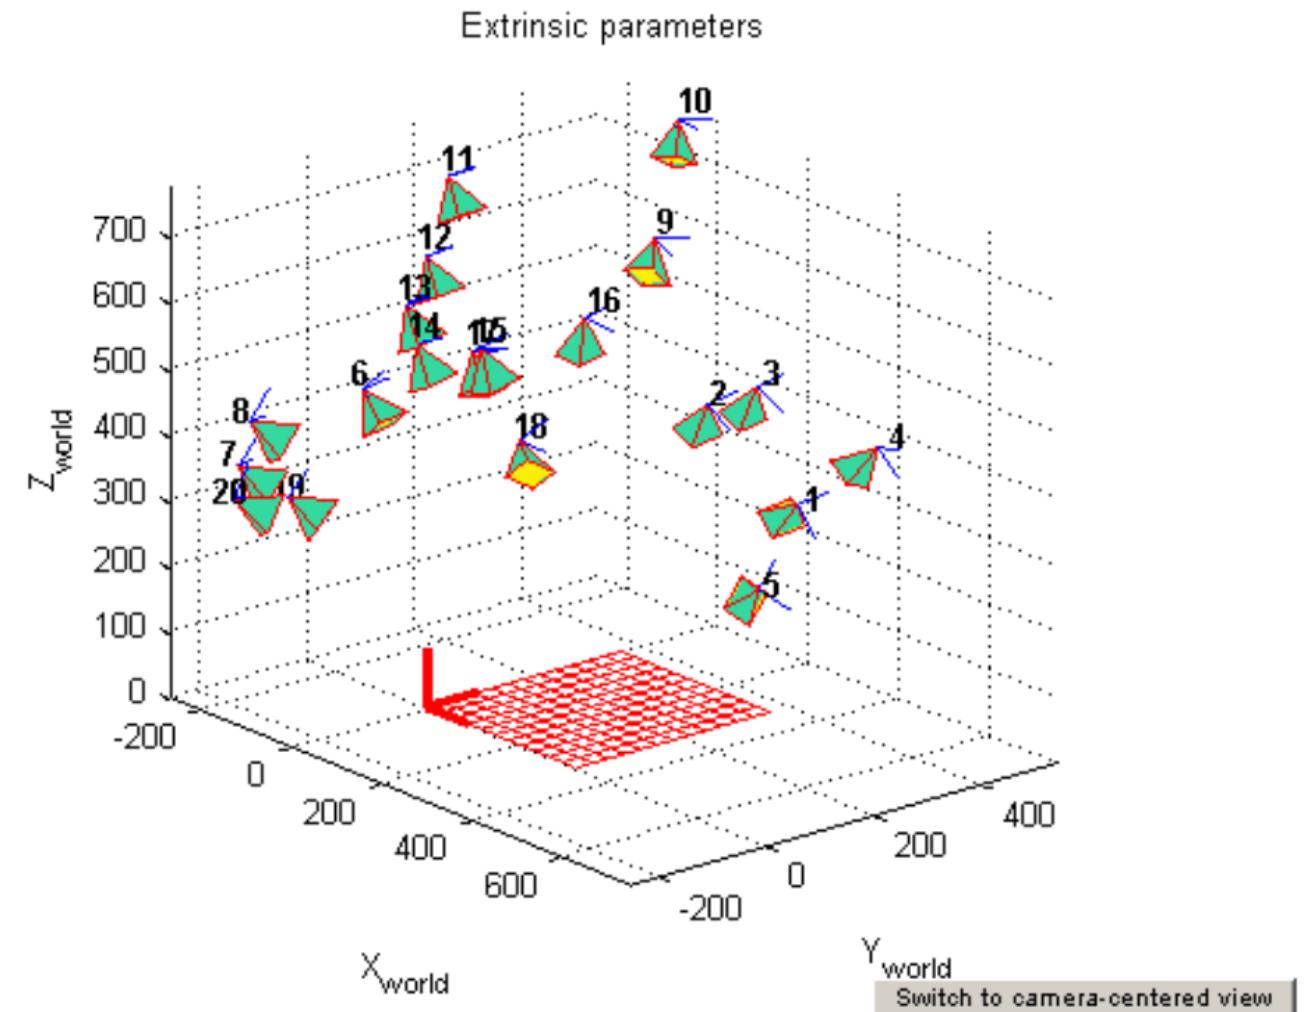
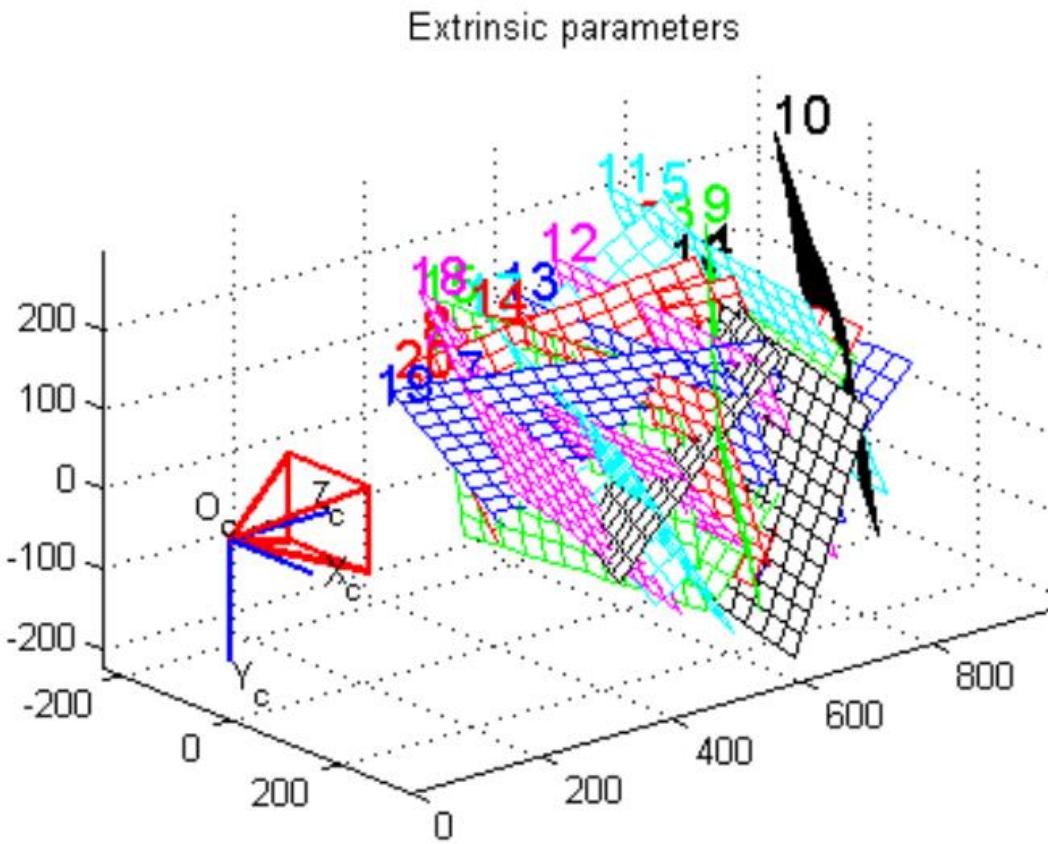
Step-by-step demonstration



Step-by-step demonstration



Step-by-step demonstration



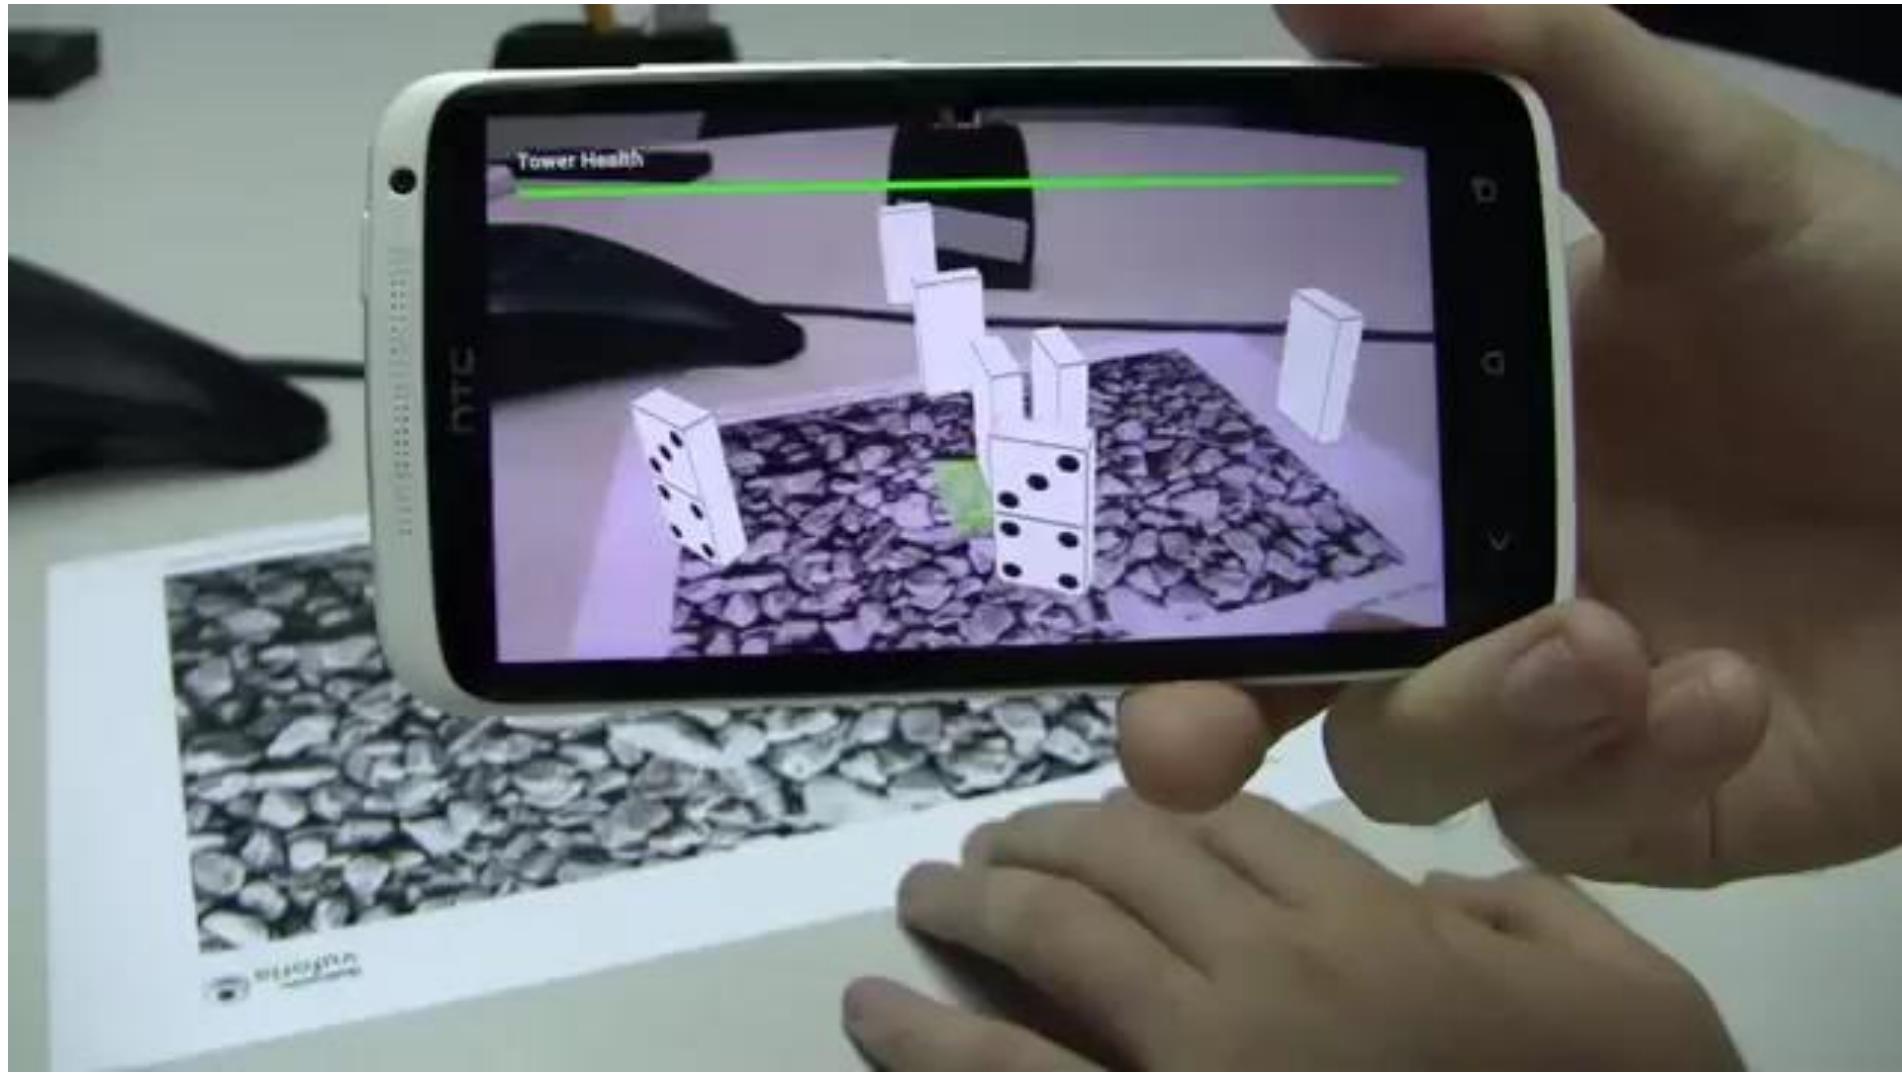
What does it mean to “calibrate a camera”?

What does it mean to “calibrate a camera”?

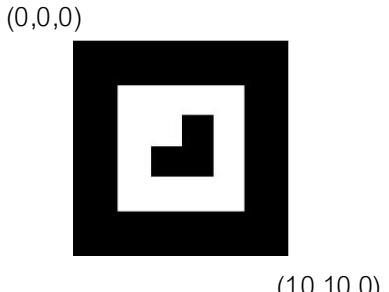
Many different ways to calibrate a camera:

- Radiometric calibration.
- Color calibration.
- Geometric calibration.
- Noise calibration.
- Lens (or aberration) calibration.

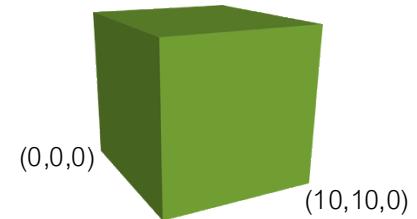
We'll briefly discuss radiometric and color calibration in later lectures. For the rest, see CS 73 (computational photography).



3D locations of planar marker features
are known in advance



3D content prepared in advance



Simple AR program

1. Compute point correspondences (2D and AR tag)
2. Estimate the pose of the camera \mathbf{P}
3. Project 3D content to image plane using \mathbf{P}

